

ERDŐS PROBLEM 1112, FROM FIRST PRINCIPLES: $r_k(d_1, d_2)$ EXISTS IF AND ONLY IF $d_2 \geq k + 1$

JOHAN LAND

ABSTRACT. Let $k \geq 3$ and $1 \leq d_1 < d_2$ be integers. Erdős problem 1112 asks whether there is an integer r such that every lacunary sequence $B = \{b_1 < b_2 < \dots\}$ of positive integers with $b_{i+1} \geq r b_i$ admits a sequence $A = \{a_1 < a_2 < \dots\}$ of positive integers with all gaps in $[d_1, d_2]$ and $(kA) \cap B = \emptyset$, where kA is the k -fold sumset; write $r_k(d_1, d_2)$ for the least such r .

We prove that $r_k(d_1, d_2)$ exists if and only if $d_2 \geq k + 1$ — so the answer to the recorded question is yes precisely when $d_2 \geq k + 1$ — with the explicit bound $r_k(d_1, d_2) \leq 192 d_2$ when it exists. Below the threshold, non-existence holds in a strong form: for every prescribed sequence of pointwise lower bounds on the successive ratios of B , however rapidly growing, a single B meets the k -fold sumset of every admissible A . The existence half uses a Beatty construction whose k -fold sumset threads the gaps a lacunary B must leave; the non-existence half reduces, through a word-combinatorial analysis of the gap sequence, to a bounded subset-sum covering lemma of possible independent interest. The proof is computer-assisted, with a finite certificate layer, and the full dichotomy is formally verified in Lean 4 in the problem’s literal integer phrasing.

This paper is also an experiment in form. It is written to be *read* rather than merely checked: everything the proof needs is built from the ground up, the governing intuition is stated before each construction and separated from the estimates that pay for it, and the argument is illustrated throughout. Nothing is assumed beyond comfort with elementary algebra and some experience reading proofs. §1.1 explains the choice and what it costs.

CONTENTS

Part I. The problem	2
1. The problem, from scratch	2
2. The one picture that explains the threshold	11
Part II. Existence: winning when $d_2 \geq k + 1$	15
3. Floors and Beatty walks	15
4. Where the k -fold sums live	18
5. The free-gap lemma	21
6. Nested intervals, and the ratio $192 d_2$	25
Part III. Non-existence: the reduction	29
7. Interlude: a mini-course on remainders	29

Date: July 2026.

2020 Mathematics Subject Classification. 11B13, 11B75, 11P70, 68V20.

Key words and phrases. Erdős problems, lacunary sequences, sumsets, Sturmian words, subset sums, formal verification, Lean, mathematical exposition.

Source, data, and the Lean 4 development: https://github.com/beetree/math_erdos_1112. §29 gives the release and reproduction instructions; §27 states how this work was produced.

8.	The certificate trick: one B against every A	32
9.	Gap words, tail alphabets, and the plan	34
10.	Eventually periodic gap words	37
11.	Interlude: a mini-course on words, and a little analysis	38
12.	Two letters: the reachable band, and the sweep	39
13.	Balanced words: the Morse–Hedlund classification	44
14.	The Sturmian endgame	50
15.	Three or more letters: the slot lemma	57
Part IV. The subset-sum lemma		60
16.	The subset-sum lemma: statement, sharpness, and the map	60
17.	Machinery: grids, frames, staircases, and the lift	62
18.	Reduction to the hard core	67
19.	Cases D and L: two clean formulas	72
20.	Case P: pairs as movers	73
21.	Case E: the η -box	76
22.	Case T: thin μ , and the first table	78
23.	Case B: finite bases, lifted forever	80
24.	The proof of SHARP, and its sharpness	82
Part V. Assembly, and what the computer checked		83
25.	Assembly: proof of the dichotomy	83
26.	What the computer checked, and how you can too	83
27.	Provenance and authorship	87
28.	Declarations	87
29.	Data availability and reproduction	88
Appendix A. Every expansion, in full		88
Appendix B. Hints and answers for the exercises		89
Appendix C. Symbol glossary		93
Appendix D. The certificate tables		93
Appendix E. Correspondence with the Lean development		96
References		97

Part I. The problem

1. THE PROBLEM, FROM SCRATCH

1.1. On the form of this paper. This paper settles a question posed by Paul Erdős and Ronald Graham in 1980 [ErGr80] and recorded as problem 1112 on erdosproblems.com [E1112], where the case $k \geq 3$ is listed as open. The result is the dichotomy stated as Theorem A in §1.8: the ratio $r_k(d_1, d_2)$ exists precisely when $d_2 \geq k + 1$.

It is written unconventionally, and deliberately so. The conventions of research mathematics reward compression: a proof is presented in the smallest number of symbols that a specialist can expand, the intuition that generated it is discarded once it has done its work, and the reader is trusted to reconstruct what was left out. That is an efficient way to communicate with fifty people and an effective way to exclude everyone else. It also, in our experience, hides errors — the compressed step is exactly where the reader stops checking.

So this paper tries the opposite. Every prerequisite is built in, from remainder arithmetic to the completeness of the real line. The governing idea of each construction is stated in plain language and then *separated* from the estimates that make it work, so that the reader can see which lines carry the mathematics and which merely pay for constants. The argument is illustrated throughout. Where a routine expansion would bury the idea it serves, the expansion goes to an appendix rather than into the reader’s way. The result is roughly three times the length of a conventional write-up of the same theorems.

We think that is a good trade, and we would like to see the experiment repeated by others; but the reader should know what it costs. It costs length. It costs the particular pleasure of a dense argument admired by people who can already read it. And it forgoes the presumption — comfortable but rarely examined — that difficulty of exposition is evidence of depth of content. What it does not cost is rigour: nothing here is asserted without proof, the finite layer is machine-checked three independent ways, and the entire dichotomy is verified in Lean 4 (§26). A conventional presentation of the same results is available in the repository for readers who prefer one.

Some ground rules, so you know what you are reading.

- *Every mathematical subject is defined.* If a word like “sumset” or “lacunary” appears, it was defined earlier in this paper. Three self-contained interludes build the background the proof needs — remainders and divisibility (§7), infinite words and a little analysis (§11), and balanced words (§13). What is *not* redeveloped from zero is the general craft of reading proofs — contradiction, induction, “it suffices to show” — so §1.2 is a short guide to the conventions this paper uses. If you hit an undefined *mathematical term*, that is our mistake; if you hit an unfamiliar *proof move*, §1.2 should name it.
- *Every proof is complete, and one table travels separately.* When a step is genuinely hard, we slow down and build the tools — we do not say “it can be shown”. Where a routine algebraic expansion would drown the idea it serves, the idea stays in the text and the expansion is written out, line by line, in Appendix A; nothing is left to “clearly”. One exception, worth knowing from the start: the finite layer of Part IV rests on 350 certificate rows — six small integers each — and every one of them is printed here, in Appendix D. Each is checkable by hand in a few minutes; the method, and the three independent systems that have verified all of them, are in §§22–26.
- *Exercises are part of the deal, and they are labelled.* Scattered through the text are *Exercises*: small exercises, chosen so that if you can do them, you have understood the last few pages. Every one has a hint or answer in Appendix B, and no later proof secretly depends on your solutions. Each carries its difficulty on its face: ROUTINE (a few minutes — if it does not come, reread), CORE (the idea of the section, worth real time), CHALLENGE (genuinely hard; skipping costs nothing), and EXPLORATION (open-ended, with no single clean answer). Stalling on a CHALLENGE is not evidence you misunderstood the section.
- *Some passages are skippable.* A few stretches of pure verification are marked (*first pass: skippable*) — skip them, keep your momentum, and return when you want the details. The parts can also be read somewhat independently (a map, with three suggested routes, is in §1.10).
- *It is long.* A conventional paper compresses; this one decompresses. The length is the price of the experiment.

Two conventions of the genre are kept, because they exist for reasons that survive scrutiny. §1.9 states precisely what is new here and what is owed to others; and §§26–29 say exactly how the claims can be checked, by whom, and what remains trusted. Clarity is not an excuse for vagueness about credit or verifiability.

One notational convention up front: $\mathbb{N} = \{1, 2, 3, \dots\}$ denotes the positive integers (natural numbers), and when we say *number* without qualification we mean a positive integer.

1.2. How to read this proof. Mathematics has a craft as well as a content, and this paper assumes a little of the craft while building all of the content. Here is the whole of what is assumed, in two pages. Skim it now; come back when a phrase in a proof feels like a move you were not taught.

Direct proof, contradiction, contrapositive. To prove “if P then Q ”, you may assume P and derive Q (direct); or assume P together with “not Q ” and derive something absurd (contradiction); or prove “if not Q then not P ”, which says exactly the same thing (contrapositive). All three appear here, and we name which one we are using whenever it might not be obvious.

“It suffices to show.” A proof often replaces its goal with a stronger or more convenient one: *it suffices to show R* , meaning “ R implies what we want, so let us prove R instead”. This is not a weakening — proving more than you need is always allowed. §4 does this deliberately and says so.

“Fix an arbitrary x .” To prove something about *every* x , we take one, refuse to assume anything special about it, and argue. Because nothing special was assumed, the argument covers all of them. The mirror move, “there exists”, is discharged by *exhibiting* one — which is why so many proofs here consist of writing down a construction and then checking it works.

Induction, ordinary and strong. To prove a statement $S(n)$ for every n : prove $S(1)$, then prove that $S(n)$ implies $S(n + 1)$. *Strong* induction allows the stronger assumption that $S(1), \dots, S(n)$ *all* hold when proving $S(n + 1)$ — used once, in §24, where a statement about a number M is proved using the same statement for all smaller values.

The well-ordering principle. Every nonempty set of positive integers has a smallest element. Obvious-sounding, genuinely load-bearing: it is what licenses phrases like “let n^* be the *first* index such that...”, which appear in §15 and §7.

Chains of inequalities. A display like

$$A \leq B = C \leq D$$

asserts each link separately and concludes $A \leq D$. Read it one link at a time, and ask of each: which hypothesis makes *this* step true? Where a link is not immediate we say why, either inline or in Appendix A.

Mechanism versus bookkeeping. This is the distinction most worth internalizing before Part II. Nearly every proof here has two ingredients: a *mechanism* — the construction or idea that makes the thing work — and *bookkeeping* — the inequalities confirming that the particular constants chosen are large enough. The mechanism is what you should remember; the bookkeeping is what you should be able to check but need not carry in your head. We flag the seam explicitly, and passages of pure verification are marked (*first pass: skippable*).

Passing to a tail. Several arguments discard a finite initial segment of an infinite sequence and work only with what remains — “pass to the tail”. This is harmless when the property in question concerns only large numbers, and we always say why it is harmless when we do it.

What is not assumed. No calculus, no set theory beyond “ $\in$ ” and “ $\subseteq$ ”, no combinatorics, no number theory beyond what §7 builds, no analysis beyond what §11 builds. If you have met proof by induction and can manipulate an inequality without panic, you have enough.

1.3. Sequences and gaps. The two main characters of this story are infinite sequences of positive integers. Let us fix the vocabulary.

Definition 1.1 (Sequence, gaps). A *sequence* is an infinite list of positive integers

$$a_1 < a_2 < a_3 < \cdots,$$

strictly increasing, going on forever. We write $A = \{a_1 < a_2 < \cdots\}$ and think of A interchangeably as a list or as an infinite set of positive integers. The *gaps* of A are the differences between consecutive terms:

$$g_n := a_{n+1} - a_n \quad (n = 1, 2, 3, \dots).$$

Example 1.2. The even numbers $A = \{2, 4, 6, 8, \dots\}$ have all gaps equal to 2. The squares $\{1, 4, 9, 16, 25, \dots\}$ have gaps $3, 5, 7, 9, \dots$, growing without bound. The powers of ten $\{10, 100, 1000, \dots\}$ have gaps $90, 900, 9000, \dots$, growing very fast indeed.

Our first character is a sequence whose gaps are *constrained to a window*: they may wobble, but only between two fixed bounds.

Definition 1.3 ((d_1, d_2) -sequence). Let $d_1 \leq d_2$ be positive integers. A (d_1, d_2) -*sequence* is a sequence A all of whose gaps lie between d_1 and d_2 :


$$d_1 \leq a_{n+1} - a_n \leq d_2 \quad \text{for every } n \geq 1.$$

Example 1.4. Take $(d_1, d_2) = (2, 4)$. Then

$$A = \{1, 3, 7, 9, 13, 15, 19, \dots\}$$

(gaps $2, 4, 2, 4, 2, 4, \dots$) is a $(2, 4)$ -sequence, and so is $\{5, 8, 11, 14, \dots\}$ (all gaps 3), and so is any sequence you build by walking right along the number line taking steps of size 2, 3 or 4, in any order you like, forever.

That last phrasing is the right mental picture: a (d_1, d_2) -*sequence is a walk along the number line whose steps are all between d_1 and d_2* . You choose the starting point and you choose each step, subject only to the window.



Exercise 1.5 (CORE). (a) Give an explicit *rule* generating an infinite $(2, 4)$ -sequence whose gap pattern never settles into a repeating cycle, and list the first six resulting terms. (Six terms alone cannot certify non-periodicity — the rule is the answer.) (b) Show that if A is a (d_1, d_2) -sequence then $a_1 + (n - 1)d_1 \leq a_n \leq a_1 + (n - 1)d_2$ for every n . (Add up the gaps.)

1.4. Sumsets. Our second ingredient is addition: we will be interested not in A itself but in the set of all numbers you can make by adding up k members of A .

Definition 1.6 (k -fold sumset). Let A be a sequence and $k \geq 1$ an integer. The k -*fold sumset* of A is

$$kA := \{x_1 + x_2 + \cdots + x_k : x_1, \dots, x_k \in A\},$$

where the x_i need *not* be distinct — you may use the same member of A several times.




FIGURE 1. A sumset, seen. Above: the $(2, 4)$ -walk A of Example 1.4. Below: its 3-fold sumset $3A$, every element a sum of three members of A with repetitions allowed (dashed lines trace one such sum). The whole problem is whether the lower row can be steered to miss a given sparse sequence.

Two warnings about this (standard) notation. First, kA is *not* the set $\{k \cdot a : a \in A\}$ of multiples — though it contains all of those (take $x_1 = \dots = x_k = a$). Second, repetitions are allowed: for $A = \{1, 3, 7, \dots\}$ and $k = 3$, the sums $1 + 1 + 1 = 3$ and $1 + 1 + 7 = 9$ and $3 + 7 + 7 = 17$ all belong to $3A$.

Example 1.7. Let $A = \{5, 8, 11, 14, \dots\}$, the numbers of the form $5 + 3j$ with $j \geq 0$. What is $3A$? Each element is a sum

$$(5 + 3j_1) + (5 + 3j_2) + (5 + 3j_3) = 15 + 3(j_1 + j_2 + j_3),$$

and since $j_1 + j_2 + j_3$ can be any integer ≥ 0 , we get exactly $3A = \{15, 18, 21, 24, \dots\}$: again an arithmetic progression with the same common difference 3, starting at 15. This tidiness is special to arithmetic progressions — and, as we will see much later, it is precisely what makes them *bad* at dodging.

Exercise 1.8 (ROUTINE). (a) For $A = \{1, 3, 7, 9, 13, \dots\}$ of Example 1.4, list every element of $3A$ up to 25. (b) Show that for any sequence A , the smallest element of kA is ka_1 . (c) Show that if A is a (d_1, d_2) -sequence, then between consecutive “diagonal” elements ka_n and ka_{n+1} of kA the distance is at most kd_2 .

1.5. Lacunary sequences. The sequence to be dodged grows *multiplicatively*: each term is at least a fixed multiple of the one before.

Definition 1.9 (Lacunary sequence). Let $r \geq 1$ be an integer. A sequence $B = \{b_1 < b_2 < \dots\}$ is *lacunary with ratio r* if

$$b_{i+1} \geq r b_i \quad \text{for every } i \geq 1.$$

Example 1.10. The powers $\{r, r^2, r^3, \dots\}$ are lacunary with ratio r — the minimal example, growing as slowly as the definition allows. But a lacunary sequence need not be this regular: $\{1, 7, 100, 1000, 10^{12}, \dots\}$ is lacunary with ratio 7, as long as every step multiplies by at least 7. The word *lacunary* comes from the Latin *lacuna*, a gap or hollow: on the number line, such a sequence is almost all gap.

The picture to hold: a lacunary sequence with a large ratio is *extremely sparse*. Below any number N it has at most about $\log_r N$ terms. Between consecutive terms b_i and $b_{i+1} \geq r b_i$ yawns an empty stretch r times longer than everything that came before. Sparseness is what should make B *easy* to dodge — that is the intuition the problem tests, and Figure 2 is what that intuition looks like.




FIGURE 2. The powers of ten, drawn twice. On a linear axis they look like almost nothing at all — every term is dwarfed by the next. On a logarithmic axis they are perfectly evenly spaced. Read the picture carefully: the powers of ten are the *densest* sequence a ratio of 10 permits, and they mark the boundary. A general ratio-10 lacunary sequence is at least this widely separated on the log axis and may be far more irregular — lacunarity is a lower bound on separation, not a spacing law. What the picture does show is that no *scale* is special: sparseness of this kind offers the walk no foothold, which is why it will turn out to be no help below the threshold.

Exercise 1.11 (ROUTINE). Let B be lacunary with ratio $r \geq 2$. Show that $b_i \geq b_1 r^{i-1}$ for all i , and conclude that B has at most $1 + \log_2 N$ terms below any $N \geq b_1$.

Exercise 1.12 (ROUTINE). Show that if $r' \geq r$, then every sequence that is lacunary with ratio r' is also lacunary with ratio r . Deduce that if announcing r wins, so does announcing any $r' \geq r$ — which is the parenthetical in Theorem A(1). (One line, but it is the line that turns Theorem 6.1's single ratio into a half-line of them.)


1.6. The question, as a game. We can now state the problem. Here it is verbatim, as recorded on erdosproblems.com as problem 1112:

Let $1 \leq d_1 < d_2$ and $k \geq 3$. Does there exist an integer r such that if $B = \{b_1 < \dots\}$ is a lacunary sequence of positive integers with $b_{i+1} \geq rb_i$ then there exists a sequence of positive integers $A = \{a_1 < \dots\}$ such that $d_1 \leq a_{i+1} - a_i \leq d_2$ for all $i \geq 1$ and $(kA) \cap B = \emptyset$, where kA is the k -fold sumset?

Every phrase in this is now something we know. But sentences with three quantifiers (“there exists $r \dots$ such that if [every] $B \dots$ then there exists A ”) are hard to hold in the head, so let us restate it as a *game* between you and an Adversary. The numbers $k \geq 3$ and $1 \leq d_1 < d_2$ are fixed in advance and known to both players.

- (1) **You** announce an integer r — once and for all.
- (2) **The Adversary**, knowing r , builds a lacunary sequence B with ratio r : an infinite sequence with $b_{i+1} \geq rb_i$. The Adversary is infinitely clever and builds B specifically to hurt you.
- (3) **You**, knowing B completely, build a (d_1, d_2) -sequence A : a walk with steps in $[d_1, d_2]$.

You win if kA avoids B entirely: no sum of k members of A ever equals a member of B . In symbols, $(kA) \cap B = \emptyset$.



The row of symbols underneath is how a logician would compress the timeline: *there exists* an r such that *for every* B *there exists* an A that wins. Whenever a quantifier statement threatens to blur, come back to this timeline: later moves see earlier moves, never the reverse.

The problem asks: *do you have a winning strategy?* That is, is there a single choice of r such that, whatever B the Adversary produces, you can answer with a winning A ?

The order of moves is the soul of the statement, so let us dwell on it.

- You choose A *after* seeing all of B . You may tailor your walk to the particular sequence you must dodge. (If instead you had to choose A first, the Adversary would simply pick some element of kA — it is an infinite set — and start B there. You would lose always, and the problem would be boring.)
- But you choose r *before* seeing B . One fixed r must defeat every lacunary-ratio- r sequence at once. Making r enormous makes the Adversary's sequence extremely sparse — intuitively, easier to dodge. The question is whether *any* r , however large, is enough.

Why would this be hard? Your walk A is dense — steps at most d_2 , so you have lots of members — and it is only kA , the k -fold sums, that must dodge B . The tension: A 's gaps are *bounded*, so kA is forced to be rather evenly spread, and an evenly spread set has a hard time missing even a very sparse one, because “evenly spread” means it shows up everywhere at every scale. Whether the k -fold sums have enough slack to be steered around a sparse B — that is exactly what the problem is asking, and the answer will turn out to hinge on a precise comparison between k and d_2 .

1.7. A first fight: tiny r loses. To get blood flowing, let us play the game and prove something: announcing $r = 1$ loses. This tiny result will not be needed later, but it makes the rules concrete, and it is our first complete proof.

Proposition 1.13. *Let $k \geq 3$ and $1 \leq d_1 < d_2$. If you announce $r = 1$, the Adversary wins: there is a lacunary sequence B with ratio 1 such that every (d_1, d_2) -sequence A has $(kA) \cap B \neq \emptyset$.*

Proof. With $r = 1$ the lacunarity condition reads $b_{i+1} \geq b_i$, which any increasing sequence satisfies. So the Adversary may play

$$B = \{1, 2, 3, 4, 5, \dots\},$$

all positive integers. Now let A be any (d_1, d_2) -sequence whatsoever. The sumset kA is not empty — it contains, for instance, ka_1 (add a_1 to itself k times), which is a positive integer, hence a member of B . So $(kA) \cap B \neq \emptyset$: the Adversary wins. $\square$

So any winning r must be at least 2, and the real question starts there: the problem page asks about “an integer r ”, and we now see this means, in effect, an integer $r \geq 2$.

Exercise 1.14 (CORE). Suppose you announce $r = 2$ and, feeling lazy, you resolve to always play an arithmetic progression: $A = \{c, c + d, c + 2d, \dots\}$ for some starting point c and fixed gap $d \in [d_1, d_2]$ of your choosing. (a) Show that for each *single* such progression, the Adversary — if they knew which one you would play — could build a legal B meeting kA . (Hint: Example 1.7 computes kA exactly for one progression; imitate it, and let the Adversary pick elements of kA far apart.) (b) The Adversary must in fact commit to B *before* seeing your progression — can one B defeat every progression at once? This is genuinely harder, and it is the seed of Part III's “certificate trick”; for now, just convince yourself that (a) alone does not settle it.

1.8. The answer. Here is the theorem this paper proves. It answers the question completely, and the answer is a genuine dichotomy: a sharp line through parameter space, with a winning strategy on one side and a losing verdict on the other.

Theorem A (The dichotomy). Let $k \geq 3$ and $1 \leq d_1 < d_2$ be integers.

- (1) (Existence.) If $d_2 \geq k + 1$, you have a winning strategy: announcing $r = 192 d_2$ works — every lacunary B with that ratio can be dodged by some (d_1, d_2) -sequence A with $(kA) \cap B = \emptyset$. (Any larger announcement works too, for a reason worth pausing on: raising r only *shrinks* the set of sequences the Adversary may legally play, so a strategy that beats them all at $r = 192d_2$ beats the survivors at any $r \geq 192d_2$. This is Exercise 1.12.)
- (2) (Non-existence, in a strong form.) If $d_2 \leq k$, no announcement wins. Worse: even if the fixed ratio r is replaced by any prescribed sequence $r_1, r_2, r_3, \dots$ of lower bounds — the Adversary now being required to satisfy $b_{i+1} \geq r_i b_i$ at every step, with the r_i growing as fast as you like — the Adversary can still build a *single* sequence B that defeats *every* (d_1, d_2) -sequence A at once.

So the answer to Erdős problem 1112 is: *yes precisely when $d_2 \geq k + 1$.*

Pause on how lopsided the two conclusions are. Above the threshold, one fixed ratio $192 d_2$ tames every lacunary B . Below it, the collapse is total

B can rescue the walker. Second, the combinatorial lemma the non-existence half consumes at its binding corner, which we isolate as Theorem 16.1 in Part IV and discuss immediately below. Third, a complete machine-checked verification (§26).

The subset-sum lemma, against the literature. Theorem 16.1 says: every finite set G of at least three positive integers with $\gcd(G) = 1$ and $\max(G) = M$ admits a multiset of at most $M - 1$ of its elements whose subset sums contain M consecutive integers. We use it only as a tool, but it may be of independent interest, so we place it against three traditions.

Dense-set subset-sum theorems. The results producing long intervals or progressions among subset sums — Sárközy [Sa89, Sa94], Lev [Lev03], Szemerédi–Vu [SV06], Conlon–Fox–Pham [CFP21], Alon–Freiman [AF88], and constructively Chen–Mao–Zhang [CMZ25] — all work with a *dense* set ($|A| \gtrsim \sqrt{n}$) and yield an arithmetic progression whose common difference need not be 1. The closest antecedent is Lev’s consecutive-subset-sum theorem [Lev98, Thm. 1], which does give a genuine interval of consecutive integers — but again only for a dense set ($|A| \gtrsim \frac{2}{3}M$), which is vacuous when $|G| = 3$. Our lemma is the *density-free, multiset* analogue: it trades density for repetition, covering an interval of length exactly M at uniform cost $|S| \leq M - 1$, from an arbitrary gcd-one set, down to three elements.


Higher-sumset structure. The structure theorems for m -fold sumsets (Nathanson [Na72], Granville–Walker [GW21], Lev [Lev22]) are density-free but describe mG for large m — a total-degree simplex of coefficients, of which the subset sums of a bounded multiset form a per-generator box inside — so their conclusion is a weaker statement about a larger set.

loss of thread, and Part IV is the steepest sustained stretch. Exercises are the footholds: every one has a hint or answer in Appendix B, routine expansions live in Appendix A, and Appendix C is the map of the notation.

Two parameter sets, carried throughout. Rather than meet a fresh example in every section, it helps to keep two concrete cases in your pocket and ask what each new result says about them:

$$\boxed{k = 3, (d_1, d_2) = (1, 4)} \quad (\text{above the threshold: you win}) \quad \boxed{k = 3, (d_1, d_2) = (1, 3)} \quad (\text{at the corner: you lose})$$

The winning case is the one §2.2 pictures, §5 computes with ($t = 10000$), and §6.4 plays to a finish (against $b_i = 800^i$, yielding $\gamma^* \approx 3.74999556$ and the walk $804, 808, 812, \dots$). The losing case is where the two-letter analysis bites: $d_2 = 3 = k$ is the corner, so every admissible walk's gap word uses letters from $\{1, 2, 3\}$, the sweep needs width ≥ 2 , and the Sturmian endgame handles what is left — and if the walk's tail alphabet has three letters, the slot lemma consumes SHARP at exactly $m(G) \leq M - 1$



windows $(3.5s - 3, 3.5s]$ for $k = 3$, $\gamma = 3.5$, $c = 0$; circled integers sit in daylight

FIGURE 3. The window picture for the walk of Example 2.1. Each 3-fold sum lands in a shaded window; the integers 18 and 25 fall in the daylight between windows, so they are *unreachable*: no three members of A sum to them. (The window $(18, 21]$ is open at its left end, which is why 18 escapes.)

rounds down by less than 1 — $\lfloor x \rfloor$ always lies in the half-open interval $(x - 1, x]$ — so k floors together round down by less than k :




FIGURE 4. The dichotomy in one triptych ($k = 3$ throughout): windows of width 3 at stride γ . Above the threshold ($\gamma > k$) daylight recurs forever; at $\gamma = k$ the windows tile the line exactly; below ($\gamma < k$) they overlap. The theorem says the picture tells the truth: the game is winnable precisely when a stride $\gamma > k$ is available, i.e. when $d_2 \geq k + 1$. (One caution the proof will respect: daylight makes dodging *possible*, never automatic — integer strides

Example 2.4 (Steady walks are individually beatable). Let A be the arithmetic progression $a_i = c + id$ with integer gap $d \leq k$. A k -fold sum is

$$(c + i_1 d) + \cdots + (c + i_k d) = kc + sd, \quad s = i_1 + \cdots + i_k,$$

and s ranges over *all* integers $\geq k$. So

$$kA = \{ kc + sd : s \geq k \} :$$

from $kc + kd$ onwards, kA contains *every* number that leaves the same remainder as kc upon division by d — in the standard notation, every number $\equiv kc \pmod{d}$. (Read “ $x \equiv y \pmod{m}$ ” as “ x and y leave the same remainder when divided by m ”; this *remainder arithmetic* is developed properly in §7, and until then the informal reading is all we use.) We say kA contains a *tail of a congruence class*. A set this fat cannot dodge: however sparse the Adversary is forced to be, the class contains arbitrarily




FIGURE 5. The proof as a map: arrows mean “feeds into”. The four debts of the honest ledger annotate the boxes that pay them. Parts III and IV open with this map’s relevant half redrawn in words (§9.3, §16.2).

- (4) *The corner $d_2 = k$ is razor-thin.* When the largest step equals k exactly, the windows touch without overlapping, and the walk has just barely too little room; converting “barely” into a proof costs a genuinely new piece of finite mathematics about subset sums, stated in §16 and proved across Part IV.

One more promissory note. In Example 2.4 the offset “ kc ” depended on the walk, and in gap (2) we waved at defeating all offsets and classes simultaneously; when we do this properly, the notion of *tail-covering* will be defined to absorb all of it. From here on, the document simply executes the ledger: Part II closes gap (1); Part III closes gaps (2) and (3), reducing the corner case to a finite lemma; Part IV proves that lemma, closing gap (4); Part V assembles the theorem and describes the computer verification. Figure 5 is the same ledger drawn as a map — worth a bookmark: every later section is one of its boxes.

Part II. Existence: winning when $d_2 \geq k + 1$

3. FLOORS AND BEATTY WALKS

Part II proves the existence half of the dichotomy: if $d_2 \geq k + 1$, announcing $r = 192d_2$ wins the game. The plan was sketched in §2.2; from here on we execute it rigorously. **Standing assumptions for all of Part II:** k, d_1, d_2 are integers with

$$k \geq 2, \quad 1 \leq d_1 < d_2, \quad d_2 \geq k + 1.$$

(The problem poses $k \geq 3$, but the existence proof works verbatim from $k = 2$, so we prove the slightly more general statement at no extra cost. Note $d_2 \geq 3$ follows.)

This section builds the two tools everything else stands on: the floor function, and the fact that a Beatty walk with slope between $d_2 - 1$ and d_2 is a legal (d_1, d_2) -sequence.

3.1. The floor function.

Definition 3.1 (Floor). For a real number x , the *floor* $\lfloor x \rfloor$ is the unique integer m satisfying

$$m \leq x < m + 1.$$

So $\lfloor 3.5 \rfloor = 3$, $\lfloor 7 \rfloor = 7$, and (watch this one) $\lfloor -0.5 \rfloor = -1$: the floor is the integer just *below* x (or x itself), not “the integer part”. That an integer $m \leq x < m+1$ exists is visually clear — walk down the number line from above x until you first pass an integer — and we take it as known; but uniqueness deserves its one-line proof: if $m \leq x < m+1$ and $m' \leq x < m'+1$ with, say, $m < m'$, then $m+1 \leq m'$ (they are integers), so $x < m+1 \leq m' \leq x$, a contradiction.

Everything we ever use about floors is in the next lemma.

Lemma 3.2 (Floor toolbox). *For all real x, y and every integer n :*

- (1) $x - 1 < \lfloor x \rfloor \leq x$;
- (2) $\lfloor x + n \rfloor = \lfloor x \rfloor + n$;
- (3) if $x \leq y$ then $\lfloor x \rfloor \leq \lfloor y \rfloor$;
- (4) for reals $x_1, \dots, x_k$: $x_1 + \dots + x_k - k < \lfloor x_1 \rfloor + \dots + \lfloor x_k \rfloor \leq x_1 + \dots + x_k$;
- (5) every real interval $[u, v]$ with $v - u \geq 1$ contains an integer.

Proof. Write $m = \lfloor x \rfloor$, so $m \leq x < m+1$ by definition.

(1) The right inequality is $m \leq x$ directly; the left is $x < m+1$ rearranged.

(2) The integer $m+n$ satisfies $m+n \leq x+n < (m+n)+1$, which is exactly the defining property of $\lfloor x+n \rfloor$; by uniqueness, $\lfloor x+n \rfloor = m+n$.

(3) Suppose $x \leq y$. Then $\lfloor x \rfloor \leq x \leq y$, so $\lfloor x \rfloor$ is an integer $\leq y$. But $\lfloor y \rfloor$ is the *greatest* integer $\leq y$: any integer $m' \leq y$ has $m' \leq \lfloor y \rfloor$, since otherwise $m' \geq \lfloor y \rfloor + 1 > y$. Hence $\lfloor x \rfloor \leq \lfloor y \rfloor$.

(4) Add the k instances of (1): summing $x_t - 1 < \lfloor x_t \rfloor \leq x_t$ over $t = 1, \dots, k$ gives the claim. (This is Exercise 2.2, now on the official record.)

(5) The integer $m := \lfloor u \rfloor + 1$ satisfies $u < m \leq u+1 \leq v$. □

Exercise 3.3 (CORE). (a) Compute $\lfloor 5.999 \rfloor$, $\lfloor -3 \rfloor$, $\lfloor -3.01 \rfloor$. (b) Show by example that $\lfloor x+y \rfloor$ is *not* always $\lfloor x \rfloor + \lfloor y \rfloor$ — and prove that it always equals $\lfloor x \rfloor + \lfloor y \rfloor$ or $\lfloor x \rfloor + \lfloor y \rfloor + 1$. (This is the fact Lemma 3.5 runs on; the proof there does it for you, so try it first.) (c) Where does the proof of Lemma 3.2(2) break if n is not an integer?

3.2. Beatty walks, officially.

Definition 3.4 (Beatty walk). For a real number $\gamma > 1$ and an integer $c \geq 0$, the *Beatty walk* with slope γ and offset c is the sequence

$$A_{\gamma,c} := \{c + \lfloor i\gamma \rfloor : i = 1, 2, 3, \dots\}, \quad \text{i.e. } a_i = c + \lfloor i\gamma \rfloor.$$

We met $A_{3.5,0} = \{3, 7, 10, 14, 17, 21, \dots\}$ in Example 2.1; Figure 6 shows the geometry: the walk is the line $y = \gamma i$ digitized downward to the integer grid, and the gaps are the vertical jumps of the resulting staircase.

One more, with a slope that is not a half-integer: for $\gamma = \frac{10}{3}$ and $c = 0$,

$$\lfloor \frac{10}{3} \rfloor = 3, \quad \lfloor \frac{20}{3} \rfloor = 6, \quad \lfloor \frac{30}{3} \rfloor = 10, \quad \lfloor \frac{40}{3} \rfloor = 13, \dots$$

so $A_{10/3,0} = \{3, 6, 10, 13, 16, 20, 23, 26, 30, \dots\}$, with gap pattern 3, 4, 3, 3, 4, 3, 3, 4, $\dots$: the gap 4 arrives once every three steps, keeping the average speed at exactly $\frac{10}{3}$.

The single structural fact we need is that a Beatty walk’s gaps take at most two values, and we can say which ones.

Lemma 3.5 (Two-value gaps). *Let $\gamma > 1$ be real. Every gap of the Beatty walk $A_{\gamma,c}$ lies in $\{\lfloor \gamma \rfloor, \lfloor \gamma \rfloor + 1\}$:*

$$\lfloor (i+1)\gamma \rfloor - \lfloor i\gamma \rfloor \in \{\lfloor \gamma \rfloor, \lfloor \gamma \rfloor + 1\} \quad \text{for every } i \geq 1.$$




FIGURE 6. The Beatty walk $a_i = \lfloor 3.5i \rfloor$ as a digitized line: each dot sits at the integer just below the line $y = 3.5i$, and consecutive dots differ by $\lfloor \gamma \rfloor = 3$ or $\lfloor \gamma \rfloor + 1 = 4$ — the two-value gap phenomenon of Lemma 3.5, seen rather than computed.

Proof. The offset c cancels from every gap, so consider $g := \lfloor (i+1)\gamma \rfloor - \lfloor i\gamma \rfloor$. Write $x := i\gamma$ and $y := \gamma$, so that $g = \lfloor x + y \rfloor - \lfloor x \rfloor$; the claim is then exactly Exercise 3.3, which we now record as part of the toolbox:

$$\lfloor x + y \rfloor - \lfloor x \rfloor \in \{ \lfloor y \rfloor, \lfloor y \rfloor + 1 \} \quad \text{for all real } x, y.$$

Proof of the display. Split each number into integer and fractional parts, $x = \lfloor x \rfloor + \{x\}$ and $y = \lfloor y \rfloor + \{y\}$, with $\{x\}, \{y\} \in [0, 1)$. Then

$$x + y = (\lfloor x \rfloor + \lfloor y \rfloor) + (\{x\} + \{y\}),$$

where the first bracket is an integer and the second lies in $[0, 2)$. If $\{x\} + \{y\} < 1$ the floor of the whole is $\lfloor x \rfloor + \lfloor y \rfloor$; if $\{x\} + \{y\} \geq 1$ it is $\lfloor x \rfloor + \lfloor y \rfloor + 1$ (Lemma 3.2(2), moving the integer bracket through the floor). Subtracting $\lfloor x \rfloor$ gives the display, and substituting $y = \gamma$ gives the lemma. $\square$

(If γ is an integer, only the value $\lfloor \gamma \rfloor = \gamma$ actually occurs, and the walk is an arithmetic progression — Exercise below. For non-integer γ , both values occur infinitely often, though we will not need that.)

3.3. Beatty walks are admissible. Now aim the slope at the top of our permitted window.

Proposition 3.6 (Admissibility). *Let γ be any real number with $d_2 - 1 < \gamma < d_2$, and $c \geq 0$ any integer. Then $A_{\gamma, c}$ is a (d_1, d_2) -sequence.*

Proof. Three things to check: the terms are positive, the sequence is strictly increasing, and the gaps lie in $[d_1, d_2]$.

Since $d_2 - 1 < \gamma < d_2$ and $d_2 - 1, d_2$ are integers, the defining property of the floor gives $\lfloor \gamma \rfloor = d_2 - 1$. By Lemma 3.5, every gap of $A_{\gamma, c}$ is $d_2 - 1$ or d_2 . Both lie in $[d_1, d_2]$: indeed

$d_1 \leq d_2 - 1$ because $d_1 < d_2$ are integers. As $d_2 \geq k + 1 \geq 3$, every gap is at least $2 > 0$, so the sequence is strictly increasing; and its first term is $a_1 = c + \lfloor \gamma \rfloor \geq 0 + (d_2 - 1) \geq 2 > 0$. All three checks pass. $\square$

Two remarks worth pausing on. First, *the proofs of Part II never ask whether the slope is rational*: the final winning slope will simply be some real number produced by the nested-interval theorem, and no argument here tests its rationality. (This is not the same as irrationality being *irrelevant*: a closing remark in §6 shows that against the strongest adversary the produced slope is *forced* to be irrational — the proof just never needs to check.) Second, notice how little of the freedom of (d_1, d_2) -sequences we are using: only steps $d_2 - 1$ and d_2 , the two largest. The lower bound d_1 plays no role beyond $d_1 \leq d_2 - 1$, which is automatic. This is the promised explanation of why d_1 is absent from the threshold: the existence strategy never has any reason to take a small step.

Exercise 3.7 (ROUTINE). (a) Show that if $\gamma = d$ is an integer, $A_{d,c}$ is the arithmetic progression $c + d, c + 2d, c + 3d, \dots$ (b) For $\gamma = \frac{7}{2}$, which i have gap $g_i = 4$ and which have 3? Find the pattern and prove it. (Hint: $\lfloor (i+1)\gamma \rfloor - \lfloor i\gamma \rfloor$ depends on whether i is even or odd.) (c) Show that for any real $\gamma > 1$, among the first n gaps of $A_{\gamma,c}$, the number of gaps equal to $\lfloor \gamma \rfloor + 1$ is $\lfloor (n+1)\gamma \rfloor - \lfloor \gamma \rfloor - n\lfloor \gamma \rfloor$. (Telescope the sum of the gaps.)

4. WHERE THE k -FOLD SUMS LIVE

This section turns the window picture of §2 into theorems, and then performs the change of viewpoint on which the whole existence proof runs: instead of asking *which numbers* the walk's sums can reach, we ask *which dial positions* γ are safe for a given target. By the end of the section, dodging a single element of B will be a clean geometric condition on γ alone.

4.1. The window containment.

Proposition 4.1 (Windows). *Let $\gamma > 1$ be real, $c \geq 0$ an integer, and $k \geq 1$. Then*

$$kA_{\gamma,c} \subseteq \bigcup_{s \geq k} (kc + s\gamma - k, kc + s\gamma],$$

the union over integers $s \geq k$.

Proof. A member of $kA_{\gamma,c}$ is a sum of k terms of the walk, say with indices $i_1, \dots, i_k \geq 1$ (repetitions allowed):

$$\sum_{t=1}^k (c + \lfloor i_t \gamma \rfloor) = kc + \sum_{t=1}^k \lfloor i_t \gamma \rfloor.$$

Set $s := i_1 + \dots + i_k$; since each $i_t \geq 1$, we have $s \geq k$. By the sum-of-floors bound (Lemma 3.2(4)) applied to the numbers $x_t = i_t \gamma$, whose sum is $s\gamma$,

$$s\gamma - k < \sum_{t=1}^k \lfloor i_t \gamma \rfloor \leq s\gamma,$$

so the member lies in $(kc + s\gamma - k, kc + s\gamma]$. $\square$

Alongside the windows we record where the sums *start*: nothing in $kA_{\gamma,c}$ is small.

Lemma 4.2 (Sums start high). *If $d_2 - 1 < \gamma < d_2$, the smallest element of $kA_{\gamma,c}$ is $k(c + d_2 - 1)$. In particular every integer $b \leq kc$ satisfies $b \notin kA_{\gamma,c}$: targets at or below kc are dodged automatically.*

Proof. Every term of the walk is at least its first term $a_1 = c + \lfloor \gamma \rfloor = c + d_2 - 1$, so every k -fold sum is at least $ka_1 = k(c + d_2 - 1)$ — and this value is attained, by summing a_1 with itself k times. Finally $k(c + d_2 - 1) > kc$, since $d_2 \geq 2$. $\square$

4.2. The avoidance criterion. Fix a target $b > kc$ that we wish to keep out of $kA_{\gamma,c}$, and write

$$t := b - kc \geq 1$$

for the *shifted target*. Subtracting kc throughout, Proposition 4.1 says: if $b \in kA_{\gamma,c}$, then $t \in (s\gamma - k, s\gamma]$ for some integer $s \geq k$ — equivalently, $s\gamma \in [t, t + k)$ for some $s \geq k$. Turning this around gives the criterion we will use forever after.

Definition 4.3 (Certified-safe slope). Let $t \geq 1$ be an integer. A real $\gamma > 0$ is *certified safe* for t (*safe* for short) if

$$s\gamma \notin [t, t + k] \quad \text{for every integer } s \geq 1.$$

The word *certified* is doing honest work: the condition is a checkable guarantee, deliberately *stronger* than the mere fact of avoidance — a slope can dodge b without earning the certificate, and we will see that happen. First, the guarantee itself:

Proposition 4.4 (Avoidance criterion). *Let $d_2 - 1 < \gamma < d_2$, let $c \geq 0$, and let $b > kc$ be an integer. If γ is certified safe for $t = b - kc$, then $b \notin kA_{\gamma,c}$.*

Proof. Contrapositive. If $b \in kA_{\gamma,c}$, the discussion above produces an integer $s \geq k \geq 1$ with $s\gamma \in [t, t + k) \subseteq [t, t + k]$, so γ is not certified safe for t . $\square$

And see it succeed once, before any fine print. Take $k = 3$, $d_2 = 4$, and the target $t = 46$: the slope $\gamma = 3.8$ is certified safe, because the products $s\gamma$ march upward in steps of 3.8 and the two nearest the window are

$$12 \cdot 3.8 = 45.6 < 46 \quad \text{and} \quad 13 \cdot 3.8 = 49.4 > 49,$$

with every smaller s falling shorter and every larger s overshooting further. So for any offset c , the number $b = 46 + 3c$ is *provably* absent from $3A_{3.8,c}$: the criterion converts one two-line check into an exclusion from an infinite sumset.

Now look closely: Definition 4.3 demands *more* than the proof of Proposition 4.4 needs, in two ways, and both are deliberate.

- It forbids $s\gamma \in [t, t + k]$ for *all* $s \geq 1$, though only $s \geq k$ can actually occur in a window. Ruling out extra values of s is a stronger demand on γ — harmless for the conclusion, and it frees every later argument from tracking which s are realizable.
- It uses the *closed* interval $[t, t + k]$, though the windows are half-open. Again this only makes “safe” harder to achieve, never easier — and closed intervals are more comfortable to work with (their endpoints belong to them, so “the least s with...” arguments have no edge cases).

The cost of this generosity is real but affordable: some slopes that would in fact dodge b are not certified safe. You will meet a concrete instance in Exercise 4.5 — the walk of Example 2.1 really does miss 25, yet the criterion declines to vouch for it. We simply will not care: the free-gap lemma of §5 finds slopes that pass even the stricter test.

Big idea. For a Beatty walk, “the sum of k members never equals b ” has become a statement about the dial position alone: no multiple $s\gamma$ may land in the closed window $[t, t + k]$, where $t = b - kc$. The game of dodging an infinite sequence B is now a game of placing one real number γ clear of infinitely many forbidden zones.



FIGURE 7. The unsafe set U_{100} for $k = 3$, on the stretch $2.95 \leq \gamma \leq 4.18$ of the dial: the interval $[\frac{100}{s}, \frac{103}{s}]$ for $s = 25, 26, \dots, 34$. Safe gaps (two are marked) open up between consecutive unsafe intervals, wider toward the right; toward the left the intervals crowd together and, below $\gamma = 3 = k$, merge into a solid unsafe block (Proposition 4.6).

Exercise 4.5 (CORE). Take the walk of Example 2.1: $\gamma = 3.5$, $c = 0$, $k = 3$. (a) Show that γ is *not* safe for $t = 25$: find the integer s with $3.5s \in [25, 28]$. (b) Yet $25 \notin 3A$ (Exercise 2.3). Explain, using the half-open window $(25, 28]$, why the point $3.5 \cdot 8 = 28$ landing exactly on the closed window's right endpoint does not put 25 into $3A$ — so the criterion's verdict is merely conservative, not wrong. (c) Stronger: show that *no* integer $t \geq 1$ is certified at $\gamma = 3.5$. (Consider the smallest multiple u of 3.5 with $u \geq t$; multiples of 3.5 are integers or half-integers, so show $u \leq t + 3$.) So at this particular slope the criterion certifies nothing at all — even though the walk misses 18, 25, 32, $\dots$. The slopes the free-gap lemma will select are precisely engineered to avoid this fate. (d) The text certified $\gamma = 3.8$ for $t = 46$. Find *all* certified slopes in $(3, 4)$ for that target, by listing the unsafe intervals $[\frac{46}{s}, \frac{49}{s}]$ that meet $(3, 4)$.

4.3. The dial, and its unsafe zones. Fix the shifted target t and look at the condition of Definition 4.3 from the dial's point of view. For each single $s \geq 1$,

$$s\gamma \in [t, t+k] \iff \gamma \in \left[\frac{t}{s}, \frac{t+k}{s}\right],$$

so the slopes ruled out by index s form a closed interval, and γ is safe for t exactly when it avoids the *unsafe set*

$$U_t := \bigcup_{s \geq 1} \left[\frac{t}{s}, \frac{t+k}{s}\right].$$

As s grows, these unsafe intervals march down the dial toward 0, shrinking as they go (both endpoints are proportional to $1/s$). Figure 7 shows the portion of U_{100} near the dial region we shop in, for $k = 3$.

Two features of this picture are the pivot of the whole story. First, the safe gaps — the white space between consecutive unsafe intervals — *exist* in the region $\gamma > k$, and we will learn to find them (§5). Second, below $\gamma = k$ there is nothing to find:

Proposition 4.6 (No safety below k). *Let $t \geq 1$ be an integer. Every real γ with $0 < \gamma \leq k$ is unsafe for t .*

Proof. We must find an integer $s \geq 1$ with $s\gamma \in [t, t+k]$, i.e. with

$$s \in \left[\frac{t}{\gamma}, \frac{t+k}{\gamma}\right].$$

This is a real interval of length $k/\gamma \geq 1$, so it contains an integer s by Lemma 3.2(5). That integer is positive, since $s \geq t/\gamma > 0$. $\square$

This is the threshold, seen from the dial. A walk whose stride is at most k cannot be certified to dodge *anything*: every single target's unsafe set swallows the entire dial region

$(0, k]$. Our standing assumption $d_2 \geq k + 1$ is exactly what places the shopping district $(d_2 - 1, d_2)$ of Proposition 3.6 inside the living zone $\gamma > k$. (That the picture “no safety below k ” hardens into the actual non-existence theorem below the threshold — not merely a failure of this particular criterion — is the business of Parts III and IV.)

What remains for the existence half is quantitative: given the target t , *how wide* are the safe gaps near the top of the dial, and how little of the dial does each new target destroy? That is the free-gap lemma, to which we now turn.

Exercise 4.7 (CORE). Let $k = 3$ and $t = 20$. (a) List the unsafe intervals $[\frac{20}{s}, \frac{23}{s}]$ that meet the dial stretch $(3, 4)$, and conclude that the safe slopes there form exactly the two gaps $(\frac{23}{7}, \frac{20}{6})$ and $(\frac{23}{6}, 4)$. (b) Pick a slope in one of them and write down the first eight terms of a $(1, 4)$ -walk none of whose 3-fold sums equals $20 + 3c$ for your chosen offset c .

Exercise 4.8 (ROUTINE). Show that the unsafe intervals for indices s and $s + 1$ overlap or touch — that is, $\frac{t+k}{s+1} \geq \frac{t}{s}$ — if and only if $s \geq t/k$. (This is the precise sense in which the unsafe intervals “merge into a solid block” at the left of Figure 7: from $s = \lceil t/k \rceil$ onward, consecutive intervals chain together without gaps.)

5. THE FREE-GAP LEMMA

We now build the engine of the existence proof. The task left over from §4.3: given a target t , find a safe stretch of dial — and find it *quantitatively*, because we will have to do this infinitely often, once per element of B , each time inside whatever room the previous rounds left us. The lemma of this section says: inside any dial interval that is not too narrow (compared to $1/t$), there is a safe subinterval that is not too much narrower. The two “not too”s are what the constant 192 will be made of in §6.

One more tool for the toolbox first.

Definition 5.1 (Ceiling). For a real x , the *ceiling* $\lceil x \rceil$ is the unique integer m with $m - 1 < x \leq m$ — the least integer $\geq x$. (Uniqueness: as for the floor. Note $\lceil x \rceil = \lfloor x \rfloor$ when x is an integer, and $\lfloor x \rfloor + 1$ otherwise.)

5.1. The statement.

Lemma 5.2 (Free gap). Let $t \geq 32d_2^2$ be an integer, and let $I = [lo, hi]$ be a closed real interval with

$$I \subseteq [d_2 - \frac{1}{2}, d_2] \quad \text{and} \quad |I| = hi - lo \geq \frac{4d_2^2}{t}.$$

Then I contains a closed subinterval I' with

$$|I'| \geq \frac{d_2}{48t}$$

such that every $\gamma \in I'$ is safe for t (Definition 4.3): $s\gamma \notin [t, t + k]$ for every integer $s \geq 1$.

Read the shape of this before the proof. The input interval must have width at least $4d_2^2/t$; the output is guaranteed width at least $d_2/(48t)$. Both scale like $1/t$: a bigger target is dodged with proportionally finer dial adjustments. The ratio between them,

$$\frac{4d_2^2/t}{d_2/(48t)} = 192d_2,$$

is the number of times the output of one round fits into the input the next round demands — hold that thought until §6, where it becomes the theorem’s constant.

A word on the constants before they arrive, because none of them is magic. The proof needs exactly one quantitative fact: the gap J it will name has width $t - k s_1$ over roughly s_1^2 , and since $s_1 \approx t/d_2$ while $k \leq d_2 - 1$, the numerator should be about s_1 -sized — there is a whole “spare 1” in $d_2 - k$. Every fraction with a 32 in it below $(\frac{15}{32}, \frac{17}{32})$, and the hypothesis $t \geq 32d_2^2$ is just a concrete way of banking part of that spare 1 so that the chain of inequalities closes with room at each step; any similar choices would work and would only change the final 192. *Throughout the proof, a bracketed line [like this] tracks one running instance: $d_2 = 4$, $k = 3$, $t = 10000$, $I = [3.5, 4]$ — so every symbolic step has a number attached.*

5.2. The idea: name one gap, count nothing. The unsafe set $U_t = \bigcup_{s \geq 1} [t/s, (t+k)/s]$ is an infinite union, and a first instinct is to estimate how many of its intervals meet I and how much room they eat. That accounting can be made to work, but there is something better: we can point at *one specific gap* between two *specific* consecutive unsafe intervals, and prove that this single gap already avoids every unsafe interval at once, has the width we need, and sits inside I . No counting, no union bounds — three clean checks on one explicitly named object.

Which gap? Look at Figure 7 again: near the top of the dial the unsafe intervals are the sparsest and the gaps the widest, so we aim as high in I as possible. Let

$$s_1 := \left\lceil \frac{t}{hi} \right\rceil \quad (\text{the least integer } s \text{ with } t/s \leq hi),$$

so that the unsafe interval of index s_1 is the highest-placed one whose left endpoint has not yet escaped past the top of I . Our gap is the space just below that interval and just above the next one:

$$L := \frac{t+k}{s_1+1}, \quad R := \frac{t}{s_1}, \quad J := (L, R)$$

— the open interval between the right end of unsafe interval $s_1 + 1$ and the left end of unsafe interval s_1 .

[Instance: $s_1 = \lceil 10000/4 \rceil = 2500$, $R = 10000/2500 = 4$, $L = 10003/2501 \approx 3.99960$; so $J \approx (3.99960, 4)$.]

5.3. The mechanism, alone. Before any estimate, here is the entire idea, as a lemma that mentions no constants at all.

Lemma 5.3 (Geometric free gap). *Let $t \geq 1$ and let $s_1 \geq 1$ be any integer. Then every γ in the open interval*

$$J = (L, R), \quad L := \frac{t+k}{s_1+1}, \quad R := \frac{t}{s_1}$$

is certified safe for t — provided J is nonempty.

Proof. Let $\gamma \in J$ and let $s \geq 1$ be any integer; we show $s\gamma \notin [t, t+k]$, splitting on the size of s .

- If $s \leq s_1$: then $t/s \geq t/s_1 = R > \gamma$, so $s\gamma < t$. The product undershoots the window.
- If $s \geq s_1 + 1$: then $(t+k)/s \leq (t+k)/(s_1+1) = L < \gamma$, so $s\gamma > t+k$. The product overshoots it.

□

That is the whole invention, and it is worth pausing over: the unsafe intervals $[t/s, (t+k)/s]$ march monotonically leftward as s grows, so the single gap between indices s_1 and s_1+1 is automatically clear of *all* of them — those with small s lie entirely to its right, those with large s entirely to its left. **One comparison per side replaces an infinite enumeration.**




FIGURE 8. Anatomy of the free gap. The unsafe intervals $[t/s, (t+k)/s]$ march leftward and shrink as s grows, so naming the single gap J between indices s_1 and s_1+1 dodges *all* of them at once — no enumeration. The three checks of the proof are the three annotations: J is safe, J is wide enough that its closed middle third I' still has width $\geq d_2/(48t)$, and J lies inside the dial room I we still own. Numerically, for $d_2 = 4$, $k = 3$, $t = 10000$, $I = [3.5, 4]$: $s_1 = 2500$, $R = 4$, $L = 10003/2501 \approx 3.99960$, so $|J| = 1/2501 \approx 4.0 \times 10^{-4}$ — some 16 times the promised $d_2/(16t) = 2.5 \times 10^{-5}$.

[Instance: at $\gamma = 3.9998 \in J$, index $s = 2500$ gives $s\gamma = 9999.5 < 10000$, and $s = 2501$ gives $10003.5 > 10003$; every smaller s undershoots further, every larger s overshoots further.]

Everything that follows is bookkeeping: choosing s_1 so that J is nonempty, lands inside I , and is wide enough to keep the construction alive. If you are reading for the idea, you now have it, and may rejoin at §6.

5.4. The audit: size and placement. (*This subsection is where every constant in the lemma comes from — and where the 192 of the main theorem is ultimately made. It proves nothing new about the mechanism.*)

Throughout, remember the standing facts: $2 \leq k \leq d_2 - 1$ (from $d_2 \geq k + 1$), $d_2 \geq 3$, $t \geq 32d_2^2$, and $d_2 - \frac{1}{2} \leq lo \leq hi \leq d_2$. We take

$$s_1 := \left\lceil \frac{t}{hi} \right\rceil$$

as in §5.2, so that Lemma 5.3 supplies safety the moment we know $J \neq \emptyset$ — which check (ii) below establishes by giving J a positive width.

We first pin down the size of s_1 . Its defining property gives the two-sided estimate

$$\frac{t}{hi} \leq s_1 < \frac{t}{hi} + 1. \quad (1)$$

From the left half and $hi \leq d_2$:

$$s_1 \geq \frac{t}{d_2} \geq \frac{32d_2^2}{d_2} = 32d_2. \quad (2)$$

From the right half and $hi \geq d_2 - \frac{1}{2} \geq \frac{d_2}{2}$ (using $d_2 \geq 1$):

$$s_1 + 1 < \frac{t}{hi} + 2 \leq \frac{2t}{d_2} + 2 \leq \frac{2t}{d_2} + \frac{t}{d_2} = \frac{3t}{d_2} \leq \frac{5t}{d_2}, \quad (3)$$

where $2 \leq t/d_2$ holds because $t/d_2 \geq 32d_2 \geq 96$.

Check (i): every point of J is safe for t . This is Lemma 5.3, already proved — it holds for any choice of s_1 whatever, and needs no hypothesis on t , I or the constants.

Check (ii): J is wide enough: $|J| = R - L \geq \frac{d_2}{16t}$. Put the difference over a common denominator:

$$R - L = \frac{t}{s_1} - \frac{t+k}{s_1+1} = \frac{t(s_1+1) - (t+k)s_1}{s_1(s_1+1)} = \frac{t - k s_1}{s_1(s_1+1)}.$$

So the width question is: how much does t exceed $k s_1$? *The goal:* a bound of the form $t - k s_1 \geq c s_1$ for some definite constant $c > 0$. *Why it should hold:* $s_1 \approx t/hi \approx t/d_2$ while $k \leq d_2 - 1$, so $k s_1 \approx t(d_2 - 1)/d_2$, leaving a margin of about $t/d_2 \approx s_1$ — a whole s_1 , of which we only need a fixed fraction. *The choice:* we take $c = \frac{15}{32}$, picked so that the chain below closes against the hypothesis $t \geq 32 d_2^2$ with a little room to spare. Now the chain. Using $s_1 < t/hi + 1 \leq t/(d_2 - \frac{1}{2}) + 1$ and $k \leq d_2 - 1$:

$$\left(k + \frac{15}{32}\right) s_1 \leq \left(d_2 - \frac{17}{32}\right) \left(\frac{t}{d_2 - \frac{1}{2}} + 1\right) = t \frac{d_2 - \frac{17}{32}}{d_2 - \frac{1}{2}} + \left(d_2 - \frac{17}{32}\right).$$

The fraction equals $1 - \frac{1/32}{d_2 - 1/2}$, so the right side is

$$t - \frac{t}{32(d_2 - \frac{1}{2})} + \left(d_2 - \frac{17}{32}\right) \leq t,$$

provided the middle term dominates the last, i.e. provided $\frac{t}{32(d_2 - 1/2)} \geq d_2 - \frac{17}{32}$. And it does: from $t \geq 32 d_2^2$,

$$\frac{t}{32(d_2 - \frac{1}{2})} \geq \frac{d_2^2}{d_2 - \frac{1}{2}} \geq d_2 + \frac{1}{2} > d_2 - \frac{17}{32},$$

where the middle inequality is $d_2^2 \geq (d_2 + \frac{1}{2})(d_2 - \frac{1}{2}) = d_2^2 - \frac{1}{4}$. Altogether $(k + \frac{15}{32}) s_1 \leq t$, i.e.

$$t - k s_1 \geq \frac{15}{32} s_1.$$

Feeding this into the width formula, then using (3):

$$R - L \geq \frac{\frac{15}{32} s_1}{s_1(s_1+1)} = \frac{15}{32} \cdot \frac{1}{s_1+1} \geq \frac{15}{32} \cdot \frac{d_2}{5t} = \frac{3}{32} \cdot \frac{d_2}{t} \geq \frac{d_2}{16t}.$$

[Instance: $R - L = 4 - 10003/2501 = 1/2501 \cdot (4 \cdot 2501 - 10003) = 1/2501 \approx 4.0 \times 10^{-4}$, comfortably above $d_2/(16t) = 4/160000 = 2.5 \times 10^{-5}$.]

Check (iii): J sits inside I . The top is easy: $R = t/s_1 \leq hi$ by the defining property of s_1 . For the bottom we need $L \geq lo$, and we get it by showing J hangs only a little below hi . First, R is close below hi : from $s_1 < t/hi + 1$,

$$R = \frac{t}{s_1} > \frac{t}{t/hi + 1} = \frac{hi t}{t + hi} = hi \left(1 - \frac{hi}{t + hi}\right) > hi - \frac{hi^2}{t} \geq hi - \frac{d_2^2}{t}.$$

Second, J itself is short: using (1),

$$R - L = \frac{t - k s_1}{s_1(s_1+1)} \leq \frac{t}{s_1^2} \leq \frac{t}{(t/hi)^2} = \frac{hi^2}{t} \leq \frac{d_2^2}{t}.$$

Adding: $L = R - (R - L) > hi - 2d_2^2/t$. Finally the hypothesis $hi - lo \geq 4d_2^2/t$ says exactly that $2d_2^2/t \leq \frac{1}{2}(hi - lo)$, so

$$L > hi - \frac{1}{2}(hi - lo) \geq lo.$$

[Instance: $R = 4 = hi$, and $L \approx 3.99960 > 4 - 2 \cdot 16/10000 = 3.9968 > 3.5 = lo$: the whole of J hangs just under the top of I .]

Conclusion. Check (ii) gives $|J| > 0$, so J is a nonempty open interval and Lemma 5.3 applies to it. Let I' be the closed middle third of J . Then $I' \subseteq J \subseteq I$ by check (iii); every $\gamma \in I'$ is certified safe for t by check (i); and $|I'| = |J|/3 \geq d_2/(48t)$ by check (ii). (This proves Lemma 5.2.) $\square$

Why the middle third, and not J itself? Because the lemma promises a *closed* subinterval — the next round of the construction will want to apply the lemma again *inside* I' , and closed intervals are what the nested-interval theorem of §6 consumes. The open gap J safely contains its closed middle third, and a third is the cheapest way to buy closedness with an explicit constant.

Exercise 5.4 (ROUTINE). The bracketed instance lines did most of the numerical work; three things remain. (a) Check the instance satisfies the hypotheses ($t \geq 32d_2^2$ and $|I| \geq 4d_2^2/t$). (b) How many unsafe intervals meet $I = [3.5, 4]$ when $t = 10000$? (Estimate — then appreciate that the proof never needed to know.) (c) Re-run the whole instance yourself for $t = 10^6$.

Exercise 5.5 (CORE). Where does the proof use $t \geq 32d_2^2$? Find all three places, and test the lemma's limits: for $d_2 = 4$, $k = 3$, find a small t for which the interval J defined above pokes out of $[d_2 - \frac{1}{2}, d_2]$ or has the wrong width.

6. NESTED INTERVALS, AND THE RATIO $192d_2$

We can now assemble the existence half. Here is the theorem, stated for the record.

Theorem 6.1 (Existence). *Let $k \geq 2$ and $1 \leq d_1 < d_2$ with $d_2 \geq k + 1$. Then announcing $r = 192d_2$ wins the game: for every lacunary B with ratio $192d_2$ there is a (d_1, d_2) -sequence A with $(kA) \cap B = \emptyset$.*

The plan: fix the Adversary's B , choose a good offset c , convert each element of B into a shifted target, and then apply the free-gap lemma over and over — each application shrinking our interval of candidate slopes, each shrunken interval still wide enough for the next application *because B is lacunary*. A limit point of the shrinking intervals is a slope safe for every target at once.

6.1. One more axiom: the nested interval theorem. The construction produces an infinite chain of intervals, each inside the last, and needs a point lying in all of them. That such a point exists is a fact about the real number line itself.

Nested interval theorem. If $I_0 \supseteq I_1 \supseteq I_2 \supseteq \dots$ are nonempty closed bounded intervals of real numbers, then there is at least one real number γ^* belonging to every I_j .

Why believe it? Write $I_j = [lo_j, hi_j]$. The left endpoints $lo_0 \leq lo_1 \leq lo_2 \leq \dots$ form a nondecreasing sequence, bounded above (by hi_0): it keeps rising yet can never pass hi_0 , so it must pile up against some limiting value γ^* — and making “pile up against some value” precise is exactly what completeness asserts. This γ^* is $\geq$ every lo_j and $\leq$ every hi_j (if $\gamma^* > hi_{j_0}$ for some j_0 , then some lo_j would eventually exceed $hi_{j_0} \geq hi_j$, impossible). The honest fine print: the existence of that limiting value is the *completeness* of the real numbers — the defining property that separates $\mathbb{R}$ from the rationals, where the theorem is false (nested intervals can squeeze down onto $\sqrt{2}$, which a rational line does not contain). We take completeness as an axiom of the real line, as every treatment must at some level; a rigorous

construction of $\mathbb{R}$ is a beautiful topic that lives one course away from this paper. Note the theorem needs *closed* intervals: the open intervals $(0, 1/n)$ are nested with empty intersection — this is why the free-gap lemma delivered a closed I' .

6.2. Clearing the small elements of B . Fix $B = \{b_1 < b_2 < \dots\}$, lacunary with ratio $r = 192 d_2$. The dial criterion of §4 handles targets above kc ; elements of B at or below kc are dodged automatically (Lemma 4.2). So we choose the offset to split B cleanly. Since B is infinite and increasing, some index j_0 has

$$b_{j_0} \geq 4d_2^2;$$

fix such a j_0 and set

$$c := b_{j_0} + 1.$$

Every element b_i with $i \leq j_0$ satisfies $b_i \leq b_{j_0} < kc$, hence is dodged automatically, whatever slope we pick. For the remaining elements define the shifted targets

$$t_i := b_i - kc \quad (i > j_0).$$

Three facts about these targets, each a short computation with the lacunarity.

(1) *The targets are positive — indeed the first one is already huge.* Using $b_{j_0+1} \geq 192 d_2 b_{j_0}$ and $kc = k(b_{j_0} + 1) \leq d_2(b_{j_0} + 1) \leq 2d_2 b_{j_0}$ (as $b_{j_0} \geq 1$):

$$t_{j_0+1} = b_{j_0+1} - kc \geq 192 d_2 b_{j_0} - 2 d_2 b_{j_0} = 190 d_2 b_{j_0} \geq 190 d_2 \cdot 4d_2^2 \geq 32 d_2^2,$$

comfortably meeting the free-gap lemma's threshold $t \geq 32d_2^2$. Later targets are larger still, since B increases.

(2) *Shifting preserves the ratio.* For $x \geq y > \tau \geq 0$ one has

$$\frac{x - \tau}{y - \tau} \geq \frac{x}{y} \quad (\text{cross-multiply: } xy - \tau x \geq xy - \tau y \iff \tau y \leq \tau x).$$

Apply this with $x = b_{i+1}$, $y = b_i$, $\tau = kc$ (valid: for $i > j_0$, $b_i \geq b_{j_0+1} \geq 192d_2 b_{j_0} > 2d_2 b_{j_0} \geq kc$, so $y > \tau$):

$$\frac{t_{i+1}}{t_i} \geq \frac{b_{i+1}}{b_i} \geq 192 d_2 \quad (i > j_0). \quad (4)$$

Subtracting the same amount from both terms of a ratio only helps the ratio — the targets inherit the lacunarity of B , undiluted.

(3) *If every target is dodged, we win.* Suppose $\gamma^* \in (d_2 - 1, d_2)$ is safe for every t_i , $i > j_0$. Then by the avoidance criterion (Proposition 4.4), no b_i with $i > j_0$ lies in $kA_{\gamma^*, c}$; and no b_i with $i \leq j_0$ does either, by the automatic clearance above. So $(kA_{\gamma^*, c}) \cap B = \emptyset$, and $A_{\gamma^*, c}$ is a legal (d_1, d_2) -sequence by Proposition 3.6.

6.3. The nest. Everything is now in place. Start from the closed interval

$$I_0 := [d_2 - \frac{1}{2}, d_2 - \frac{1}{4}],$$

of width exactly $\frac{1}{4}$. (We stop at $d_2 - \frac{1}{4}$ rather than d_2 out of tidiness only; any closed subinterval of $[d_2 - \frac{1}{2}, d_2]$ of width $\frac{1}{4}$ would do. Note $I_0 \subseteq (d_2 - 1, d_2)$ since $d_2 \geq 3$.)

Round 1. Apply the free-gap lemma to the first target t_{j_0+1} and the interval I_0 . Its hypotheses hold: $t_{j_0+1} \geq 32d_2^2$ by fact (1), and

$$|I_0| = \frac{1}{4} \geq \frac{4d_2^2}{t_{j_0+1}} \iff t_{j_0+1} \geq 16 d_2^2,$$

true again by fact (1). Out comes a closed $I_1 \subseteq I_0$, safe for t_{j_0+1} , with $|I_1| \geq d_2/(48 t_{j_0+1})$.




FIGURE 9. The nest, schematically. The widths are wildly not to scale, and cannot be: at $d_2 = 4$ each interval is 768 times narrower than its parent, so a true-scale second round would already be thinner than the ink. That violence is the message — the free-gap lemma spends a factor $192 d_2$ of width every round, and the Adversary's forced growth refills exactly that factor before the next target arrives. Completeness of the reals guarantees a point γ^* surviving all trims.

Round $j + 1$, given round j . Suppose inductively that the closed interval I_j is safe for the targets $t_{j_0+1}, \dots, t_{j_0+j}$ and satisfies

$$|I_j| \geq \frac{d_2}{48 t_{j_0+j}}.$$

We want to apply the free-gap lemma to the next target $t := t_{j_0+j+1}$ inside I_j . The only hypothesis in doubt is the width demand $|I_j| \geq 4d_2^2/t$, which in view of the inductive width bound follows from

$$\frac{d_2}{48 t_{j_0+j}} \geq \frac{4d_2^2}{t_{j_0+j+1}} \iff t_{j_0+j+1} \geq 192 d_2 \cdot t_{j_0+j},$$

which is precisely (4). *This is the entire ledger of the constant:* $192 = 4 \times 48$, the free-gap lemma's input demand ($4d_2^2/t$) divided by its output guarantee ($d_2/(48t)$). The lacunarity ratio of B is exactly what refills the width the previous round spent. The lemma returns a closed $I_{j+1} \subseteq I_j$, safe also for t_{j_0+j+1} , of width $\geq d_2/(48 t_{j_0+j+1})$ — the inductive hypothesis for the next round.

The bookkeeping, as a ledger — each row is one round:

Round	target	width needed	width delivered	who pays the gap
1	t_{j_0+1}	$4d_2^2/t_{j_0+1}$	$\geq d_2/(48 t_{j_0+1})$	$ I_0 = \frac{1}{4}$: ample
$j+1$	t_{j_0+j+1}	$4d_2^2/t_{j_0+j+1}$	$\geq d_2/(48 t_{j_0+j+1})$	lacunarity: $t_{j_0+j+1} \geq 192 d_2 t_{j_0+j}$

Each round *spends* a factor $192 d_2$ of width (input demand over output guarantee) and the Adversary's forced growth *refills* exactly that factor before the next target arrives. Figure 9 draws the nest.

The limit. The rounds produce nonempty closed bounded nested intervals $I_0 \supseteq I_1 \supseteq I_2 \supseteq \dots$; the nested interval theorem yields a real γ^* in all of them. For every $i > j_0$, the slope γ^* lies in I_{i-j_0} , which is safe for t_i ; so γ^* is safe for every target simultaneously. By fact (3) of §6.2, the Beatty walk $A_{\gamma^*,c}$ is a legal (d_1, d_2) -sequence with $(kA_{\gamma^*,c}) \cap B = \emptyset$. (This proves Theorem 6.1.) $\square$

6.4. One win, worked end to end. Before the summary, let us actually play a round of the game and win it, with numbers throughout. Take

$$k = 3, \quad (d_1, d_2) = (1, 4), \quad r = 192 d_2 = 768,$$

and let the Adversary answer with the powers of 800:

$$B = \{800, 640,000, 512,000,000, 409,600,000,000, \dots\}, \quad b_i = 800^i.$$

This is legal — $b_{i+1}/b_i = 800 \geq 768$ — and it is close to the fastest-shrinking play available at this ratio, so the bookkeeping will be tight. Now we answer.

The offset. We need $b_{j_0} \geq 4d_2^2 = 64$: already $b_1 = 800$ qualifies, so $j_0 = 1$ and

$$c = b_1 + 1 = 801, \quad kc = 2403.$$

The single element $b_1 = 800$ is now below $\min kA \geq k(c + d_2 - 1) = 3 \cdot 804 = 2412$, so it is dodged for free, whatever slope we choose. Every later element becomes a shifted target:

$$t_2 = 800^2 - 2403 = 637,597, \quad t_3 = 800^3 - 2403 = 511,997,597, \quad t_4 = 409,599,997,597.$$

(Check the threshold: $t_2 = 637,597 \geq 32d_2^2 = 512$, amply. And check the ratios survive the shift: $t_3/t_2 \approx 803.0$ and $t_4/t_3 \approx 800.004$, both ≥ 768 , as (4) promised.)

The rounds. Start at $I_0 = [3.5, 3.75]$, of width $\frac{1}{4}$, and turn the crank:

Round	target t	width needed $\frac{4d_2^2}{t}$	width we have	width delivered $ I_j $
1	637,597	1.00×10^{-4}	2.5×10^{-1}	1.47×10^{-6}
2	511,997,597	1.25×10^{-7}	1.47×10^{-6}	1.83×10^{-9}
3	409,599,997,597	1.56×10^{-10}	1.83×10^{-9}	2.29×10^{-12}

Read the middle two columns down: *what we have always exceeds what the next round needs.* Round 1's numbers, in full: $s_1 = \lceil 637597/3.75 \rceil = 170,026$, so

$$R = \frac{637597}{170026} \approx 3.749997, \quad L = \frac{637600}{170027} \approx 3.749993, \quad |J| = \frac{127519}{28909010702} \approx 4.41 \times 10^{-6},$$

and I_1 is the closed middle third of J — width 1.47×10^{-6} , comfortably above the guaranteed $d_2/(48t_2) \approx 1.31 \times 10^{-7}$. After three rounds the surviving stretch is

$$I_3 \approx [3.7499955633047, 3.7499955633071],$$

and the rounds continue forever, each one pinning down a few more digits. The common point is a specific real number $\gamma^* \approx 3.74999556330\dots$, and the walk

$$A = \{801 + \lfloor i\gamma^* \rfloor : i \geq 1\} = \{804, 808, 812, 815, 819, 823, 827, 830, \dots\}$$

— steps of 4 and 3, all inside $[1, 4]$ — has $(3A) \cap B = \emptyset$. That is a win, in full.

Where the constant lives. Compare each round's *delivery* against the next round's *demand*:

$$\frac{d_2/(48t_j)}{4d_2^2/t_{j+1}} = \frac{1}{192d_2} \cdot \frac{t_{j+1}}{t_j} \approx \frac{800}{768} = 1.042.$$

Four percent to spare, every round, forever — and that four percent is exactly the margin by which the Adversary's ratio 800 exceeds the demanded $192d_2 = 768$. Had the ratio been 760, each round's delivery would have fallen short of the next round's demand, and *this* induction would no longer close. (Note precisely what that does and does not say: it is a statement about the bookkeeping of this proof, not about the game. Whether 760 is in fact a winning announcement at these parameters is a question this document does not answer — the constant 192 is what our estimates can pay for, and §6's closing remark already disclaimed any belief that it is optimal.)

Big idea. The existence proof in one breath: aim a Beatty walk’s slope dial at the top of the permitted window; each element of B forbids a sparse family of dial positions, but the free-gap lemma always finds a safe sub-stretch losing only a factor $192d_2$ of width; lacunarity of B regains exactly that factor before the next element arrives; completeness of the reals hands us a single slope surviving every round.

Two closing remarks, and then a promised debt about irrationality. First, nothing was random and nothing was irrational-by-fiat: the winning slope γ^* is whatever real number the nest closes on, and no argument above ever asked anything else about it. Second, the constant $192d_2$ is an artifact of our bookkeeping (the middle third, the choice $\frac{15}{32}$, the crude bound (3)); we make no claim that it is anywhere near optimal.

Big idea. *Daylight is necessary but nowhere near sufficient — the true resource is the continuum.* An arithmetic progression with step $d_2 \geq k + 1$ has splendid daylight, yet it loses: its gap word is periodic, §10 will show that every eventually-periodic walk’s sumset contains a full congruence-class tail *whatever the parameters*, and §8 builds a single legal B — available at any prescribed ratio, including our own $192d_2$ — that meets every such sumset. More: a *rational* slope $\gamma = u/v$ makes the Beatty walk’s gap word periodic with period v (Lemma 3.2(2) gives $\lfloor (i+v)\gamma \rfloor = \lfloor i\gamma \rfloor + u$), so when the Adversary plays that certificate B , the slope γ^* our nest produces *cannot be rational*. (Two citations forward, both legitimate: Lemma 10.1 is proved without using $d_2 \leq k$ — see Remark 10.2 — and what we borrow from §8 is the *construction* of the certificate B , which is likewise parameter-free, not the conditional statement of Lemma 8.6(2).) The nested-interval theorem is therefore not a bookkeeping convenience: the completeness of the real line — an entire interval of dial positions, almost all of them irrational — is precisely the resource the existence proof spends, and the threshold $d_2 \geq k + 1$ is precisely the condition under which that interval exists.

State of play. We now know: above the threshold, $r = 192d_2$ wins — part (1) of Theorem A is proved. What remains: part (2), non-existence below the threshold, in its strong form. Next: convert “defeat every walk” into a property checked one walk at a time.

Exercise 6.2 (CHALLENGE). Trace the construction for $k = 3$, $d_2 = 4$, against the concrete lacunary sequence $b_i = 800^i$ (ratio $800 > 192 \cdot 4 \dots$ check!). Find j_0 and c , compute t_{j_0+1} , and carry out Round 1 numerically with the recipe of §5: exhibit an explicit closed interval of slopes, every one of which yields a $(1, 4)$ -walk whose 3-fold sums miss b_{j_0+1} .

Exercise 6.3 (CORE). The proof used lacunarity *only* through (4), one ratio per consecutive pair. Convince yourself the argument survives if the Adversary’s B merely satisfies $b_{i+1} \geq 192d_2 b_i$ *eventually* (for all $i \geq \text{some } i_0$) — and locate the single place in §6.2 that needs adjusting.

Part III. Non-existence: the reduction

7. INTERLUDE: A MINI-COURSE ON REMAINDERS

Why does the mathematics change character here? Part II lived on the real number line: when $d_2 \geq k + 1$ the windows leave daylight, and a continuum of dial positions was the resource we spent. Below the threshold that resource is gone — the windows crush into a solid block, and no real-valued tuning can help. But look *inside* the block: the integers a

walk's sums actually reach arrange themselves into rigid repeating patterns — arithmetic progressions, congruence classes. Tracking rigid repetition is a job for the arithmetic of clock faces, and that is the toolkit we now build.

From here to the end, the natural language of the proof is the arithmetic of *remainders* — what mathematicians call modular arithmetic. Part II never needed it; Parts III and IV use it constantly. This interlude builds all of it, with proofs, in a few pages. If you have seen modular arithmetic before, skim the statements — especially the covering lemma (Lemma 7.6), which the rest of the document uses more than a dozen times — and move on.

7.1. Divisibility, and division with remainder. For integers $d \neq 0$ and x , say d *divides* x , written $d \mid x$, if $x = dc$ for some integer c . So $7 \mid 42$, $7 \nmid 40$, every d divides 0, and 1 divides everything.

Lemma 7.1 (Division with remainder). *Let $m \geq 1$ and let x be any integer. There are unique integers q (the quotient) and r (the remainder) with*

$$x = qm + r, \quad 0 \leq r < m.$$

Proof. Existence: take $q := \lfloor x/m \rfloor$ and $r := x - qm$; then $q \leq x/m < q + 1$ (Definition 3.1) gives $0 \leq r < m$ after multiplying by m . Uniqueness: if $qm + r = q'm + r'$ with both r, r' in $[0, m)$, then $m(q - q') = r' - r$ is a multiple of m of absolute value $< m$, hence 0. $\square$

7.2. Greatest common divisors, and Bézout's lemma. The *greatest common divisor* $\gcd(a, b)$ of integers a, b (not both 0) is the largest integer dividing both. Call a and b *coprime* if $\gcd(a, b) = 1$: they share no factor beyond 1. For a finite set G , $\gcd(G)$ is the largest integer dividing every element.

The single deepest fact we need about gcds — and it is not very deep — is that the gcd can be *built* out of a and b :

Lemma 7.2 (Bézout). *Let a, b be integers, not both 0. Then there exist integers x, y with*

$$ax + by = \gcd(a, b).$$

Proof. Consider all integers of the form $ax + by$ (x, y ranging over all integers, positive and negative) and let d be the *smallest positive* one. Two things justify that phrase: at least one such integer is positive (take $x = a, y = b$, giving $a^2 + b^2 > 0$), and *every nonempty set of positive integers has a least element* — the well-ordering principle of §1.2, which is what makes “the smallest” meaningful. We claim $d = \gcd(a, b)$.

First, $d \mid a$: divide with remainder, $a = qd + r$ with $0 \leq r < d$; then

$$r = a - qd = a - q(ax + by) = a(1 - qx) + b(-qy)$$

is itself of the form $ax' + by'$, nonnegative and smaller than d — by minimality of d , r cannot be positive, so $r = 0$. Likewise $d \mid b$. So d is a common divisor. Finally, any common divisor c of a and b divides $ax + by = d$; in particular $\gcd(a, b) \leq d$ and $\gcd(a, b) \mid d$, while $d \leq \gcd(a, b)$ since d is a common divisor. Hence $d = \gcd(a, b)$. $\square$

Lemma 7.3 (Euclid's cancellation). *If $m \mid u\alpha$ and $\gcd(m, \alpha) = 1$, then $m \mid u$.*

Proof. Bézout gives $\alpha x + my = 1$. Multiply by u : $u = u\alpha x + umy$. Both terms on the right are divisible by m (the first because $m \mid u\alpha$), so $m \mid u$. $\square$

7.3. Congruence.

Definition 7.4. Fix $m \geq 1$. Integers x and y are *congruent modulo m* , written

$$x \equiv y \pmod{m}, \quad \text{if} \quad m \mid x - y.$$

Equivalently: x and y leave the same remainder upon division by m (if $x = qm + r$ and $y = q'm + r'$ with $r, r' \in [0, m)$, then $m \mid x - y$ iff $m \mid r - r'$ iff $r = r'$, since $|r - r'| < m$). Congruence behaves like a looser equality; two operations are always legal:

Lemma 7.5. *If $x \equiv y$ and $u \equiv v \pmod{m}$, then $x + u \equiv y + v$ and $xu \equiv yv \pmod{m}$.*

Proof. $(x + u) - (y + v) = (x - y) + (u - v)$: a sum of multiples of m . And $xu - yv = (x - y)u + y(u - v)$: likewise. $\square$

Division, by contrast, is *not* generally legal: $6 \equiv 16 \pmod{10}$, i.e. $2 \cdot 3 \equiv 2 \cdot 8$, yet $3 \not\equiv 8 \pmod{10}$. The correct substitute for division is coming in Lemma 7.7.

Fix m . The *congruence class* (or *residue class*) of r modulo m is the set of all integers congruent to r :

$$\{\dots, r - 2m, r - m, r, r + m, r + 2m, \dots\}$$

— an arithmetic progression with gap m , extending both ways. There are exactly m classes (one per remainder $0, 1, \dots, m - 1$), and every integer belongs to exactly one (Lemma 7.1). Two reflexes to internalize: *a class is a progression*, so it contains arbitrarily large numbers; and *knowing x 's class* means knowing x up to adding multiples of m .

7.4. The covering lemma, and inverses. Here is the workhorse. It will be cited constantly in Part IV, and silently used whenever we say “the multiples of α sweep out all residues”.

Lemma 7.6 (Covering lemma). *Let $m \geq 1$ and let α be coprime to m . Then the m numbers*

$$0 \cdot \alpha, \quad 1 \cdot \alpha, \quad 2 \cdot \alpha, \quad \dots, \quad (m - 1) \cdot \alpha$$

lie in m different congruence classes modulo m — hence in all m classes, each exactly once.

Proof. Suppose $i\alpha \equiv j\alpha \pmod{m}$ with $0 \leq i, j \leq m - 1$. Then $m \mid (i - j)\alpha$, so $m \mid i - j$ by Euclid's cancellation (Lemma 7.3); since $|i - j| < m$, this forces $i = j$. So the m listed numbers occupy m distinct classes; as there are only m classes in total, every class is hit exactly once. $\square$

Lemma 7.7 (Modular inverse). *If $\gcd(\alpha, m) = 1$, there is an integer $\bar{\alpha} \in [0, m - 1]$ with $\alpha \bar{\alpha} \equiv 1 \pmod{m}$; it is unique in that range, and we call it the inverse of α modulo m . Multiplying by $\bar{\alpha}$ undoes multiplying by α — the legal substitute for division.*

Proof. By the covering lemma, some $i \cdot \alpha$ with $i \in [0, m - 1]$ lies in the class of 1: take $\bar{\alpha} := i$. Uniqueness: if $\alpha i \equiv \alpha i' \equiv 1 \pmod{m}$, then $i \equiv i \cdot (\alpha i') = (i\alpha) \cdot i' \equiv i' \pmod{m}$ (using Lemma 7.5 twice), and two congruent numbers in $[0, m - 1]$ are equal. $\square$

Example 7.8 (Solving a congruence). Solve $5s \equiv 2 \pmod{7}$. The inverse of 5 modulo 7 is 3 (as $5 \cdot 3 = 15 \equiv 1$), so multiply both sides by 3: $s \equiv 3 \cdot 2 = 6 \pmod{7}$. The solutions are exactly the class $\{\dots, 6, 13, 20, 27, \dots\}$ — one class, an arithmetic progression. This “pin the variable to one class” move is used at a crucial point in §12 and §14, and throughout Part IV.

Exercise 7.9 (ROUTINE). (a) Compute $\gcd(2026, 715)$ by repeatedly replacing the larger number by its remainder upon division by the smaller (this is *Euclid's algorithm* — explain why each step preserves the gcd). (b) Find Bézout integers x, y for $a = 21, b = 8$. (c) Find the inverse of 8 modulo 21, and solve $8s \equiv 5 \pmod{21}$.




FIGURE 10. The covering lemma on two clock faces. Left: $\gcd(3, 7) = 1$, and the multiples $0, 3, 6, 9, \dots$ visit all seven pegs before returning — the star polygon is the lemma. Right: $\gcd(2, 6) = 2$, and the multiples of 2 never leave the even pegs. Coprimality is exactly the difference.

Exercise 7.10 (ROUTINE). (a) List the numbers $0, 5, 10, 15, 20, 25, 30$ modulo 7 and watch the covering lemma happen. (b) Show the covering lemma *fails* without coprimality: list $0, 6, 12, 18$ modulo 8. (c) Prove: if $\gcd(\alpha, m) = 1$, then for any c the congruence $\alpha s \equiv c \pmod{m}$ has exactly one solution class.

8. THE CERTIFICATE TRICK: ONE B AGAINST EVERY A

Part III begins the non-existence half. **Standing assumptions for Parts III and IV:** k, d_1, d_2 are integers with

$$k \geq 3, \quad 1 \leq d_1 < d_2, \quad d_2 \leq k.$$

The goal is part (2) of Theorem A in its full strength: for *any* prescribed sequence $R(1), R(2), \dots$ of ratio lower bounds, a single B with $b_{i+1} \geq R(i) b_i$ whose intersection with kA is nonempty for *every* (d_1, d_2) -sequence A .

This section performs the first great simplification: it converts “defeat every A ” — a statement about all walks at once — into a property to be checked *one walk at a time*.

Remark 8.1 (Where $k \geq 3$ is spent, and what happens at $k = 2$). Part II proudly worked from $k = 2$; Parts III–IV will not. The hypothesis $k \geq 3$ is consumed in two specific places: the width dichotomy’s even case needs $d_2 \leq k - 1$ or a boundary argument requiring $k \geq 4$ (§§12–14), and the slot lemma needs a tail alphabet of three letters, which $d_2 = 2$ forbids outright. At $k = 2, d_2 = 2$ this paper’s machinery genuinely does not decide the question, the problem as posed excludes it, and we make no claim there.

8.1. Tail-covering.

Definition 8.2 (Tail-covering). A (d_1, d_2) -sequence A is *tail-covering* (for our fixed k) if there are integers $m \geq 1, \rho$, and X_0 such that

$$kA \supseteq \{x \geq X_0 : x \equiv \rho \pmod{m}\} :$$

the sumset contains *every* sufficiently large number in some fixed congruence class.

We met the phenomenon in Example 2.4: for an arithmetic progression with gap d , the sumset kA is a full tail of a class modulo d . The definition allows any modulus m , and this generosity is necessary — the modulus that works can exceed d_2 , and kA need not contain all large integers.

Example 8.3. Let the gaps of A alternate $3, 5, 3, 5, \dots$ and take $k = 5$ (so $d_2 = 5 \leq k$). The prefix sums (partial sums of the gaps) are $0, 3, 8, 11, 16, 19, 24, \dots$: the numbers $8m$ and $8m + 3$. A 5-fold sum of members of A is $5a_1$ plus a sum of five prefix sums, and modulo 8 such a sum is $3j$ where $j \in \{0, 1, \dots, 5\}$ counts how many of the five had the form $8m+3$. Since $3j \bmod 8$ takes the values $\{0, 3, 6, 1, 4, 7\}$, the sumset $5A$ lives in six of the eight classes modulo 8 and misses the classes $5a_1 + 2$ and $5a_1 + 5$ entirely. It is tail-covering — it eventually contains all of $5a_1 + 8\mathbb{Z}$, for instance — but with modulus $8 > d_2$, and it is far from containing all large integers.

Exercise 8.4 (CORE). Verify the claims of Example 8.3: (a) the prefix sums are exactly the numbers $8m$ and $8m + 3$ ($m \geq 0$); (b) $5A$ contains every sufficiently large multiple of 8, shifted by $5a_1$ (how large is “sufficiently”); (c) $5A$ misses the class $5a_1 + 2 \pmod{8}$ completely.

8.2. One sequence to meet them all. Suppose — this is what the rest of Part III will prove — that *every* admissible walk is tail-covering. Each walk hands us a triple (m, ρ, X_0) : its sumset owns the class ρ modulo m from X_0 onward. The Adversary does not know which walk you will play, hence does not know which class; but the candidate pairs (m, ρ) can be arranged in a single infinite list, and the Adversary has infinitely many moves. So the Adversary hedges: build B to visit *every* class of *every* modulus, infinitely often, while growing as fast as anyone demands. Whatever walk you play, its class eventually contains an element of B beyond its X_0 — and that element is a collision.

To organize the hedging we need a list of all pairs.

Lemma 8.5 (The recurring list). *There is a sequence of pairs $(m_1, \rho_1), (m_2, \rho_2), \dots$ in which every pair (m, ρ) with $m \geq 1$ and $0 \leq \rho < m$ appears infinitely often.*

Proof. List in blocks. Block N (for $N = 1, 2, 3, \dots$) lists all pairs with $m \leq N$, say in order of m then ρ : block 2 is $(1, 0), (2, 0), (2, 1)$; block 3 is $(1, 0), (2, 0), (2, 1), (3, 0), (3, 1), (3, 2)$; and so on. Concatenate the blocks. Each block is finite, so the concatenation is a legitimate infinite list, and a given pair (m, ρ) appears in every block from the m -th onward — infinitely often. $\square$

Lemma 8.6 (The certificate). *Fix the recurring list of Lemma 8.5 and let $R : \mathbb{N} \rightarrow \mathbb{N}$ be an arbitrary prescribed ratio sequence — growing as fast as you like, with no regularity assumed. Define B by $b_1 := 1$ and, for $i \geq 1$,*

$$b_{i+1} := \text{the least integer} > \max(R(i)b_i, b_i, i) \text{ that is } \equiv \rho_{i+1} \pmod{m_{i+1}}.$$

Then:

- (1) B is a strictly increasing sequence of positive integers with $b_{i+1} \geq R(i)b_i$ for every i ;
- (2) if every (d_1, d_2) -sequence is tail-covering, then $(kA) \cap B \neq \emptyset$ for every (d_1, d_2) -sequence A .

Proof. (1) Each b_{i+1} exists: the numbers $\equiv \rho_{i+1} \pmod{m_{i+1}}$ form an arithmetic progression, which contains numbers beyond any bound. By construction $b_{i+1} > b_i$ (strictly increasing, hence positive throughout) and $b_{i+1} > R(i)b_i$. Note also $b_{i+1} > i$, so $b_i \geq i$ for every $i \geq 2$: the sequence grows at least linearly, a small fact we need in a moment.

(2) Let A be any (d_1, d_2) -sequence, and let (m, ρ, X_0) witness its tail-covering. We must find an element of B that is $\equiv \rho \pmod{m}$ and $\geq X_0$ — such an element lies in kA by Definition 8.2. The pair (m, ρ) appears infinitely often in the recurring list, say at positions $i^{(1)} < i^{(2)} < \dots$; at each such position $i \geq 2$ the construction made $b_i \equiv \rho \pmod{m}$. Since $b_i \geq i \rightarrow \infty$, some such position has $b_i \geq X_0$ as well. That b_i belongs to both B and kA . $\square$




FIGURE 11. The certificate's weave. Rows are the pairs (m, ρ) of the recurring list; the Adversary's terms $b_1, b_2, \dots$ (dots, marching rightward ever faster) visit every row again and again, forever. Whatever walk you play, its sumset owns some row from some point X_0 on (shaded) — and that row still receives visits beyond any X_0 . Each such visit is a collision.

Big idea. Sparseness cannot save the walker, because congruence classes are everywhere: however fast B is forced to grow, it can always take its next step *inside* any prescribed class. The certificate B cycles through all classes of all moduli forever, so it eventually lands, arbitrarily late, in whichever class a given walk's sumset owns. One sequence, built blind, defeats every tail-covering walk.

Notice what the certificate does *not* depend on: k , d_1 and d_2 never entered the construction — only the ratio sequence R and the recurring list (Figure 11). So the very same B works simultaneously for all parameter triples below the threshold. This is stronger than the theorem demands; it comes for free.

The non-existence half is now reduced to a statement about individual walks:

Goal of Part III. *Every (d_1, d_2) -sequence with $d_2 \leq k$ (and $k \geq 3$) is tail-covering.*

The rest of Part III proves this, walk by walk, according to how the walk wobbles; one case will lean on the subset-sum lemma of Part IV.

State of play. We now know: one B , built blind, defeats every tail-covering walk — gap (2) of the ledger is closed. What remains: proving that below the threshold *every* admissible walk is tail-covering (gap (3)). Next: coordinates adapted to a single walk, and the case split that organizes the rest of the part.

Exercise 8.7 (ROUTINE). Take $R(i) = 10^i$. Compute b_2, b_3, b_4 of the certificate (with the block list of Lemma 8.5: $(m_2, \rho_2) = (1, 0)$, $(m_3, \rho_3) = (2, 0)$, $(m_4, \rho_4) = (2, 1)$). Check the ratios.

9. GAP WORDS, TAIL ALPHABETS, AND THE PLAN

To analyze a single walk we need coordinates adapted to its steps. This section sets them up, dispatches the two easiest kinds of walk, and lays out the battle plan for the rest.

9.1. Prefix sums, and the sumset in step coordinates. Let A be a (d_1, d_2) -sequence with gaps $g_n = a_{n+1} - a_n \in [d_1, d_2]$. Define the *prefix sums*

$$p_0 := 0, \quad p_n := g_1 + g_2 + \dots + g_n \quad (n \geq 1),$$




FIGURE 12. Four coordinate systems for one walk (gaps 3, 2, 3, 2, 2, 3, 2, ...; letters $\delta = 2$, $\delta' = 3$, so $e = 1$): the terms a_n on the line, the gaps g_n as arches, the binary word h (large step = 1), and the prefix sums $p_n = a_{n+1} - a_1$. §§9–14 translate between these constantly; this figure is the dictionary.

so that $a_{n+1} = a_1 + p_n$ for every $n \geq 0$: each term of the walk is the start plus the steps taken so far.

Lemma 9.1 (The sumset in step coordinates).

$$kA = ka_1 + \{p_{i_1} + p_{i_2} + \cdots + p_{i_k} : i_1, \dots, i_k \geq 0\}.$$

Proof. A member of kA is $a_{j_1} + \cdots + a_{j_k}$ with each $j_t \geq 1$; substituting $a_{j_t} = a_1 + p_{j_t-1}$ and setting $i_t := j_t - 1 \geq 0$ gives one direction, and the substitution reverses. $\square$

From now on we study the set $kA - ka_1$ of *sums of k prefix sums*. Note that tail-covering is insensitive to the shift: kA contains a class tail if and only if $kA - ka_1$ does (with the residue shifted by ka_1).

9.2. The gap word and the tail alphabet. The *gap word* of A is the infinite sequence $g_1, g_2, g_3, \dots$ read as a word — a string of “letters” drawn from the finite alphabet $\{d_1, \dots, d_2\}$ (Figure 12 shows the coordinate systems this part will move between). Since the alphabet is finite, some letters occur infinitely often in the word and some (perhaps) only finitely often. Define the *tail alphabet*

$$G_\infty := \{\delta : g_n = \delta \text{ for infinitely many } n\},$$

which is nonempty, and fix an index T beyond which *only* letters of G_∞ occur (possible: each of the finitely many non-recurring letters has a last occurrence; take T past all of them).

Lemma 9.2 (Tail reduction). *Let $A' := \{a_{T+1} < a_{T+2} < \cdots\}$ be the tail of A beyond T — itself a (d_1, d_2) -sequence, whose gap word uses only letters of G_∞ . If A' is tail-covering, so is A .*

Proof. Every sum of k members of A' is a sum of k members of A , so $kA \supseteq kA'$; a class tail inside kA' is a class tail inside kA . $\square$

We will invoke this silently: “pass to the tail” means replace A by A' , so that every gap lies in G_∞ and every letter of G_∞ still occurs infinitely often.

The analysis now splits on the *size* of the tail alphabet, and the first case is immediate:

Lemma 9.3 (One letter). *If $G_\infty = \{\delta\}$, then A is tail-covering.*

Proof. Pass to the tail: all gaps equal δ , so the tail is an arithmetic progression with gap δ , and Example 2.4 (which is now official: the computation there is exact) shows its k -fold sumset is a full class tail modulo δ . $\square$

One more normalization. If all the recurring letters share a common factor, we can divide it out:

Lemma 9.4 (Rescaling). *Suppose (after passing to the tail) every gap of A lies in G_∞ , and let $g := \gcd(G_\infty) > 1$. Define the rescaled walk A' by*

$$a'_n := \frac{a_n - a_1}{g} + 1 \quad (n \geq 1),$$

a sequence of positive integers with gaps $g'_n = g_n/g$, alphabet $G_\infty/g := \{\delta/g : \delta \in G_\infty\}$, $\gcd(G_\infty/g) = 1$, and maximum letter $\leq d_2/g$. If A' is tail-covering modulo m' , then A is tail-covering modulo gm' .

Proof. Every gap of A is a multiple of g , so each $a_n - a_1$ is one too and the a'_n are integers, strictly increasing from $a'_1 = 1$; their gaps are g_n/g by construction. Unwinding the definition, $a_n = (a_1 - g) + g a'_n$, so a sum of k members of A is $k(a_1 - g) + g \cdot$ (a sum of k members of A'):

$$kA = k(a_1 - g) + g \cdot kA'.$$

Now suppose $kA' \supseteq \{x \geq X'_0 : x \equiv \rho' \pmod{m'}\}$. Consider any $y \equiv k(a_1 - g) + g\rho' \pmod{gm'}$ with y large, and put $w := y - k(a_1 - g)$. Three small deductions. First, the congruence says $w = g\rho' + gm'c$ for some integer c — both terms visibly multiples of g — so $g \mid w$. Second, $x := w/g = \rho' + m'c \equiv \rho' \pmod{m'}$. Third, x is large along with y . So $x \in kA'$, and $y = k(a_1 - g) + gx \in kA$. Hence kA contains a full class tail modulo gm' . $\square$

So from now on we may and do assume $\gcd(G_\infty) = 1$ (rescaling first if necessary, and remembering that the final modulus picks up a factor g — harmless, since tail-covering permits any modulus).

9.3. The battle plan. Here is the case tree for the Goal; each node names the section that closes it.

Tail alphabet (after rescaling)	Closed by
$ G_\infty = 1$	Lemma 9.3 (arithmetic progression)
$ G_\infty = 2$, gap word eventually periodic	§10
$ G_\infty = 2$, k even and $d_2 = k$ (the boundary)	§14 (its own routing)
$ G_\infty = 2$, otherwise:	
— irregular enough (wide band)	the sweep (§12)
— maximally regular (balanced)	Morse–Hedlund (§13), then the Sturmian endgame (§14)
$ G_\infty \geq 3$	§15 (slot lemma) + Part IV

This is the *canonical* tree, drawn as a flowchart in Figure 13: the assembly in §25 will retrace exactly these branches, in exactly this order, with the closing lemma numbers filled in. (The third row may look unmotivated now — it is the one configuration the width machinery of §12 cannot reach, and it gets a dedicated argument.)

The two-letter case is the deep one: a two-letter walk has exactly one degree of freedom per step, and the walks that use that freedom *most evenly* — neither settling into a period nor drifting irregularly — are precisely the Sturmian walks, where the hardest analysis lives. Three or more letters, counterintuitively, is *easier*: more letters means more kinds of step to combine, and a bounded combinatorial gadget (the slot lemma) plus one finite theorem (Part IV) covers all such walks at once.




FIGURE 13. Routing a walk, after passing to the tail and rescaling to $\gcd(G_\infty) = 1$. Every admissible walk enters at the top and leaves through exactly one branch, each of which ends in *tail-covering*. The two-letter branches (the lower four) are where the work is; the assembly table of §25 is this picture in words.

10. EVENTUALLY PERIODIC GAP WORDS

A gap word is *eventually periodic* if beyond some point it repeats with a fixed period: there are $N_0 \geq 0$ and $\ell \geq 1$ with

$$g_{n+\ell} = g_n \quad \text{for all } n > N_0.$$

Lemma 10.1 (Eventually periodic). *If the gap word of A is eventually periodic, then A is tail-covering.*

Remark 10.2 (This one escapes the standing assumption). Part III runs under the standing hypothesis $d_2 \leq k$, but the proof below never uses it — no relation between d_2 and k is invoked anywhere in it. Lemma 10.1 is therefore the one result of this part that holds *above* the threshold as well, which is exactly what the closing remark of §6 needed when it argued that a rational slope loses whatever the parameters. Worth flagging, because everywhere else in Parts III and IV the standing assumption is genuinely in force.

Proof. Pass to the tail beyond N_0 (Lemma 9.2); relabel so that the gap word satisfies $g_{n+\ell} = g_n$ for all $n \geq 1$. Let

$$P := g_1 + g_2 + \cdots + g_\ell$$

be the *period sum* — the distance the walk advances in each full period. Then the prefix sums repeat with drift P : for $0 \leq j < \ell$ and $m \geq 0$,

$$p_{m\ell+j} = mP + p_j,$$

by grouping the first $m\ell$ gaps into m full periods (each summing to P) plus the leftover prefix.

Now build sums of k prefix sums using only the multiples of the period: choosing $i_t = m_t\ell$ (i.e. $j = 0$) for all t gives

$$p_{m_1\ell} + \cdots + p_{m_k\ell} = (m_1 + \cdots + m_k)P,$$

and $m_1 + \cdots + m_k$ ranges over every integer ≥ 0 (Exercise 2.5 showed this ranging argument for $\geq k$; here the i_t may be 0, so even every integer ≥ 0). By Lemma 9.1, $kA \supseteq ka_1 + P\{0, 1, 2, \dots\}$: a full tail of the class of ka_1 modulo P . $\square$

Note how little of the period structure we used — only the sub-walk at the period marks, which is an arithmetic progression with gap P . The other residues $j \neq 0$ would give further




FIGURE 14. Balance, by sliding windows. Above, a balanced word: any two windows of the same length hold 1-counts differing by at most one, wherever you place them. Below, 00110011... fails at once — the windows 00 and 11 differ by two. Balance is the strongest possible “spread the 1s evenly” condition short of periodicity, and §13 shows it forces the word to track a straight line.

classes (as in Example 8.3, which is the case $\ell = 2$, $P = 8$ of this proof), but one class is all tail-covering asks.

Exercise 10.3 (CORE). The walk with repeating gap block $(2, 3, 2)$ (period $\ell = 3$, $P = 7$) and $k = 3$: which classes modulo 7 does $3A - 3a_1$ contain in full, beyond the class of 0? Compute the possible values of $p_{i_1} + p_{i_2} + p_{i_3} \bmod 7$.

11. INTERLUDE: A MINI-COURSE ON WORDS, AND A LITTLE ANALYSIS

The two-letter case will occupy us for four sections, and it runs on a small amount of theory — about *words* (infinite strings of symbols) and about *limits* — that deserves its own unhurried introduction. Nothing here mentions the problem; it is pure toolbox.

11.1. Words. An *alphabet* is a finite set of symbols; a *word* is a finite or infinite string of them. We care about *binary* words, alphabet $\{0, 1\}$, written $h = h_1h_2h_3\cdots$ with $h_n \in \{0, 1\}$. A *window* of h is a run of consecutive letters $h_{i+1}h_{i+2}\cdots h_j$ (which we address by its endpoints $i < j$; its *length* is $j - i$).

The *counting function* of h is

$$q_0 := 0, \quad q_n := h_1 + h_2 + \cdots + h_n \quad (\text{the number of 1s among the first } n \text{ letters}).$$

The count of 1s in the window from i to j is then $q_j - q_i$. Two properties tie h and q together: q is nondecreasing with steps $q_n - q_{n-1} = h_n \in \{0, 1\}$, and conversely any such staircase function comes from exactly one word.

Definition 11.1 (Balanced). A binary word h is *balanced* if any two windows of equal length contain numbers of 1s differing by at most one:

$$j - i = j' - i' \implies |(q_j - q_i) - (q_{j'} - q_{i'})| \leq 1.$$

Balance is a strong uniformity requirement: wherever you cut out two equal stretches of the word, they must contain almost exactly the same number of 1s (Figure 14).

Example 11.2. (a) The constant words $000\cdots$ and $111\cdots$ are balanced, as is the alternating word $010101\cdots$ (a window of length $2m$ always has exactly m ones; length $2m+1$ has m or $m+1$). (b) The periodic word $001001001\cdots$ is balanced (check windows of each length modulo 3). (c) The periodic word $00110011\cdots$ is *not* balanced: the windows 00 and 11 have 1-counts 0 and 2. (d) The *Fibonacci word* $0100101001001010010\cdots$ — built by iterating the

substitution $0 \mapsto 01$, $1 \mapsto 0$ — is balanced and *not* eventually periodic. It will reappear in §14 as the paradigm of the hardest case. (Proving these two claims about it is not needed for our proof; the classification below works from balancedness alone, whatever word exhibits it.)

Exercise 11.3 (ROUTINE). Verify (b) and (c) of Example 11.2. Then show that in a balanced word, if some window of length Δ contains u ones, every window of length Δ contains $u - 1$, u , or $u + 1$ ones.

11.2. Suprema, and limits. Two notions from analysis, in the exact strength we need.

Supremum. A set S of real numbers that is nonempty and bounded above has a *least upper bound* $\sup S$: an upper bound of S that is $\leq$ every other upper bound. (Its existence is, again, the completeness of the real line — the same axiom behind the nested interval theorem of §6; indeed each implies the other.) Likewise a nonempty set bounded below has a greatest lower bound $\inf S$. Two reflexes to internalize: every element of S is $\leq \sup S$; and for any $\varepsilon > 0$, some element of S exceeds $\sup S - \varepsilon$ (else $\sup S - \varepsilon$ would be a smaller upper bound).

One warning that will matter: *a supremum need not belong to its set.* The set $\{\frac{1}{2}, \frac{2}{3}, \frac{3}{4}, \frac{4}{5}, \dots\}$ — the numbers $1 - \frac{1}{n}$ — has supremum 1, yet no element equals 1. When §13 defines a certain slope as a supremum, this distinction is load-bearing: read “sup” as *least upper bound*, never as *maximum*.

To see the exact supremum that section will use, take the alternating word $0101\dots$, whose counting function we tabulated above, and list the quantities $(q_n - 1)/n$:

$$\frac{-1}{1}, \quad \frac{0}{2}, \quad \frac{0}{3}, \quad \frac{1}{4}, \quad \frac{1}{5}, \quad \frac{2}{6}, \quad \frac{2}{7}, \quad \frac{3}{8}, \dots = -1, 0, 0, 0.25, 0.2, 0.333, 0.286, 0.375, \dots$$

These creep upward toward $\frac{1}{2}$ without ever reaching it (the even terms are $\frac{n/2-1}{n} = \frac{1}{2} - \frac{1}{n}$). So the supremum is $\frac{1}{2}$ — the density of 1s in the word, exactly as it should be — and it is attained by no term. That is why §13 must build the slope as a supremum and only afterwards prove it is the limit.

Limits. A sequence $x_1, x_2, \dots$ of reals *converges* to x (written $x_n \rightarrow x$) if for every $\varepsilon > 0$, all but finitely many n have $|x_n - x| < \varepsilon$. We need only three facts, each immediate from the definition: (i) $1/n \rightarrow 0$; (ii) if $x_n \rightarrow x$ and $x_n \leq C$ for all n , then $x \leq C$ (a limit cannot escape a closed bound); (iii) if $x_n \rightarrow x$ and $x_n \geq c$ for all n , then $x \geq c$.

11.3. The discrete intermediate value theorem.

Lemma 11.4 (Discrete IVT). *Let $v_0, v_1, \dots, v_N$ be integers with $|v_{i+1} - v_i| \leq 1$ for every i . Then every integer between v_0 and v_N (inclusive) equals some v_i .*

Proof. Say $v_0 \leq v_N$ (else reverse the list) and let u be an integer with $v_0 \leq u \leq v_N$. Let j be the *first* index with $v_j \geq u$ (one exists: $j = N$ qualifies). If $j = 0$, then $v_0 \geq u \geq v_0$, so $v_0 = u$. Otherwise $v_{j-1} < u$, so $v_{j-1} \leq u - 1$; and $v_j \leq v_{j-1} + 1 \leq u$, which with $v_j \geq u$ gives $v_j = u$. $\square$

A sequence of integers that never jumps by more than 1 cannot step over an integer. We will use this constantly: it is the discrete shadow of the “no-skipping” arguments that ran through Part II.

12. TWO LETTERS: THE REACHABLE BAND, AND THE SWEEP

We now take up the main case. **Standing setup for §§12–14:** after tail-passing and rescaling (§9), the walk A has tail alphabet $G_\infty = \{\delta < \delta'\}$ with

$$\gcd(\delta, \delta') = 1, \quad e := \delta' - \delta \geq 1, \quad \delta' = \delta + e \leq d_2 \leq k,$$

and both letters occur infinitely often. (The bound $\delta' \leq d_2$ holds because every gap is $\leq d_2$; after rescaling it could only improve. Note $\gcd(\delta, e) = \gcd(\delta, \delta') = 1$.) Encode the gap word in binary: write

$$g_n = \delta + e h_n, \quad h_n \in \{0, 1\},$$

so small steps are 0s and large steps are 1s, and let q_n be the counting function of the binary word h .

12.1. The reachable band. The prefix sums split along the same lines: $p_n = \sum_{j \leq n} (\delta + e h_j) = \delta n + e q_n$. So a sum of k prefix sums, with indices $i_1, \dots, i_k \geq 0$ summing to s , is

$$p_{i_1} + \dots + p_{i_k} = \delta s + e (q_{i_1} + \dots + q_{i_k}),$$

and by Lemma 9.1,

$$kA - k a_1 = \{ \delta s + e w : s \geq 0, w \in W_s \}, \quad W_s := \left\{ \sum_{t=1}^k q_{i_t} : i_t \geq 0, \sum_{t=1}^k i_t = s \right\}. \quad (5)$$

Everything now hangs on the sets W_s : for each “index budget” s , which totals w of 1-counts are achievable? Before the general lemma, compute a few by hand. Take the alternating word $h = 0101 \dots$ — so $h_1 = 0, h_2 = 1, h_3 = 0, \dots$, and the counting function runs

$$q_0 = 0, \quad q_1 = 0, \quad q_2 = 1, \quad q_3 = 1, \quad q_4 = 2, \quad q_5 = 2, \quad q_6 = 3, \dots$$

— and let $k = 3$. At budget $s = 4$ the configurations (unordered, since order never matters) and their values $q_{i_1} + q_{i_2} + q_{i_3}$ are

$$(4, 0, 0) \mapsto 2, \quad (3, 1, 0) \mapsto 1, \quad (2, 2, 0) \mapsto 2, \quad (2, 1, 1) \mapsto 1, \quad (1, 1, 2) \mapsto 1,$$

so $W_4 = \{1, 2\}$. Running the same computation at each budget gives the band:

s	0	1	2	3	4	5	6	7	8
W_s	$\{0\}$	$\{0\}$	$\{0, 1\}$	$\{0, 1\}$	$\{1, 2\}$	$\{1, 2\}$	$\{2, 3\}$	$\{2, 3\}$	$\{3, 4\}$
$w^-(s)$	0	0	0	0	1	1	2	2	3
$w^+(s)$	0	0	1	1	2	2	3	3	4

Two features leap out, and they are exactly the two halves of the next lemma: *no row has a hole*, and *each endpoint row climbs by 0 or 1 at a time*. Neither is an accident of this example.

Lemma 12.1 (The band is an interval). *For every $s \geq 0$, the set W_s is a full interval of integers: $W_s = [w^-(s), w^+(s)] \cap \mathbb{Z}$ for some integers $w^-(s) \leq w^+(s)$. Moreover both endpoints climb gently:*

$$w^\pm(s+1) - w^\pm(s) \in \{0, 1\}.$$

Proof. Call a list $(i_1, \dots, i_k)$ of nonnegative integers with sum s a *configuration* (order will never matter — what counts is which indices appear, with repetition allowed), and its *value* the total $\sum_t q_{i_t}$. A *move* takes one unit from one slot and gives it to another: it replaces some pair (i_a, i_b) by $(i_a - 1, i_b + 1)$ (allowed when $i_a \geq 1$). A move keeps the sum s and changes the value by

$$(q_{i_b+1} - q_{i_b}) - (q_{i_a} - q_{i_a-1}) = h_{i_b+1} - h_{i_a} \in \{-1, 0, 1\}.$$

Any configuration can be walked to any other by such moves. Let $(i_1, \dots, i_k)$ and $(i'_1, \dots, i'_k)$ be two configurations, and measure their disagreement slot by slot: $D := \sum_t |i_t - i'_t|$. Suppose $D > 0$. Since both configurations have the same total s , the slots where i_t exceeds i'_t and the slots where it falls short must both exist: an excess somewhere forces a deficit elsewhere. Pick a slot a with $i_a > i'_a$ and a slot b with $i_b < i'_b$. The move $(i_a, i_b) \mapsto (i_a - 1, i_b + 1)$ is



FIGURE 15. A configuration move. The k index-slots hold a total budget s ; shifting one unit from a high slot to a low one keeps the total fixed and changes the value by a single letter's difference, hence by at most one. Since any configuration reaches any other through such moves, the achievable values step by at most one along the way — so they cannot skip an integer, and the reachable set W_s is a solid interval.

legal ($i_a > i'_a \geq 0$, so $i_a \geq 1$), and it shrinks both $|i_a - i'_a|$ and $|i_b - i'_b|$ by one: D drops by 2. After finitely many such moves $D = 0$ and the configurations agree.

In particular there is a chain of moves from a minimum-value configuration to a maximum-value one; the values along the chain form an integer sequence with steps in $\{-1, 0, 1\}$ from $w^-(s)$ to $w^+(s)$, so by the discrete IVT (Lemma 11.4) every integer in $[w^-(s), w^+(s)]$ is a value. That is the interval claim, and Figure 15 is the mechanism.

For the endpoint increments, compare budgets s and $s + 1$ in both directions. Two elementary moves are available. *Adding* a unit to some slot of a configuration for s produces a configuration for $s + 1$ whose value has grown by $h_{i_a+1} \in \{0, 1\}$ — it never falls and never gains more than one. *Removing* a unit from a positive slot of a configuration for $s + 1$ produces one for s whose value has dropped by $h_{i_a} \in \{0, 1\}$. Apply each move to each kind of extremal configuration:

Upper endpoint. Adding a unit to a maximizer for s exhibits a configuration for $s + 1$ of value $\geq w^+(s)$, so $w^+(s + 1) \geq w^+(s)$: the top never falls. Removing a unit from a maximizer for $s + 1$ exhibits a configuration for s of value $\geq w^+(s + 1) - 1$, so $w^+(s) \geq w^+(s + 1) - 1$: the top never rises by more than one. Together, $w^+(s + 1) - w^+(s) \in \{0, 1\}$.

Lower endpoint. Symmetrically: removing a unit from a minimizer for $s + 1$ gives $w^-(s) \leq w^-(s + 1)$ (the floor never falls), and adding a unit to a minimizer for s gives $w^-(s + 1) \leq w^-(s) + 1$ (it never rises by more than one). Together, $w^-(s + 1) - w^-(s) \in \{0, 1\}$. $\square$

So for each s the reachable pairs (s, w) form a vertical integer segment, and as s grows these segments drift upward, each endpoint by 0 or 1 per step: a *band* across the (s, w) -plane. Write

$$\text{width}(s) := w^+(s) - w^-(s)$$

for the band's height. One more elementary bound, used twice below: since $q_i \leq i$ always, every configuration's value is at most $\sum_t i_t = s$, so

$$w^+(s) \leq s; \tag{6}$$

and taking the configuration $(s, 0, \dots, 0)$ shows $w^+(s) \geq q_s$, which tends to infinity because the letter 1 occurs infinitely often.

12.2. The sweep. Now the payoff. Recall from (5) that realizing a target integer x means writing $x = \delta s + ew$ with $w \in W_s$. If the band is *wide* — as wide as the alphabet spread requires — then every large x is realized:




FIGURE 16. The sweep. The reachable band (shaded, endpoints $w^\pm$) rises gently as the budget s grows; the required height $w(s)$ for a fixed target x descends along the admissible budgets. If the band is wide, the descending line cannot skip it: some admissible s^* lands inside, and there x is realized.

Lemma 12.2 (Sweep). *Suppose there is an S_0 with $\text{width}(s) \geq d_2 - 1$ for all $s \geq S_0$. Then $kA - ka_1$ contains every sufficiently large integer; in particular A is tail-covering with modulus $m = 1$.*

Proof. Fix a large integer x (thresholds at the end). Which budgets s could realize x ? From $x = \delta s + ew$ we need $\delta s \equiv x \pmod{e}$, and since $\gcd(\delta, e) = 1$ this pins s to a single residue class modulo e ; call such s *admissible*, and for admissible s set

$$w(s) := \frac{x - \delta s}{e} \in \mathbb{Z},$$

the height at which the band would have to be hit. As s steps up by e (to the next admissible budget), $w(s)$ steps *down* by exactly δ . Meanwhile, by Lemma 12.1, each band endpoint rises by between 0 and e over those e unit steps. A descending target against a rising band: the picture from Figure 16 is that they must cross, and that a wide band cannot be jumped over in a single stride. We make both halves precise.

The target starts above the band and ends below it. At the small end: let $s_{\min}$ be the smallest admissible $s \geq S_0$; then $w(s_{\min}) \geq (x - \delta(S_0 + e))/e$, which for large x exceeds $s_{\min}$, hence exceeds $w^+(s_{\min})$ by (6) — explicitly, $x > (\delta + e)(S_0 + e)$ suffices. At the large end: enlarge S_0 once and for all so that $w^+(s) \geq \delta$ for every $s \geq S_0$ (possible since $w^+(s) \geq q_s \rightarrow \infty$; this depends only on the word, not on x). For the largest admissible s with $w(s) \geq 0$, the next admissible step would take w negative, so $0 \leq w(s) < \delta \leq w^+(s)$: at that budget the target sits at or below the band's top.

No skipping. Suppose, for contradiction, that no admissible $s \geq S_0$ has $w(s) \in [w^-(s), w^+(s)]$. The target descends (by δ per admissible step) from above the band to below its top; since it never lands inside, there must be two *consecutive* admissible budgets s' and $s' + e$ with

$$w(s') > w^+(s') \quad \text{and} \quad w(s' + e) < w^-(s' + e) :$$

the leap over the band happens between them. First, s' is large: from $w(s'+e) \leq w^-(s'+e) \leq w^+(s'+e) \leq s'+e$ (again (6)),

$$\frac{x - \delta(s'+e)}{e} \leq s'+e \implies s'+e \geq \frac{x}{\delta+e},$$

so $s' \geq S_0$ holds automatically once $x \geq (\delta+e)(S_0+e)$ — the same threshold as before, so no new constraint. Now stack the inequalities. The target drops by exactly δ across the leap, and the band's floor rises by at most e :

$$w^+(s') < w(s') = w(s'+e) + \delta \leq w^-(s'+e) - 1 + \delta \leq w^-(s') + e - 1 + \delta.$$

Hence

$$\text{width}(s') = w^+(s') - w^-(s') \leq \delta + e - 2 = \delta' - 2 \leq d_2 - 2,$$

contradicting the hypothesis $\text{width}(s') \geq d_2 - 1$. So some admissible $s^* \geq S_0$ has $w(s^*) \in [w^-(s^*), w^+(s^*)]$; by Lemma 12.1 the integer $w(s^*)$ is actually attained, and $x = \delta s^* + e w(s^*) \in kA - ka_1$. $\square$

Look where the alphabet bound was spent: the contradiction needed $\text{width} \geq \delta' - 1$, and the hypothesis supplies $d_2 - 1 \geq \delta' - 1$. A band exactly one narrower could be leapt over — the sweep operates with no room to spare, a first taste of how tight the corner $d_2 = k$ will be.

12.3. Wide or balanced: the width dichotomy. The sweep shifts the burden: we must now show the band is wide. Width comes from *irregularity* of the word h — if some stretches are 1-rich and others 1-poor, configurations can exploit both. The cleanest measure uses just two indices. For $\sigma \geq 0$ define the *two-index width*

$$W_2(\sigma) := \max\{q_i + q_j : i + j = \sigma, i, j \geq 0\} - \min\{q_i + q_j : i + j = \sigma, i, j \geq 0\}.$$

One bookkeeping lemma first — trivial to prove, used four times in this part, and worth having on the record rather than waving at.

Lemma 12.3 (Splitting). *Split the k slots of a configuration into groups, and prescribe for each group the subtotal of its indices (subtotals summing to s). The groups then choose their indices independently, so among prescribed configurations the largest (smallest) achievable value is the sum of the groups' separate maxima (minima). And since prescribed configurations are particular configurations with total s ,*

$$\text{width}(s) \geq \sum_{\text{groups}} (\text{group max} - \text{group min}).$$

Proof. The value $\sum_t q_{it}$ is the sum of the groups' values, and no constraint couples one group's index choices to another's — so extremes add. For the display: $w^+(s)$ is a maximum over *all* configurations, hence at least the prescribed maximum; symmetrically $w^-(s)$ is at most the prescribed minimum. $\square$

Lemma 12.4 (Width dichotomy).

- (a) Suppose $W_2(\sigma_0) \geq 2$ for some σ_0 , and suppose k is odd or $d_2 \leq k-1$. Then $\text{width}(s) \geq d_2 - 1$ for all sufficiently large s — and the sweep applies.
- (b) If instead $W_2(\sigma) \leq 1$ for every σ , then the word h is balanced.

Proof. (a) Both parities are one application of the splitting lemma. Note that a two-slot group with prescribed subtotal σ_0 has group max minus group min exactly $W_2(\sigma_0) \geq 2$.

k odd. Write $k = 2u + 1$. For any $s \geq u\sigma_0$, split into u two-slot groups prescribed at σ_0 and one single-slot group at $j := s - u\sigma_0 \geq 0$ (a single slot has spread 0). By Lemma 12.3,

$$\text{width}(s) \geq u W_2(\sigma_0) \geq 2u = k - 1 \geq d_2 - 1,$$

using $d_2 \leq k$.

k even, $d_2 \leq k - 1$. Write $k = 2u$. For $s \geq (u - 1)\sigma_0$, split into $u - 1$ two-slot groups at σ_0 and one two-slot group at $\tau := s - (u - 1)\sigma_0 \geq 0$, whose spread $W_2(\tau)$ is at least 0. By Lemma 12.3,

$$\text{width}(s) \geq (u - 1) W_2(\sigma_0) + W_2(\tau) \geq 2u - 2 = k - 2 \geq d_2 - 1.$$

(b) Take two equal-length windows, from i to i' and from j' to j (so $i' - i = j - j'$, i.e. $i + j = i' + j' =: \sigma$). Then

$$|(q_{i'} - q_i) - (q_j - q_{j'})| = |(q_{i'} + q_{j'}) - (q_i + q_j)| \leq W_2(\sigma) \leq 1,$$

which is precisely balancedness. $\square$

Part (a)'s parity split is not laziness. When k is even and $d_2 = k$, the budget of k indices is exactly wrong: u full pairs use all k indices and give width $\geq k$ — but only when s has the right form — while $u - 1$ pairs guarantee only $k - 2 = d_2 - 2$, one short of what the sweep needs. This single missing unit of width is a genuine boundary phenomenon, and §14 will close it with a separate argument (Lemma 14.5) rather than a sharper pairing.

Where we stand. If some $W_2(\sigma_0) \geq 2$ and (k odd or $d_2 \leq k - 1$): done, by (a) and the sweep. If every $W_2(\sigma) \leq 1$: the word is balanced, and the next section classifies it — eventually periodic (done, by §10) or Sturmian (handled in §14). The one remaining configuration, k even with $d_2 = k$ and the word unbalanced, is also settled in §14. The ledger stays honest.

Exercise 12.5 (CORE). For the periodic word $h = 0011 0011 \dots$ (which Example 11.2 showed is unbalanced), find a σ_0 with $W_2(\sigma_0) \geq 2$, and exhibit, for $k = 3$, $s = 12$, two configurations whose values differ by at least 2.

Exercise 12.6 (CORE). W_2 uses exactly two indices while width uses k ; compare them. Prove that $\text{width}(s) \geq W_2(\sigma)$ whenever $s \geq \sigma$ and $k \geq 2$ (pad the two-index configurations with zeros — $q_0 = 0$ costs nothing), and find a word and an s with $\text{width}(s) > W_2(s)$.

13. BALANCED WORDS: THE MORSE–HEDLUND CLASSIFICATION

The width dichotomy left us with the balanced words, and this section proves the structure theorem that tames them: *a balanced binary word is either eventually periodic, or — after discarding at most one initial stretch — is “mechanical”: its counting function is $q_n = \lfloor n\alpha + \beta \rfloor$ for a fixed irrational slope α .* This is (one direction of) a celebrated 1940 theorem of Morse and Hedlund. The proof below is elementary and self-contained: its only imported ingredient is the completeness of the reals (through sup), and the section closes with the one approximation tool (Dirichlet, by pigeonhole) that the endgame will need.

Throughout, h is a balanced binary word with counting function q ($q_0 = 0$).

13.1. Step 1: a balanced word is nearly additive.

Lemma 13.1 (Near-additivity). *For all $n, m \geq 0$: $|q_{n+m} - q_n - q_m| \leq 1$.*

Proof. $q_{n+m} - q_n$ counts the 1s in the window of length m starting after position n ; q_m counts them in the window of length m at the very start. Balance bounds the difference by 1. $\square$

Iterating pins down long-range behavior:

Lemma 13.2. *For all $n \geq 0, t \geq 1$: $tq_n - (t-1) \leq q_{tn} \leq tq_n + (t-1)$.*

Proof. Induct on t . Trivial at $t = 1$. If it holds at t , then by Lemma 13.1 (with the pair tn, n), $q_{(t+1)n} \leq q_{tn} + q_n + 1 \leq (t+1)q_n + t$, and similarly from below. $\square$

13.2. Step 2: the slope. The slope we are after *ought* to be the density of 1s in the word — the limit of q_n/n . But we cannot speak of “the limit” before proving one exists, so we build the number as a supremum and only afterwards verify that it deserves the name. For orientation, the values on our running examples: the alternating word $0101 \dots$ will get $\alpha = \frac{1}{2}$; the word $001001 \dots$ gets $\alpha = \frac{1}{3}$; the Fibonacci word gets $\alpha = (3 - \sqrt{5})/2 \approx 0.382$.

Lemma 13.3 (The slope exists). *There is a real number $\alpha \in [0, 1]$ such that*

$$|D_n| \leq 1 \quad \text{for every } n \geq 0, \quad \text{where } D_n := q_n - n\alpha.$$

Moreover $q_n/n \rightarrow \alpha$.

Proof. The key inequality: for all $n, m \geq 1$,

$$mq_n - nq_m \leq n + m - 2. \tag{7}$$

The idea behind it: *count the 1s among the first nm letters two ways* — as m blocks of length n , and as n blocks of length m — and let near-additivity certify that both counts are accurate to small errors. Formally, bound q_{nm} twice by Lemma 13.2: viewing nm as m copies of n , $q_{nm} \geq mq_n - (m-1)$; viewing it as n copies of m , $q_{nm} \leq nq_m + (n-1)$. Chaining the two gives (7).

Dividing (7) by nm :

$$\frac{q_n - 1}{n} - \frac{q_m + 1}{m} = \frac{mq_n - nq_m - (n+m)}{nm} \leq \frac{-2}{nm} < 0.$$

So *every* number of the form $(q_n - 1)/n$ is below *every* number of the form $(q_m + 1)/m$. The set $\{(q_n - 1)/n : n \geq 1\}$ is therefore bounded above; let

$$\alpha := \sup_{n \geq 1} \frac{q_n - 1}{n}.$$

By construction $(q_n - 1)/n \leq \alpha$ for every n ; and $\alpha \leq (q_m + 1)/m$ for every m , because $(q_m + 1)/m$ is an upper bound of the set. The two together read $q_n - 1 \leq n\alpha \leq q_n + 1$, i.e. $|D_n| \leq 1$ (also at $n = 0$, where $D_0 = 0$). Since $0 \leq q_n \leq n$, we get $(q_n - 1)/n < 1$ for all n , so $\alpha \leq 1$; and $\alpha \geq (q_n - 1)/n \geq -1/n$ for every n , so $\alpha \geq 0$. Finally $|q_n/n - \alpha| = |D_n|/n \leq 1/n \rightarrow 0$. $\square$

The number α is the *slope* of h : the limiting density of 1s. The inequality $|D_n| \leq 1$ says the word tracks the line $y = \alpha x$ with error never exceeding 1 — already a strong rigidity statement, extracted from balance alone, and worth seeing (Figure 17).




FIGURE 17. The counting function of the Fibonacci word ($\alpha = (3 - \sqrt{5})/2$), trapped forever in the shaded band $\alpha n \pm 1$: Lemma 13.3 drawn. The dots drift within the band — that drift is the discrepancy D_n — but Lemma 13.4 will confine even the drift to a window of height 1.

13.3. Step 3: the discrepancy barely oscillates.

Lemma 13.4 (Oscillation bound). *For all $n, m \geq 0$: $|D_n - D_m| \leq 1$.*

Proof. Say $m = n + \Delta$ with $\Delta \geq 1$. For $j \geq 0$ let $W(j) := q_{j+\Delta} - q_j$ count the 1s in the length- Δ window starting after j ; balance gives $|W(j) - W(n)| \leq 1$ for every j . Sum W over the t consecutive disjoint windows starting at $0, \Delta, 2\Delta, \dots, (t-1)\Delta$; the sum telescopes to $q_{t\Delta}$, and each summand lies within 1 of $W(n)$:

$$t(W(n) - 1) \leq q_{t\Delta} \leq t(W(n) + 1).$$

Divide by $t\Delta$ and let $t \rightarrow \infty$: the middle tends to α (Lemma 13.3), so by the limit facts of §11,

$$\frac{W(n) - 1}{\Delta} \leq \alpha \leq \frac{W(n) + 1}{\Delta}, \quad \text{i.e.} \quad |W(n) - \Delta\alpha| \leq 1.$$

And $D_{n+\Delta} - D_n = (q_{n+\Delta} - q_n) - \Delta\alpha = W(n) - \Delta\alpha$. □

So all the discrepancies $D_0, D_1, D_2, \dots$ live in a band of height 1 — they may drift within it, but no two ever differ by more than 1. Everything below flows from this single fact.

13.4. Step 4: rational slope forces eventual periodicity.

Proposition 13.5. *If α is rational, h is eventually periodic.*

Proof. Write $\alpha = u/v$ in lowest terms: u, v integers, $v \geq 1$, $\gcd(u, v) = 1$ (divide numerator and denominator by their gcd). We avoid the letters p, q here — q_n is already the counting function. Set

$$c_n := v q_n - u n = v D_n,$$




FIGURE 18. The second costume. The staircase $q_n = \lfloor n\alpha + \beta \rfloor$ tracks the line, and the word beneath records where it *steps* (a 1) versus where it *waits* (a 0). Figure 6 showed the same object as a walk on the number line: a Beatty walk's gaps and a mechanical word's letters are two readings of one digitized line.

an integer with $|c_n| \leq v$ (Lemma 13.3). Let $w_n := q_{n+v} - q_n$ (the 1-count of the length- v window after n), so that

$$c_{n+v} - c_n = v(w_n - u).$$

Balance confines the w_n to two consecutive values $\{W, W+1\}$ (all length- v windows agree to within one). We claim $u \in \{W, W+1\}$: if every $w_n \geq u+1$, then $c_{n+v} - c_n \geq v$ for every n , so following the indices $n, n+v, n+2v, \dots$ the integer c climbs by at least v per step, exiting $[-v, v]$ — impossible; symmetrically, $w_n \leq u-1$ for all n is impossible. So either $w_n \in \{u, u+1\}$ for all n , or $w_n \in \{u-1, u\}$ for all n .

In the first case, $c_{n+v} - c_n \in \{0, v\}$: along each of the v index classes modulo v , the sequence c is nondecreasing and bounded, hence eventually constant; from some point on, no step of $+v$ ever occurs, i.e. $w_n = u$ for all large n . The second case is the mirror image (nonincreasing), with the same conclusion. So $q_{n+v} = q_n + u$ for all large n , and then

$$h_{n+v+1} = q_{n+v+1} - q_{n+v} = (q_{n+1} + u) - (q_n + u) = h_{n+1}$$

for all large n : the word repeats with period v from some point on. $\square$

13.5. Step 5: irrational slope forces a mechanical tail.

Definition 13.6 (Mechanical word). A binary word h is *mechanical* with slope α and intercept β if

$$q_n = \lfloor n\alpha + \beta \rfloor \quad \text{for every } n \geq 1.$$

Mechanical words with irrational slope are called *Sturmian*.

A mechanical word is the “digitized straight line” of slope α : its n -th letter records whether the line $y = \alpha x + \beta$ crosses an integer height in the n -th unit interval. Beatty walks and mechanical words are two costumes of the same actor — compare Definition 3.4 and Figure 6 with Figure 18. In Part II we *chose* such a line and rode it to victory; here the Adversary's walk turns out to be riding one whether it meant to or not, and that rigidity is what we will exploit against it.

Proposition 13.7. *If α is irrational, then there is an index $n_0 \geq 0$ such that the shifted word $h'_j := h_{n_0+j}$ is mechanical with slope α (and some intercept β').*

Proof. Set $\beta := \sup_n D_n$, finite since $|D_n| \leq 1$. By the oscillation bound, every $D_n \geq \beta - 1$: for any m , $D_m - D_n \leq 1$, and taking the supremum over m gives $\beta - D_n \leq 1$.

When does the mechanical formula hold at n ? Unwinding the floor's definition,

$$q_n = \lfloor n\alpha + \beta \rfloor \iff n\alpha + \beta - 1 < q_n \leq n\alpha + \beta \iff \beta - 1 < D_n \leq \beta.$$

The upper half is the definition of β ; the lower half holds at least weakly ($D_n \geq \beta - 1$), so the formula can fail only where $D_n = \beta - 1$ exactly — equivalently where $n\alpha + \beta = q_n + 1$ is an integer. If two distinct indices $n_1 \neq n_2$ both had $n_i\alpha + \beta \in \mathbb{Z}$, subtracting would give $(n_1 - n_2)\alpha \in \mathbb{Z}$, impossible for irrational α . So there is at most one exceptional index; call it n_0 (and set $n_0 := 0$ if there is none).


Discard everything through n_0 : the word $h'_j := h_{n_0+j}$ has counting function $q'_j = q_{n_0+j} - q_{n_0}$, and for every $j \geq 1$ the index $n_0 + j$ is unexceptional, so

$$q'_j = \lfloor (n_0 + j)\alpha + \beta \rfloor - q_{n_0} = \lfloor j\alpha + (n_0\alpha + \beta - q_{n_0}) \rfloor = \lfloor j\alpha + \beta' \rfloor, \quad \beta' := \beta - D_{n_0},$$

where the middle equality moves the *integer* q_{n_0} through the floor (Lemma 3.2(2)). $\square$

Assembling Steps 4 and 5:

Theorem 13.8 (Balanced classification). *A balanced binary word is either eventually periodic, or agrees beyond some index with a Sturmian word: a mechanical word of irrational slope $\alpha \in (0, 1)$.*




The five steps at a glance. Balance alone forces the word to track a line; the line's slope then decides everything, and the two verdicts are exactly the two branches §12 needs — one already closed by §10, the other handed to §14.

Proof. Steps 2–5; when α is irrational it lies in $(0, 1)$, being in $[0, 1]$ and not equal to 0 or 1. $\square$

For our purposes the theorem is a perfect fork: the periodic branch was closed in §10, so *the only balanced words still standing are the Sturmian tails*, with an explicit formula $q_n = \lfloor n\alpha + \beta \rfloor$ to attack. The attack needs one lever: the ability to find indices n where the fractional part $\{n\alpha + \beta\}$ lands in any prescribed arc. That is a statement about *rotations of a circle*, and it closes this section.

13.6. Rotations, and how to hit an arc on schedule. For a real x write $\{x\} := x - \lfloor x \rfloor \in [0, 1)$ for its *fractional part*. Picture the interval $[0, 1)$ bent into a circle of circumference 1; the map $x \mapsto \{x + \alpha\}$ is rotation by α . The numbers $\{n\alpha + \beta\}$ are the successive positions of a point rotating by α per tick. Write $\|x\|$ for the distance from x to the nearest integer — the circle distance from $\{x\}$ to 0.



equal steps $\theta < |J|$: the walk laps the circle and cannot jump the arc

FIGURE 19. The circle walk of Lemma 13.10. Successive points x_j advance by the same short step θ in the same direction; after $\lceil 1/\theta \rceil$ steps the total distance walked is at least a full lap, and a step shorter than the arc cannot cross it without landing in it.

Lemma 13.9 (Dirichlet's approximation theorem). *For every real α and integer $Q \geq 1$ there are integers n_0, m with $1 \leq n_0 \leq Q$ and*

$$|n_0\alpha - m| < \frac{1}{Q}.$$

Proof. Pigeonhole. The $Q + 1$ numbers $\{0 \cdot \alpha\}, \{1 \cdot \alpha\}, \dots, \{Q\alpha\}$ all lie in $[0, 1)$, which we cut into Q boxes $[i/Q, (i + 1)/Q)$. Two of them share a box: $\{j\alpha\}$ and $\{j'\alpha\}$ with $j < j'$, differing by less than $1/Q$. Set $n_0 := j' - j \in [1, Q]$; then $n_0\alpha = (j'\alpha - \lfloor j'\alpha \rfloor) - (j\alpha - \lfloor j\alpha \rfloor) + (\lfloor j'\alpha \rfloor - \lfloor j\alpha \rfloor) = \{j'\alpha\} - \{j\alpha\} + m$ for the integer $m := \lfloor j'\alpha \rfloor - \lfloor j\alpha \rfloor$, and $|n_0\alpha - m| = |\{j'\alpha\} - \{j\alpha\}| < 1/Q$. $\square$

Lemma 13.10 (Uniform hitting). *Let $\alpha \in (0, 1)$ be irrational and $\beta \in \mathbb{R}$. For every $\varepsilon > 0$ there is an integer $L = L(\alpha, \varepsilon)$ — depending only on α and ε , not on β and not on the arc's position — such that for every arc $J \subseteq (0, 1)$ of length $\geq \varepsilon$ and every $N \geq 0$, some index $n \in \{N + 1, \dots, N + L\}$ has $\{n\alpha + \beta\} \in J$.*

Proof. An explicit circle walk. Choose an integer $Q > 2/\varepsilon$ and apply Dirichlet: there are $1 \leq n_0 \leq Q$ and m with $\theta := |n_0\alpha - m| < 1/Q < \varepsilon/2$; and $\theta \neq 0$ since α is irrational. Advancing the index by n_0 rotates the point $\{n\alpha + \beta\}$ by exactly θ around the circle — always in the same direction (the sign of $n_0\alpha - m$ is fixed).

Now fix N , and follow the $\lceil 1/\theta \rceil + 1$ indices

$$n_j := N + (j + 1)n_0 \quad (j = 0, 1, \dots, \lceil 1/\theta \rceil),$$

all lying in $\{N + 1, \dots, N + L\}$ for $L := n_0(\lceil 1/\theta \rceil + 1)$. The successive positions $x_j := \{n_j\alpha + \beta\}$ each step a fixed distance θ around the circle, in a fixed direction, and over $\lceil 1/\theta \rceil$ steps the total distance walked is at least 1: the walk laps the whole circle. If none of the x_j lay in the arc J , then some single step would have to jump clean over J — but each step has length $\theta < \varepsilon/2 < |J|$, too short to clear it. So some $x_j \in J$ (Figure 19).

Finally, audit the dependence: Q came from ε ; n_0, m, θ from α and Q ; and $L = n_0(\lceil 1/\theta \rceil + 1) \leq Q(\lceil 1/\theta \rceil + 1)$ from these. Neither β , nor N , nor where J sits was used — only $|J| \geq \varepsilon$. $\square$

The uniformity is the point: in the endgame the arc J will *move* as the target changes, and we need a single window length L serving every position. A subtler equidistribution theorem would say more; this crude walking bound, powered by one pigeonhole, says exactly enough.

Exercise 13.11 (CORE). (a) For the periodic balanced word $001001\cdots$, compute α and check $|D_n| \leq 1$ and the conclusion of Step 4 directly. (b) For $\alpha = \sqrt{2} - 1$ and $Q = 5$, find the $n_0 \leq 5$ and θ that Dirichlet’s pigeonhole produces, and the resulting L for $\varepsilon = \frac{1}{5}$.

Exercise 13.12 (CORE). Where exactly did Step 5 use irrationality? (Two places. Find both, and check the argument really does collapse for rational α — as it must, since periodic words are not mechanical-with-irrational-slope.)

14. THE STURMIAN ENDGAME

This section finishes the two-letter case: first the Sturmian walks (the hardest single proof in this paper — take it slowly), then the boundary configuration that the width dichotomy could not reach, then the assembled two-letter theorem.

14.1. The Sturmian lemma. We keep the two-letter coordinates of §12 — letters $\delta < \delta' = \delta + e$ with $\gcd(\delta, e) = 1$ and $\delta' \leq d_2 \leq k$, binary word h , counting function q — and now assume, after Theorem 13.8 and a final tail-pass, that the word is Sturmian:

$$q_n = \lfloor n\alpha + \beta \rfloor \quad (n \geq 1)$$

for an irrational $\alpha \in (0, 1)$ and some real β .

Lemma 14.1 (Sturmian). *In this situation $kA - ka_1$ contains every sufficiently large integer; in particular A is tail-covering with modulus 1.*

This is stronger than the sweep gave us — *all* large integers, no residue class needed — and it has to be, because Sturmian words are exactly the words for which the band of §12 is too narrow for the sweep. The proof replaces “wide band” with something better: *exact control*. For a mechanical word the value of a configuration is not merely trapped in a window; it can be steered onto a precise target, because choosing indices means choosing points of an irrational rotation, and the hitting lemma lets us place those points wherever we like.

Proof. Recall (Lemma 9.1 and (5)) that it suffices to realize each large integer x as $x = \delta s + e w$ where w is a sum of k counting-function values $q_{i_1} + \cdots + q_{i_k}$ with $i_t \geq 1$ and $i_1 + \cdots + i_k = s$. The proof has four movements.

*Step 0: where the values live.*³ Write, for $n \geq 1$,

$$q_n = n\alpha + \beta - \{n\alpha + \beta\}$$

(floor = number minus fractional part). Summing over k indices with total s :

$$\sum_{t=1}^k q_{i_t} = X_s - \sum_{t=1}^k \{i_t\alpha + \beta\}, \quad X_s := s\alpha + k\beta. \quad (8)$$

Each fractional part lies in $[0, 1)$, so the sum of k of them lies in $[0, k)$: every achievable value sits in the real window $(X_s - k, X_s]$ — the two-letter avatar of Part II’s window picture. But now we run the logic *forward*: to achieve a prescribed value w at depth $\theta := X_s - w$ below the window’s top, it suffices to make the fractional parts sum to *exactly* θ .

³One convention shifts here, and it is deliberate. Equation (5) let the summand indices i_t range over *all* integers ≥ 0 ; from this proof onward we use only indices $i_t \geq 1$. The reason is that the Sturmian formula $q_n = \lfloor n\alpha + \beta \rfloor$ was established for $n \geq 1$ only (Proposition 13.7 says nothing at $n = 0$), so index 0 is not covered by it. Restricting is harmless: we are *constructing* representations, so using fewer indices than we are entitled to can only make the conclusion stronger.

Step 1: the tool. Lemma 13.10: for any arc length $\varepsilon > 0$ there is a window length $L(\alpha, \varepsilon)$ such that every L consecutive indices contain one whose rotation point $\{n\alpha + \beta\}$ lies in any prescribed arc of length ε — uniformly in the arc’s position.

Step 2: hitting any depth in the middle of the window. Fix a margin $\zeta \in (0, \frac{1}{2})$, to be chosen in Step 3. **Claim:** there is a threshold S_0 (depending on ζ, α, β, k only) such that for every $s \geq S_0$ and every integer w whose depth $\theta = X_s - w$ lies in $[\zeta, k - \zeta]$, there are *distinct* indices $i_1 < \dots < i_k$ (all ≥ 1 , summing to s) with

$$\sum_{t=1}^k \{i_t \alpha + \beta\} = \theta \quad \text{exactly,} \quad \text{hence} \quad \sum_{t=1}^k q_{i_t} = w.$$

The mechanism: place the first $k - 1$ points of the rotation in a narrow arc positioned so their sum falls just below θ , then let the forced last index close the books — with an integrality argument snapping the total onto θ exactly. Because this step introduces several objects in quick succession, here is the cast before the play:

Object	Job	Depends on
ζ	margin: how deep inside the window we work	fixed in Step 3
$\kappa = \zeta/4$	safety padding inside every estimate	ζ
$\mathcal{I}$	interval of legal per-point targets	θ
T^*	the chosen per-point target	θ
J	the arc where the $k-1$ free points must land	T^*
$\varepsilon_0 = \frac{2\kappa}{k-1}$	J ’s length — position moves, length never	ζ, k
L	hitting window for arcs of length ε_0	α, ε_0
S_0	budget threshold making the last index behave	L, k

Set $\kappa := \zeta/4$, and consider the open interval of candidate “per-point targets”

$$\mathcal{I} := (\max(0, \theta - 1) + \kappa, \min(k - 1, \theta) - \kappa).$$

Its length is at least $\zeta/2$ whatever the regime of θ — three cases, by which of the max/min clauses bind:

- $\zeta \leq \theta \leq 1$: $|\mathcal{I}| = \theta - 2\kappa \geq \zeta - \frac{\zeta}{2} = \frac{\zeta}{2}$;
- $1 < \theta \leq k - 1$: $|\mathcal{I}| = 1 - 2\kappa = 1 - \frac{\zeta}{2} > \frac{\zeta}{2}$ (as $\zeta < 1$);
- $k - 1 < \theta \leq k - \zeta$: $|\mathcal{I}| = (k - \theta) - 2\kappa \geq \zeta - \frac{\zeta}{2} = \frac{\zeta}{2}$.

Pick any $T^* \in \mathcal{I}$ and set the arc

$$J := \left(\frac{T^*}{k-1} - \frac{\kappa}{k-1}, \frac{T^*}{k-1} + \frac{\kappa}{k-1} \right) \subseteq (0, 1),$$

of length $\varepsilon_0 := 2\kappa/(k-1)$, and let $L := L(\alpha, \varepsilon_0)$ be its hitting constant. The arc’s *position* depends on θ (through T^*), but its *length* does not — and the hitting lemma is uniform in position, so this single L serves every target and every s . (This is exactly why we insisted on uniformity in §13.)

Choosing the free indices. For $j = 1, \dots, k - 1$, the hitting lemma finds an index i_j in the block $\{(j-1)L + 1, \dots, jL\}$ with $\{i_j \alpha + \beta\} \in J$ (Figure 20). Different blocks, so $i_1 < i_2 < \dots < i_{k-1}$, with

$$i_j \leq jL, \quad \sum_{j=1}^{k-1} i_j \leq L \frac{k(k-1)}{2}, \quad i_{k-1} \leq (k-1)L. \quad (9)$$




FIGURE 20. How Step 2 chooses its indices. *Right*: the $k-1$ free indices are drawn one from each of $k-1$ disjoint blocks of length L — disjointness makes them distinct — and the last index is then forced by the budget s . Taking $s \geq S_0$ pushes i_k past all of them. *Left*: the hitting lemma is what lets each free index be chosen with its rotation point inside the short arc J , so their fractional parts sum to within κ of the target; the forced point's fractional part is uncontrolled, and integrality (Figure 21) cleans up the remainder.

The forced last index. What must the threshold S_0 accomplish? The last index $i_k := s - (i_1 + \dots + i_{k-1})$ is forced, and we need it both legal (≥ 1) and distinct from the others, for which $i_k > i_{k-1}$ suffices. By (9) the free indices consume at most $Lk(k-1)/2$ of the budget, and the largest of them is at most $(k-1)L$ — so it suffices to start s strictly above the sum of those two bounds. Accordingly, set

$$S_0 := 1 + L \frac{k(k-1)}{2} + (k-1)L = 1 + L(k-1)\left(\frac{k}{2} + 1\right)$$

(an integer, as $(k-1)(k+2)/2$ has one even factor). For $s \geq S_0$, by (9),

$$i_k \geq S_0 - L \frac{k(k-1)}{2} = 1 + L(k-1)\left(\frac{k}{2} + 1 - \frac{k}{2}\right) = 1 + (k-1)L > i_{k-1},$$

so all k indices are distinct and ≥ 1 , with total exactly s .

The snap. Let $F := \sum_{j \leq k-1} \{i_j \alpha + \beta\}$; since each term lies in J ,

$$F \in (T^* - \kappa, T^* + \kappa) \subseteq (\max(0, \theta - 1), \min(k-1, \theta)),$$

the inclusion because $T^* \in \mathcal{I}$ keeps distance $> \kappa$ from both endpoints.

[Instance, with Example 14.2's Fibonacci data ($\alpha \approx 0.382$, $\zeta = \alpha/4 \approx 0.095$, $\kappa \approx 0.024$): for the target of that example, $\theta \approx 0.647$, so $\mathcal{I} \approx (0.024, 0.623)$; choosing $T^* = 0.3$ puts the arc at $J \approx (0.138, 0.162)$, and the two free rotation points must land in that short arc.] Add the last point: $\Theta := F + \{i_k \alpha + \beta\}$, with the fractional part in $[0, 1)$, satisfies

$$\Theta > \max(0, \theta - 1) \geq \theta - 1 \quad \text{and} \quad \Theta < \min(k-1, \theta) + 1 \leq \theta + 1.$$

On the other hand, by (8) applied to our indices,

$$\Theta = X_s - \sum_{t \leq k} q_{i_t},$$

and both $X_s - \theta = w$ and $\sum_t q_{i_t}$ are integers, so

$$\Theta - \theta = (X_s - \sum_t q_{i_t}) - (X_s - w) = w - \sum_t q_{i_t} \in \mathbb{Z}.$$

Now finish in three short steps. The difference $\Theta - \theta$ is an integer. From $\theta - 1 < \Theta < \theta + 1$ we get $-1 < \Theta - \theta < 1$. The only integer strictly between -1 and 1 is 0 . Therefore $\Theta - \theta = 0$, i.e. $\Theta = \theta$ — and consequently $\sum_t q_{i_t} = X_s - \theta = w$, exactly as required. The Claim is proved.




FIGURE 21. The snap. The open interval $(\theta - 1, \theta + 1)$ contains exactly one point of the lattice $\theta + \mathbb{Z}$: its center. So $\Theta = \theta$, with no error at all — the estimates never had to be sharp, only sharper than the lattice spacing.

Big idea. *The snap* (Figure 21) is the step worth slowing down for, so do not let it pass in four lines: soft estimates confined Θ to an open interval of length 2, and hard integrality confined $\Theta - \theta$ to the integers; the only point obeying both masters is θ exactly. Analysis narrows the suspects to a lineup; arithmetic makes the identification.

(One aside before the ladder: why did Step 2 bother producing *distinct* indices, when Definition 1.6 happily permits repeats? Because distinctness came free — the $k - 1$ free indices live in disjoint blocks — and the argument for pinning down i_k was easiest to state as $i_k > i_{k-1}$. Nothing was lost and nothing extra was needed; distinctness is simply the stronger conclusion the construction happened to deliver.)

Step 3: the ladder. Now fix the margin

$$\zeta := \frac{1}{4} \min(\alpha, 1 - \alpha) \in (0, \frac{1}{8}],$$

and let S_0 be the resulting threshold. Let x be a large integer. As in the sweep, the budgets s able to represent $x = \delta s + ew$ with integer w form one residue class modulo e (the *admissible* s), and for admissible s we set $w(s) := (x - \delta s)/e$ and measure the depth

$$\theta(s) := X_s - w(s) = s\alpha + k\beta - \frac{x - \delta s}{e}.$$

Per admissible step $s \mapsto s + e$, the value X_s grows by $e\alpha$ and $w(s)$ falls by δ , so the depth climbs by exactly

$$\psi := e\alpha + \delta > 0$$

— an arithmetic ladder. Step 2 accepts any depth in the strip $[\zeta, k - \zeta]$, of height $k - 2\zeta$. The ladder cannot step over the strip, because its rung spacing is smaller:

$$\psi = \delta + e\alpha = \delta' - e(1 - \alpha) \leq k - (1 - \alpha) < k - 2\zeta,$$

using $\delta' \leq k$, $e \geq 1$, and $2\zeta \leq \frac{1}{2}(1 - \alpha) < 1 - \alpha$. Recognize this inequality: the rung $\psi = \delta + e\alpha$ is precisely the walk's *average step*, and the strip height is essentially k — “rungs shorter than the strip” is the triptych of Figure 4 all over again, stride against window width, now *proved* for a wobbling walk rather than pictured for a steady one. Since $\theta(s)$ is very negative for small admissible s (there $w(s) \approx x/e$ dominates) and climbs unboundedly by steps $\psi < k - 2\zeta$, some admissible s^* lands in the strip: $\theta(s^*) \in [\zeta, k - \zeta]$ (Figure 22).

The thresholds, explicitly. Multiplying the definition of θ by e and rearranging: $s\psi = x + e\theta(s) - ek\beta$. So any s^* with $\theta(s^*) \in [0, k]$ satisfies

$$\frac{x - ek|\beta|}{\psi} \leq s^* \leq \frac{x + ek(1 + |\beta|)}{\psi}.$$

Two requirements: (a) $s^* \geq S_0$, so Step 2 applies — since $\psi \leq \delta + e = \delta' \leq k$, the left bound gives this once $x \geq kS_0 + ek|\beta|$. (b) $w(s^*) \geq 0$, i.e. $\delta s^* \leq x$ — by the right bound this holds




FIGURE 22. The ladder against the strip, drawn with the Fibonacci numbers of Example 14.2: rung spacing $\psi = 1 + \alpha \approx 1.38$, strip height $k - 2\zeta \approx 2.81$. A ladder whose rungs are closer together than the strip is tall cannot pass from below it to above it without a rung landing inside.

once $\delta(x + ek(1 + |\beta|)) \leq (\delta + e\alpha)x$, i.e. once $x \geq \delta k(1 + |\beta|)/\alpha$ (here $\alpha > 0$ earns its keep). So for

$$x \geq X_0 := 1 + \max\left(kS_0 + ek|\beta|, \frac{\delta k(1 + |\beta|)}{\alpha}\right),$$

Step 2 realizes $w(s^*) \in W_{s^*}$ with legal indices, and

$$x = \delta s^* + ew(s^*) \in kA - ka_1. \quad \square$$

Big idea. For a Sturmian walk the window picture stops being an estimate and becomes an equation: the depth of a value below the window's top *equals* the sum of k rotation points on the circle. The hitting lemma lets us place $k - 1$ of those points at will; integrality snaps the total onto the exact target; and the irrationality of the slope makes the ladder's rung $\psi = \delta + e\alpha$ incommensurable-with-nothing — just strictly shorter than the acceptance strip, so no target escapes.

Example 14.2 (The Fibonacci walk, traced). Take $k = 3$, letters $\{1, 2\}$ (so $\delta = e = 1$, every s admissible), and the Fibonacci gap word: $q_n = \lfloor n\alpha + \beta \rfloor$ with $\alpha = (3 - \sqrt{5})/2 \approx 0.382$ and $\beta = \alpha$. Then $\zeta = \alpha/4 \approx 0.095$ and the ladder climbs by $\psi = 1 + \alpha \approx 1.382$ per step — comfortably below the strip height $3 - 2\zeta \approx 2.81$. To represent $x = 100$: $\theta(s) = (s + 3)\alpha - 100 + s$ climbs by ψ and first enters $[\zeta, 3 - \zeta]$ at $s^* = 72$, where $\theta(72) = 75\alpha - 28 \approx 0.647$. So $w = w(72) = 100 - 72 = 28$, and what is wanted is three distinct indices with $i_1 + i_2 + i_3 = 72$ and $q_{i_1} + q_{i_2} + q_{i_3} = 28$ — three rotation points summing to depth 0.647. One such triple is $i = (10, 26, 36)$, whose counting values are $4 + 10 + 14 = 28$ (check them against $q_n = \lfloor (n + 1)\alpha \rfloor$), so indeed $100 = 1 \cdot 72 + 1 \cdot 28 \in 3A - 3a_1$.

An honest caveat about this example. Step 2 guarantees such a triple only for budgets $s \geq S_0$, and S_0 is built from the hitting constant L of Lemma 13.10 — which, for these parameters, is enormous: the arc has length $\varepsilon_0 \approx 0.024$, Dirichlet's pigeonhole with $Q = 85$ yields a step $\theta \approx 0.0081$ and hence $L \approx 6800$, giving $S_0 \approx 34,000$. So $s^* = 72$ is far *below*




FIGURE 23. Border–period duality (Lemma 14.3) on $u = 0101010$. A prefix that is also a suffix forces the word to repeat with period “length minus border”. In Lemma 14.5 this is what converts two palindromic prefixes into a period — and arbitrarily long prefixes of the same period into global periodicity, which is the contradiction sought.

the threshold at which the theorem’s construction promises anything, and the triple above was found by direct search rather than produced by Step 2. That is exactly what one should expect: the proof’s constants are chosen for uniformity across all large x , not for sharpness at small ones. The example is here to show you the *ladder* — $\theta(s)$ climbing by ψ until it lands in the strip — with numbers small enough to check by hand, not to demonstrate Step 2 firing.

14.2. The boundary: k even, $d_2 = k$. One configuration escaped the width dichotomy: k even at the exact boundary $d_2 = k$, where the pairing argument came up one unit of width short. The rescue is a lemma that squeezes width out of *unbalancedness itself*, with a beautiful detour through palindromes.

First, a two-line classic:

Lemma 14.3 (Border–period duality). *Let u be a finite word of length τ with a border of length ℓ : its prefix of length ℓ equals its suffix of length ℓ . Then u has period $\tau - \ell$: $u_i = u_{i+\tau-\ell}$ for all $1 \leq i \leq \ell$.*

Proof. The border condition says $u_i = u_{(\tau-\ell)+i}$ for $1 \leq i \leq \ell$ — which is the period condition, re-read. $\square$

Example 14.4. $u = 0101010$, of length 7, has the border 01010 of length 5 (first five letters = last five letters). The lemma awards period $7 - 5 = 2$: indeed $u_i = u_{i+2}$ throughout. Long border, short period — the general rule of thumb.

Lemma 14.5 (Boundary width). *Let $k \geq 4$ be even. Suppose the binary word h uses both letters infinitely often, is not eventually periodic, and has the property that every tail of h is unbalanced. Then $\text{width}(s) \geq k - 1$ for all sufficiently large s .*

Proof. First, unbalancedness of every tail yields two-index width ≥ 2 infinitely often. Suppose instead $W_2(\sigma) \leq 1$ for all $\sigma \geq \sigma^*$. Take any two equal-length windows of the tail of h beyond position σ^* , say from i to i' and from j' to j ; then $\sigma := i + j = i' + j' \geq \sigma^*$, and exactly as in Lemma 12.4(b), $W_2(\sigma) \leq 1$ bounds the difference of their 1-counts by 1. So that tail would be balanced — contradiction. Hence $W_2(\sigma) \geq 2$ for infinitely many σ ; fix two such, $\sigma_0 < \sigma_1$, and set $\Delta := \sigma_1 - \sigma_0 \geq 1$.

Second, suppose the width bound fails at some large s : $\text{width}(s) \leq k - 2$. Play the splitting lemma twice, with a twist. Split one: $k/2 - 1$ two-slot groups at σ_0 plus one two-slot group at $\tau := s - (k/2 - 1)\sigma_0$; Lemma 12.3 forces $\text{width}(s) \geq (k - 2) + W_2(\tau)$, so $W_2(\tau) = 0$.

Split two ($k \geq 4$ makes room): $k/2 - 2$ groups at σ_0 , one at σ_1 , and one at $\tau - \Delta$; it forces $W_2(\tau - \Delta) = 0$ likewise.

Third, zero two-index width means palindrome. $W_2(\tau) = 0$ says $q_i + q_{\tau-i}$ is constant over $0 \leq i \leq \tau$; differencing consecutive i gives $h_{i+1} = h_{\tau-i}$ for $0 \leq i \leq \tau - 1$ — the prefix $h[1..\tau]$ reads the same forwards and backwards. So both $h[1..\tau]$ and $h[1..\tau - \Delta]$ are palindromes. A palindromic prefix of a palindrome is also its suffix (reverse the whole word: it is unchanged, and the prefix's mirror image — itself — appears at the end). So $h[1..\tau - \Delta]$ is a border of $h[1..\tau]$, and by Lemma 14.3, $h[1..\tau]$ has period Δ . The chain in one line:

$W_2 = 0 \rightarrow \text{palindrome} \rightarrow \text{border} \rightarrow \text{period } \Delta \rightarrow (\text{next paragraph}) \text{ global periodicity.}$

Finally, infinitely many failures would force global periodicity. Suppose $\text{width}(s) \leq k - 2$ for infinitely many s . The offset Δ is fixed, while $\tau(s) = s - (k/2 - 1)\sigma_0 \rightarrow \infty$ along the failing s : so we obtain prefixes $h[1..\tau]$ of period Δ with τ arbitrarily large. Now fix any position n and choose a failing s with $\tau(s) \geq n + \Delta$; the periodic prefix contains both positions n and $n + \Delta$, so its period relation gives $h_n = h_{n+\Delta}$. As n was arbitrary, the entire word has period Δ — in particular it is eventually periodic, contradicting the hypothesis. Hence $\text{width}(s) \geq k - 1$ for all large s . $\square$

Proposition 14.6 (Boundary routing). *Let k be even, $d_2 = k$, and let A have two-letter tail alphabet (the standing setup of §12). Then A is tail-covering.*

Proof. Route by the first case that applies; each closes the matter. If the gap word is eventually periodic: Lemma 10.1. Otherwise, if *some* tail of the gap word is balanced: pass to that tail (Lemma 9.2) and classify it (Theorem 13.8) — eventually periodic (Lemma 10.1 again) or Sturmian on a further tail (Lemma 14.1); note neither route used any width argument. Otherwise *every* tail is unbalanced and the word is not eventually periodic (and $k \geq 4$, being even and ≥ 3): Lemma 14.5 gives $\text{width}(s) \geq k - 1 = d_2 - 1$ for large s , and the sweep (Lemma 12.2) finishes. $\square$

14.3. The two-letter theorem.

Theorem 14.7 (Two letters). *Let $k \geq 3$ and $d_2 \leq k$. Every (d_1, d_2) -sequence whose tail alphabet has at most two letters is tail-covering.*

Proof. One letter: Lemma 9.3. Two letters: rescale to $\gcd = 1$ (Lemma 9.4). If k is even and $d_2 = k$: Proposition 14.6. Otherwise k is odd or $d_2 \leq k - 1$, and we consult the two-index widths: if $W_2(\sigma_0) \geq 2$ for some σ_0 , then Lemma 12.4(a) plus the sweep (Lemma 12.2) conclude; if $W_2(\sigma) \leq 1$ for all σ , the word is balanced (Lemma 12.4(b)), hence eventually periodic (Lemma 10.1) or Sturmian on a tail (Theorem 13.8, then Lemma 14.1). Every branch ends in tail-covering. $\square$

State of play. We now know: every walk with a tail alphabet of at most two letters is tail-covering. What remains: alphabets of three or more letters — and there the band machinery fails outright, so a new weapon is needed. Next: the slot lemma, and the exact finite statement it costs.

Exercise 14.8 (CORE). In Example 14.2, verify the claimed triple: generate the Fibonacci word ($0 \mapsto 01$, $1 \mapsto 0$, starting from 0), compute q_{10}, q_{26}, q_{36} , and check the two equations. Then find a different valid triple for the same $x = 100$.

Exercise 14.9 (CHALLENGE). Why can the Sturmian lemma not simply reuse the sweep? Locate the precise failure: show that for a Sturmian word, $\text{width}(s) \leq \text{some bound that falls}$

short of $d_2 - 1$ in the critical case $\delta' = d_2$. (Hint: Step 0 traps W_s in a real window of length k ; how many integers can it contain, and what does balancedness do to the endpoints?)

15. THREE OR MORE LETTERS: THE SLOT LEMMA

One case of the battle plan remains: tail alphabets with three or more letters. The band-and-sweep machinery does not survive the trip — with three letters, a single move between configurations changes the value by a *difference of letters*, which can exceed 1, so the reachable set W_s can have holes and Lemma 12.1 fails. Instead of repairing the band we change weapons entirely: a small combinatorial gadget — the *slot lemma* — that spends the k summands like a budget, and whose exact price is the finite theorem of Part IV.

15.1. Subset sums of multisets.

Definition 15.1. A *multiset* S drawn from a set G is a collection of elements of G with repetitions allowed ($|S|$ counts repetitions). Its *subset sums* are

$$\text{sums}(S) := \left\{ \sum_{x \in T} x : T \text{ a sub-multiset of } S \right\}$$

(including the empty sum 0). For a finite set G of positive integers with $M := \max G$, define

$$m(G) := \min \{ |S| : S \text{ a multiset drawn from } G \text{ whose subset sums contain } M \text{ consecutive integers} \},$$

with the convention $m(G) := \infty$ if no such multiset exists.

That convention is not idle. If $\gcd(G) = d > 1$ then *every* subset sum is a multiple of d , so no two consecutive integers — let alone M of them — can both be sums, and the minimum really is over an empty set. So $m(G)$ is finite exactly when $\gcd(G) = 1$ (for $M \geq 2$), which is why the coprimality hypothesis appears in every statement below rather than being an afterthought.

Example 15.2. $G = \{3, 4, 5\}$, $S = \{3, 3, 4, 5\}$: the subset sums are $\{0, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 15\}$, which contain the run $3, 4, \dots, 12$ — ten consecutive integers, amply covering $M = 5$. So $m(\{3, 4, 5\}) \leq 4$. (Part IV will show 4 is optimal here.)

Part IV is devoted to proving:

The bounded subset-sum covering lemma (SHARP). *Every finite set G of at least three positive integers with $\gcd(G) = 1$ and $\max(G) = M$ has $m(G) \leq M - 1$.*

For now we take it on credit and show exactly how the non-existence proof spends it.

15.2. The slot lemma.

Lemma 15.3 (Slot lemma). *Let A have tail alphabet G_∞ with $|G_\infty| \geq 3$ and $\gcd(G_\infty) = 1$ (after rescaling), and put $M := \max G_\infty \leq d_2$. If $k - 1 \geq m(G_\infty)$, then A is tail-covering.*

Proof. Pass to the tail (Lemma 9.2): all gaps lie in G_∞ , all letters recur infinitely often, and we work with prefix sums via Lemma 9.1. Let S be a multiset realizing $m(G_\infty) =: m \leq k - 1$: $|S| = m$, drawn from G_∞ , with

$$\text{sums}(S) \supseteq \{c, c + 1, \dots, c + M - 1\}$$

for some c . Enumerate S as $\delta_1, \dots, \delta_m$ (values in G_∞ , possibly repeating).

The fine slots. Because each letter of G_∞ occurs infinitely often in the gap word, we can choose m distinct positions $n_1 < n_2 < \dots < n_m$ (all beyond the tail start) with

$$g_{n_t} = \delta_t \quad (t = 1, \dots, m).$$

Summand t of our k -fold sum will be confined to two adjacent choices: index $n_t - 1$ (“off”), contributing the prefix sum p_{n_t-1} , or index n_t (“on”), contributing $p_{n_t} = p_{n_t-1} + \delta_t$. Toggling slot t from off to on adds exactly δ_t . Over all 2^m on/off patterns, the m slots jointly contribute

$$\left(\sum_{t=1}^m p_{n_t-1} \right) + \text{sums}(S) \supseteq \text{base}_0 + [c, c + M - 1], \quad \text{base}_0 := \sum_t p_{n_t-1} :$$

a full block of M consecutive achievable offsets — this is where the multiset’s covering run is cashed in.

The parked summands. We have used m of our k summands. Park the next $k - 1 - m$ at any fixed tail index (repetition of indices is legal in the sumset), folding their fixed contribution into the base: $\text{base} := \text{base}_0 + (\text{parked total})$. Now $k - 1$ summands realize every value of $\text{base} + [c, c + M - 1]$.

The coarse dial. The final summand is an index n of our choosing, contributing p_n . In the tail, consecutive prefix sums differ by a gap $g_{n+1} \in G_\infty \subseteq [1, M]$, so $n \mapsto p_n$ is strictly increasing with steps of size at most M , tending to infinity: a coarse dial that ascends forever but can jump by as much as M at a time.

The first-crossing argument. Fix a large target x and let

$$\phi(n) := x - \text{base} - p_n,$$

which decreases strictly in n , in steps of at most M , from a starting value $\phi(T) > c + M - 1$ (true once x is large, the base and T being fixed) down toward $-\infty$. Let n^* be the *first* index past T with $\phi(n^*) \leq c + M - 1$. By firstness, $\phi(n^* - 1) > c + M - 1$; and one step drops ϕ by at most M , so

$$\phi(n^*) > (c + M - 1) - M = c - 1, \quad \text{i.e.} \quad \phi(n^*) \geq c.$$

So $\phi(n^*)$ lands *inside* the achievable block $[c, c + M - 1]$: some on/off pattern realizes the offset $\phi(n^*)$, and then

$$x = \text{base} + \phi(n^*) + p_{n^*}$$

exhibits x as a sum of k prefix sums — m slots, $k - 1 - m$ parked, one dial. (The dial index may coincide with a slot or parked index; harmless, repetitions are allowed.) Hence $kA - ka_1$ contains every large integer, and A is tail-covering (with the factor g restored by Lemma 9.4 if rescaling was applied). $\square$

Watch the accounting at the binding corner (Figure 24). The dial can step by M , so the achievable block must be at least M long — a run of $M - 1$ consecutive offsets would let the dial leap clean over it. The block costs $m(G_\infty)$ summands, the dial costs one, and only k are available:

$$m(G_\infty) \leq k - 1.$$

When the Adversary’s walk uses the largest legal gap infinitely often and the parameters sit at the corner $k = d_2 = M$, this requirement becomes exactly

$$\underbrace{m(G_\infty)}_{\text{fine slots}} + \underbrace{1}_{\text{coarse dial}} \leq k = M \quad \Longleftrightarrow \quad m(G_\infty) \leq M - 1$$




FIGURE 24. The slot machine. The m toggles buy a block of M consecutive achievable offsets; the parked summands are dead weight folded into the base; and the coarse dial walks the remaining requirement $\phi(n)$ down in jumps of at most M . A descent by jumps $\leq M$ cannot leap over a block of width M : the first value $\leq c + M - 1$ is automatically $\geq c$.

— SHARP, with nothing to spare. The whole weight of the corner case rests on that finite statement, which is why Part IV exists. Away from the corner there is slack: $m(G_\infty) \leq M - 1 \leq d_2 - 1 \leq k - 1$ holds a fortiori.

For perspective — not needed below — a much cruder bound is easy to prove and shows the mechanism in miniature:

Lemma 15.4 (Base grid). *Let G have $|G| \geq 3$, $\gcd(G) = 1$, $a := \min G$, $M := \max G$, and suppose some $c \in G$ has $\gcd(a, c) = 1$. Then $m(G) \leq a + M - 1 \leq 2M - 3$.*

Proof. Take $a - 1$ copies of c and $p := \lceil ((a - 1)c + M - 1)/a \rceil$ copies of a . For any $n \in [(a - 1)c, pa]$: the residues $0, c, 2c, \dots, (a - 1)c$ modulo a are pairwise distinct (as $\gcd(a, c) = 1$), so a unique $i \in [0, a - 1]$ has $ic \equiv n \pmod{a}$; then $ic \leq (a - 1)c \leq n$, and $(n - ic)/a$ is an integer in $[0, p]$, so $n = ic + a \cdot \frac{n - ic}{a}$ is a subset sum. The interval has length $pa - (a - 1)c + 1 \geq M$ by the choice of p . The size: $c \leq M$ gives $p \leq \lceil ((a - 1)M + M - 1)/a \rceil = \lceil (aM - 1)/a \rceil = M$, so $(a - 1) + p \leq a + M - 1$; and $|G| \geq 3$ squeezes $a \leq M - 2$. $\square$

The gap between $2M - 3$ and SHARP's $M - 1$ is the gap between an evening's exercise and Part IV; note also the extra hypothesis, which genuinely can fail ($G = \{6, 10, 15\}$ has $\gcd 1$ but no element coprime to 6).

15.3. The non-existence half, assembled (modulo SHARP).

Proposition 15.5 (Goal of Part III). *Assume SHARP. Then every (d_1, d_2) -sequence with $d_2 \leq k$ and $k \geq 3$ is tail-covering.*

Proof. Pass to the tail; rescale so $\gcd(G_\infty) = 1$ (Lemmas 9.2, 9.4). If $|G_\infty| \leq 2$: Theorem 14.7. If $|G_\infty| \geq 3$: SHARP applies to G_∞ (three or more elements, $\gcd 1$) and gives $m(G_\infty) \leq \max(G_\infty) - 1 \leq d_2 - 1 \leq k - 1$, so the slot lemma (Lemma 15.3) concludes. $\square$

Combined with the certificate (Lemma 8.6), this proves the strong non-existence half of Theorem A — modulo SHARP. The debt is now fully explicit, and Part IV pays it.

State of play. We now know: both halves of the dichotomy, provided $m(G) \leq M - 1$ for every finite G with $|G| \geq 3$ and $\gcd(G) = 1$. What remains: exactly that one finite statement — no sequences, no walks, no lacunarity. Next: Part IV, which is about nothing but multisets of integers.

Exercise 15.6 (CORE). (a) Show $m(\{1, 2, 3\}) \leq 2$ and $m(\{2, 3, 4\}) \leq 3$, with explicit multisets. (b) Trace the slot lemma concretely: $k = 4$, $d_2 = 4$, $G_\infty = \{2, 3, 4\}$ (so $M = 4$), gap word cycling $2, 3, 4, 2, 3, 4, \dots$ — pick the multiset from (a), choose slot positions inside the first two periods, and find the dial index and slot pattern realizing $x = 57$ (with a_1 of your choosing).

Exercise 15.7 (CORE). The parked summands were placed at one fixed index. Would the proof survive parking them at index 0 (contributing $p_0 = 0$)? And why does the argument insist the block have length M rather than $M - 1$ — exhibit a concrete failure: a dial whose steps are all M leaping over a block of length $M - 1$.

Part IV. The subset-sum lemma

16. THE SUBSET-SUM LEMMA: STATEMENT, SHARPNESS, AND THE MAP

Part IV is a self-contained piece of finite mathematics. It never mentions walks, sumsets of sequences, or lacunarity; it proves exactly the statement that Part III left as a debt. Recall $m(G)$ from Definition 15.1: the smallest size of a multiset drawn from G whose subset sums contain $\max(G)$ consecutive integers.

Theorem 16.1 (Bounded subset-sum covering — “SHARP”). *Every finite set G of positive integers with $|G| \geq 3$, $\gcd(G) = 1$ and $\max(G) = M$ satisfies*

$$m(G) \leq M - 1.$$

Before proving it (§§17–24), let us be precise about what the bound claims, and check that neither hypothesis nor constant can be weakened.

16.1. First moves, and sharpness. If $1 \in G$, take $M - 1$ copies of 1: the subset sums are $0, 1, \dots, M - 1$ — exactly M consecutive integers at budget $M - 1$. So assume from now on $a := \min G \geq 2$.

The constant cannot be improved to $M - 2$. Take $G = \{3, 4, 5\}$, so $M = 5$. Example 15.2 showed $m(G) \leq 4$ via $\{3, 3, 4, 5\}$; we claim no multiset of size ≤ 3 works, so $m(\{3, 4, 5\}) = 4 = M - 1$ exactly. This is a finite check — there are 3 multisets of size one, 6 of size two, and 10 of size three — and the size-three table settles it (smaller multisets have fewer sums and do no better):

Multiset	Subset sums	Longest run
$\{3, 3, 3\}$	0, 3, 6, 9	1
$\{3, 3, 4\}$	0, 3, 4, 6, 7, 10	2
$\{3, 3, 5\}$	0, 3, 5, 6, 8, 11	2
$\{3, 4, 4\}$	0, 3, 4, 7, 8, 11	2
$\{3, 4, 5\}$	0, 3, 4, 5, 7, 8, 9, 12	3
$\{3, 5, 5\}$	0, 3, 5, 8, 10, 13	1
$\{4, 4, 4\}$	0, 4, 8, 12	1
$\{4, 4, 5\}$	0, 4, 5, 8, 9, 13	2
$\{4, 5, 5\}$	0, 4, 5, 9, 10, 14	2
$\{5, 5, 5\}$	0, 5, 10, 15	1

The best run of consecutive integers achievable with three elements is 3, short of the required $M = 5$. (And smaller multisets genuinely do no better: every multiset of size ≤ 3 sits inside one of size exactly 3, and adding elements only adds subset sums, so runs never shrink — the table’s maximum bounds them all.)

The hypothesis $|G| \geq 3$ cannot be dropped. Take $G = \{2, 3\}$: $\gcd = 1$, $M = 3$, and the target would be $m(G) \leq 2$. The five nonempty multisets of size ≤ 2 have subset sums

$$\{0, 2\}, \quad \{0, 3\}, \quad \{0, 2, 4\}, \quad \{0, 2, 3, 5\}, \quad \{0, 3, 6\},$$

and none contains 3 consecutive integers. So two-element sets genuinely need more room; three elements is where the uniform budget $M - 1$ becomes available.

Finally, a caution about what SHARP does *not* say: it is not a formula for $m(G)$, which is usually far smaller than $M - 1$ (for $G = \{1, 2, M\}$ it is tiny — see the exercise). It is a *uniform budget*, tuned to be exactly what the slot lemma's corner $k = d_2 = M$ can afford — recall §15: one dial summand plus at most $M - 1$ slot summands is all that $k = M$ pays for.

Exercise 16.2 (EXPLORATION). (a) Compute $m(\{1, 2, M\})$ for general $M \geq 3$. (b) A multiset of size s has at most 2^s subset sums, and M consecutive integers are M distinct sums; deduce the elementary lower bound $m(G) \geq \log_2 M$. Then find a three-element family of sets G (one for each large M) whose $m(G)$ is genuinely logarithmic-ish in M ... or convince yourself that with only three generators this is impossible, and explain which structural feature of G controls how far below $M - 1$ the truth sits.

16.2. The shape of the proof. The whole of Part IV runs on a single covering mechanism, previewed here and built in §17. To cover M consecutive integers with copies of elements of G :

- (1) *Cover the residues.* Using a bounded number of copies of the two largest generators, hit every residue class modulo a (the smallest generator) at least once — a *box* of representatives.
- (2) *Slide and stack.* Add $0, 1, 2, \dots, x$ copies of a : each representative slides up an arithmetic progression of step a , and the progressions interleave into a solid run once every class is represented and the run is long enough.

The cost accounting — how many copies of a suffice, and whether the total stays under $M - 1$ — is the *budget*, and the entire difficulty of Part IV is that near the corner the budget has no slack at all.


The proof organizes as follows.

- §17: the covering machinery — a two-generator interval lemma, the frame lemma (the slide-and-stack step made exact), the staircase (a refined multi-level version), and the λ -lift (a way to certify infinitely many triples with one finite check).
- §18: reduction of the general statement to *hard-core triples* $\{a, b, M\}$ with $\gcd(a, b) = 1$ and b close to M (precisely $1 \leq M - b \leq a - 2$).
- §§19–23: the hard core is carved by simple arithmetic invariants into six branches — D, P, L, E, T, B — four closed by explicit formulas valid across their whole range, two (T and B) closed by formulas *plus* a finite list of exceptional triples, each carrying a small computer-verified certificate. §26 explains exactly what was computed and how it is checked.
- §24: assembly, exhaustiveness of the six branches, and the sharpness proposition.

The six-way split is by two divisibility questions and two size questions, routed in order (evaluate top to bottom, take the first that applies):

Branch	Condition	Mechanism	Closure
D (<i>divisible max</i>)	$a \mid M$	b spans the residues alone	formula
P (<i>paired movers</i>)	$a \mid b + M$	(b, M) pairs, residue-neutral	formula
L (<i>linear spacing</i>)	$e = h$	staircase, extended form	formula
E (η - <i>box</i>)	$a \geq 12$ and $\mu \geq 12$	the η -box	formula
T (<i>thin μ</i>)	$\mu \leq 11$	staircase lines; 172 exceptions	formula + table
B (<i>bounded base</i>)	$a \leq 11, \mu \geq 12$	178 bases, lifted	table + lift

The three symbols in the conditions are just the three gaps in the triple, fixed here once so the table can be read now rather than after §17:



So “ μ small” means the generators are crowded near M , and “ a small” means the smallest generator is tiny. And the 12s and 11s are not deep: 12 is simply the cheapest constant at which Case E’s budget estimate closes — see the last line of its proof — and T, B split what E’s sizes exclude.) One more reading aid for the whole part: every case section below answers the same three questions — *how are the residues covered?* *how is the solid run produced?* *why is the budget $\leq M - 1$?* If you ever feel lost in a case, locate which of the three questions the current paragraph is answering.

17. MACHINERY: GRIDS, FRAMES, STAIRCASES, AND THE LIFT

Everything in Part IV covers runs of consecutive integers by the same two-phase mechanism sketched in §16.2. This section builds it once, with exact constants, so the case sections can spend their attention on budgets.

17.1. The two-generator grid. How much of the number line do two coprime generators cover, using boundedly many copies of each? The classic “postage stamp” fact — with stamps of two coprime denominations α, β and unlimited copies, every amount from $(\alpha - 1)(\beta - 1)$ on is payable — has the following bounded-copies version.

Lemma 17.1 (Two-generator interval). *Let $\alpha < \beta$ be coprime positive integers, and let $x \geq \beta - 1$ and $y \geq \alpha - 1$ be integers. Then, with $C := (\alpha - 1)(\beta - 1)$,*

$$\{i\alpha + j\beta : 0 \leq i \leq x, 0 \leq j \leq y\} \supseteq [C, \alpha x + \beta y - C].$$

Note the symmetry of the conclusion: the grid covers everything except a margin of C at each end of its full span $[0, \alpha x + \beta y]$.

Proof. Case $\alpha = 1$. Here $C = 0$. For $n \in [0, x + \beta y]$ put $j := \min(y, \lfloor n/\beta \rfloor)$. If $j = \lfloor n/\beta \rfloor$ then $n - j\beta$ is the remainder, in $[0, \beta - 1] \subseteq [0, x]$; if instead $j = y$ then $n - y\beta \leq x$ directly. Either way $i := n - j\beta \in [0, x]$.

Case $\alpha \geq 2$, base $y = \alpha - 1$. Let $n \in [C, \alpha x + \beta(\alpha - 1) - C]$. Since $\gcd(\alpha, \beta) = 1$, the multiples $0, \beta, 2\beta, \dots, (\alpha - 1)\beta$ hit all residues modulo α — the covering lemma (Lemma 7.6), in its first of many appearances in this part; choose the unique $j \in [0, \alpha - 1]$ with $j\beta \equiv n \pmod{\alpha}$. Then $n - j\beta$ is divisible by α and

$$n - j\beta \geq C - (\alpha - 1)\beta = -(\alpha - 1);$$




FIGURE 25. Lemma 17.1 for $(\alpha, \beta) = (3, 5)$, $x = 4$, $y = 2$ (so $C = 8$): the dots are the grid values $\{3i + 5j\}$, solid throughout the guaranteed middle $[8, 14]$, gappy only in the two margins of size C — exactly the shape the lemma promises. (Here the coverage happens to extend beyond the guarantee; the margins are worst cases.)

a multiple of α that is $\geq -(\alpha - 1)$ is ≥ 0 . And with $i := (n - j\beta)/\alpha$:

$$n - j\beta \leq (\alpha x + \beta(\alpha - 1) - C) - j\beta \leq \alpha x + (\alpha - 1)\beta - C = \alpha x + (\alpha - 1),$$

so $i \leq x + \frac{\alpha-1}{\alpha} < x + 1$, i.e. $i \leq x$.

Induction $y \rightarrow y + 1$ (for $y \geq \alpha - 1$). Only the new stretch $n \in (\alpha x + \beta y - C, \alpha x + \beta(y + 1) - C]$ needs attention. Subtract one β : certainly $n - \beta \leq \alpha x + \beta y - C$; and $n - \beta > \alpha x + \beta y - C - \beta \geq C$, because

$$\alpha x + \beta y \geq \alpha(\beta - 1) + \beta(\alpha - 1) = 2\alpha\beta - \alpha - \beta \geq 2C + \beta,$$

where the last inequality, written out, is

$$2\alpha\beta - \alpha - \beta \geq 2(\alpha - 1)(\beta - 1) + \beta = 2\alpha\beta - 2\alpha - \beta + 2 \iff \alpha \geq 2.$$

So $n - \beta$ lies in the already-covered interval: $n - \beta = i\alpha + j\beta$ with $i \leq x$, $j \leq y$, and $n = i\alpha + (j + 1)\beta$ with $j + 1 \leq y + 1$. $\square$

Exercise 17.2 (ROUTINE). Recompute Figure 25's dot set yourself and confirm the solid middle $[8, 14]$. Then: where exactly is coprimality used in the proof, and what goes wrong for $(\alpha, \beta) = (2, 4)$?

17.2. The frame lemma, and budgets. The second phase — slide and stack — in its general form:

Lemma 17.3 (Frame lemma). *Let ν, g_1, g_2 be positive integers and $L \geq 1$ a target run length. Suppose a box of size (Y, Z) covers all residues modulo ν : for every residue ρ there are $j_\rho \leq Y$, $k_\rho \leq Z$ with $j_\rho g_1 + k_\rho g_2 \equiv \rho \pmod{\nu}$. Let*

$$S := \max_{\rho} (j_{\rho} g_1 + k_{\rho} g_2), \quad x := \left\lceil \frac{L - 1 + S}{\nu} \right\rceil.$$

Then the multiset of Y copies of g_1 , Z copies of g_2 and x copies of ν — size $Y + Z + x$ — has subset sums covering the entire integer interval $[S, \nu x]$, which contains $[S, S + L - 1]$: a run of length L .

Proof. Let $n \in [S, \nu x]$ and $\rho := n \bmod \nu$. The representative $r := j_{\rho} g_1 + k_{\rho} g_2$ satisfies $r \leq S \leq n$ and $r \equiv n \pmod{\nu}$, so $t := (n - r)/\nu$ is a nonnegative integer, and $t \leq n/\nu \leq x$. Take j_{ρ} copies of g_1 , k_{ρ} of g_2 , and t of ν . Finally $\nu x \geq L - 1 + S$ by the choice of x , so the covered interval contains $[S, S + L - 1]$. $\square$




FIGURE 26. Slide and stack, on the triple $(5, 8, 9)$ with box representatives $0, 16, 17, 8, 9$ (values $8j + 9k$, one per residue class modulo 5, marked red). Stacking copies of $a = 5$ on each representative sends five arithmetic progressions climbing in step; from the height $S = 17$ of the *tallest* representative onward, every column is populated, and the shaded band $[S, S+M-1] = [17, 25]$ is covered solidly. This one picture is the mechanism of every construction in Part IV.


In every use below, $\nu = a = \min G$, $(g_1, g_2) = (b, M)$, and $L = M$ (Figure 26 shows a complete instance). The *budget* of a box (Y, Z) for the triple (a, b, M) is the resulting multiset size,

$$\mathcal{B} = \left\lceil \frac{M-1+S}{a} \right\rceil + Y + Z,$$

and a box *certifies* the triple when it covers all residues, $Y + Z \leq a - 1$, and $\mathcal{B} \leq M - 1$. (The middle condition looks unmotivated, and it is — for now. It does no work at all until §23, where it is exactly what makes the λ -lift free: a box with $Y + Z \leq a - 1$ costs at most a extra budget per lift, matching what the target gains. Carry it as bookkeeping until then.) The case analysis is, at bottom, a hunt for certifying boxes; the exceptional tables of cases T and B are nothing but lists of such boxes, checkable by the two-line computation in this proof.

Big idea. All of Part IV is one mechanism run six ways: *cover the residues modulo* a with a few copies of the big generators, then *slide and stack* with copies of a until the progressions knit into a solid run (Figure 26). The cost is the height of the tallest residue representative, divided by a , plus the box size — and the entire difficulty of the six branches is that near the corner this cost must squeeze under $M - 1$ with zero slack.

17.3. Standing notation for triples. *Notation reset.* Part IV reuses a few letters from Part III with new meanings, and from here on the old meanings never reappear: e and h now measure a *triple* (not a two-letter difference, not a binary word), and c will count box copies (not an offset). The picture to keep:



From §18 onward we work with triples $G = \{a, b, M\}$, $a < b < M$, $\gcd(a, b) = 1$. Fix, once and for all,

$$e := b - a, \quad h := M - b, \quad \mu := M - a = e + h, \quad g := \gcd(e, \mu),$$

$$e' := e/g, \quad \mu' := \mu/g, \quad C' := (e' - 1)(\mu' - 1).$$

(Note $\gcd(e, \mu) = \gcd(e, h)$, and e', μ' are coprime with $\mu' > e'$, so $\mu' \geq 2$ whenever $h \geq 1$.) One fact used repeatedly:

$$\gcd(a, g) = 1 :$$

a common prime of a and g would divide e and a , hence $b = a + e$, contradicting $\gcd(a, b) = 1$.

17.4. Levels and staircases. Consider a multiset of x copies of a , y of b , z of M , with the standing hypotheses

$$y \geq \mu' - 1, \quad z \geq e' - 1;$$

put $c := y + z$ and $V' := ye' + z\mu' - C'$. Group its subset sums by the *number of elements used*: choosing i copies of a , j of b , k of M with $i + j + k = t$ gives

$$ia + jb + kM = ta + je + k\mu,$$

so *level* t ($0 \leq t \leq x + y + z$) consists of $ta + \mathcal{O}_t$ with offset set

$$\mathcal{O}_t = \{je + k\mu : j \leq y, k \leq z, \max(0, t - x) \leq j + k \leq t\}.$$

Each level is a shifted copy of (part of) the same two-generator grid — a *frame* — and consecutive levels are offset by exactly a . Frames merging into solid runs is the entire game.

Before the general lemma, one instance in full. Take $(a, b, M) = (7, 9, 10)$, so $e = 2$, $h = 1$, $\mu = 3$, $g = 1$, $(e', \mu') = (2, 3)$ and $C' = (2 - 1)(3 - 1) = 2$. Choose $y = 2$ copies of b and $z = 2$ copies of M , so $c = y + z = 4$ and

$$V' = ye' + z\mu' - C' = 4 + 6 - 2 = 8,$$

and take $x = c + 1 = 5$ copies of a : the multiset $\{7^5, 9^2, 10^2\}$, of size $9 = M - 1$ exactly. Now watch the levels. The offsets available at any level $t \in [c, x] = [4, 5]$ are all of $\{2j + 3k : j \leq 2, k \leq 2\} = \{0, 2, 3, 4, 5, 6, 7, 8, 10\}$, which contains the guaranteed block $[C', V'] = [2, 8]$. So:

$$\text{level 4 : } 4 \cdot 7 + [2, 8] = [30, 36], \quad \text{level 5 : } 5 \cdot 7 + [2, 8] = [37, 43].$$

The two frames *abut exactly* — 36 and 37 are consecutive — and their union [30, 43] is a run of $14 \geq M = 10$ consecutive integers. Why did they abut rather than overlap or leave a hole? Because

$$V' - C' = 6 = a - 1$$

holds here with equality, which is precisely the merge condition of part (b) below. One unit less of frame and 37 would have been stranded. Keep this picture in mind: *the lemma's clauses are the several ways frames can be made to chain.*

Lemma 17.4 (Staircase). *In the setting above:*

- (a) *For every t with $c \leq t \leq x$: $\mathcal{O}_t \supseteq g \cdot [C', V']$ (meaning $\{gu : C' \leq u \leq V'\}$).*
- (b) (Two-frame merge; $g = 1$.) *If $x = c + 1$ and $V' - C' \geq a - 1$, the subset sums cover the interval $[ca + C', (c + 1)a + V']$, of length $a + (V' - C') + 1$.*
- (c) (Short merge; $g = 1$.) *If $x = c$, then $\mathcal{O}_{c+1} \supseteq [\max(C', 1), V']$, and if additionally $V' \geq a + \max(C', 1) - 1$ the subset sums cover $[ca + C', (c + 1)a + V']$.*
- (d) (Phase union; $g \geq 2$.) *If $x = c + g$, the subset sums cover $[(c + g - 1)a + gC', ca + gV']$ with no further hypothesis (the base form). If moreover $x \geq c + g$ and the frame-merge condition $V' - C' \geq a - 1$ holds, they cover $[(c + g - 1)a + gC', (x - g + 1)a + gV']$ (the extended form).*

Proof. (a) For $t \in [c, x]$, every box pair (j, k) ($j \leq y, k \leq z$) is admitted at level t : $j + k \leq c \leq t$ and $i = t - j - k$ lies in $[0, x]$. So $\mathcal{O}_t \supseteq \{je + k\mu\} = g\{je' + k\mu'\}$, and Lemma 17.1 applied to the coprime pair (e', μ') — its hypotheses are exactly $y \geq \mu' - 1, z \geq e' - 1$ — covers $[C', V']$.

(b) Levels c and $c + 1$ both lie in $[c, x]$, so each covers $ta + [C', V']$; two intervals offset by a merge into one exactly when $ca + V' \geq (c + 1)a + C' - 1$, i.e. $V' - C' \geq a - 1$.

(c) At level $c + 1 = x + 1$, the admission constraint is $j + k \geq t - x = 1$: only the pair $(0, 0)$ — offset 0 — is lost. Every offset in $[\max(C', 1), V']$ is nonzero, so all its grid representations survive. The merge then needs $ca + V' \geq (c + 1)a + \max(C', 1) - 1$.

(d) *Base form.* All levels $t \in [c, c + g]$ satisfy (a). Level t 's values $ta + g[C', V']$ lie in the residue class of ta modulo g , and since $\gcd(a, g) = 1$ the map $t \mapsto ta \bmod g$ is a bijection on any g consecutive levels — the covering lemma (Lemma 7.6), with modulus g and multiplier a : each residue class modulo g is served by exactly one $t_0 \in [c, c + g - 1]$. Given n in the stated interval, let t_0 be its server and $u := (n - t_0a)/g$; then $u \geq C'$ (since $n \geq (c + g - 1)a + gC' \geq t_0a + gC'$) and $u \leq V'$ (since $n \leq ca + gV' \leq t_0a + gV'$), so level t_0 realizes n .

Extended form. Within the class served by t_0 , the frames at levels $t_0, t_0 + g, t_0 + 2g, \dots$ (all $\leq x$) are arithmetic progressions of step g offset by ga from one another; consecutive frames leave no class-gap exactly when $(t_0 + g)a + gC' \leq t_0a + gV' + g$, i.e. when $V' - C' \geq a - 1$ — the frame-merge condition. Under it, the class is covered solidly from $t_0a + gC'$ up to $(t_0 + Jg)a + gV'$, where $J := \lfloor (x - t_0)/g \rfloor$. Over $t_0 \in [c, c + g - 1]$, all left endpoints are $\leq (c + g - 1)a + gC'$, and maximality of J ($t_0 + Jg > x - g$) makes all right endpoints $\geq (x - g + 1)a + gV'$; so every integer of the stated interval is covered within its own class. $\square$

Remark 17.5 (The merge condition is real). The frame-merge hypothesis in the extended form cannot be dropped. For $(a, b, M) = (5, 7, 9)$: $e = 2, \mu = 4, g = 2, (e', \mu') = (1, 2), C' = 0$. Take $y = 1, z = 0$, so $c = 1, V' = 1$, and $x = 3 = c + g$. The merge condition fails ($V' - C' = 1 < 4 = a - 1$) — and so does the would-be conclusion: the multiset $\{5, 5, 5, 7\}$ has subset sums $\{0, 5, 7, 10, 12, 15, 17, 22\}$, missing 11 inside the interval $[10, 12]$ that the un-repaired statement would promise. (An early draft of this staircase lemma omitted this hypothesis; the formalization caught it, and the counterexample is now permanently pinned




FIGURE 27. Left: the staircase when the frame-merge condition $V'-C' \geq a-1$ holds — consecutive levels' frames overlap and their union is one solid run. Right: the counterexample of Remark 17.5 — subset sums of $\{5, 5, 5, 7\}$, whose frames are too short, leaving the hole at 11.

in the Lean development as a guard. A cautionary tale told honestly: machine-checking is not decoration.) Figure 27 shows the merge working and failing, side by side.

17.5. The λ -lift. The final tool lets one finite check certify an infinite family. Triples (a, b, M) and $(a, b+a, M+a)$ have the same a , the same $h = M - b$, and the same residues $b, M \bmod a$; climbing this ladder is the λ -lift.

Lemma 17.6 (λ -lift). *Suppose the box (Y, Z) certifies the triple (a, b, M) : it covers all residues modulo a , has $Y + Z \leq a - 1$, and budget $\mathcal{B} \leq M - 1$. Then the same box certifies $(a, b + a, M + a)$.*

Proof. Modulo a , the generators are unchanged, so the same pairs (j_ρ, k_ρ) still cover all residues. Each representative's height grows by at most $(j_\rho + k_\rho)a \leq (Y + Z)a$, so the new maximum height satisfies $S' \leq S + (Y + Z)a$, and the new budget is

$$\left\lceil \frac{(M + a) - 1 + S'}{a} \right\rceil + Y + Z \leq \left\lceil \frac{M - 1 + S}{a} \right\rceil + 1 + (Y + Z) + Y + Z = \mathcal{B} + (Y + Z) + 1.$$

The budget grew by at most $Y + Z + 1 \leq a$; the target $M - 1$ grew by exactly a . A certificate whose margin was ≥ 0 keeps a margin ≥ 0 — forever, up the whole ladder. $\square$

Exercise 17.7 (CORE). Verify Remark 17.5's subset-sum list by hand. Then find the smallest x for which $\{5^{(x)}, 7\}$ (i.e. x copies of 5 and one 7) *does* contain a run of 5 consecutive integers among its subset sums — or prove none exists and explain which ingredient of the frame lemma is missing.

Exercise 17.8 (CORE). In the λ -lift, where exactly is $Y + Z \leq a - 1$ used, and what goes wrong without it? Construct a numerical instance: a box with $Y + Z = a$ whose budget margin degrades as it climbs.

18. REDUCTION TO THE HARD CORE

This section shrinks the problem from all alphabets to a two-parameter family of triples — a funnel with five stages, each stage one lemma:

$$\boxed{|G| \geq 3, \gcd = 1}_{\text{any } G} \xrightarrow{18.2} \boxed{\text{irreducible}} \xrightarrow{18.8} \boxed{\text{triples}} \xrightarrow{18.6} \boxed{\gcd(a, b) = 1} \xrightarrow{18.3, 18.4} \boxed{h \leq a - 2}_{\text{hard core}}$$

The tools are one growth lemma of Graham and four reduction lemmas, each an explicit construction.

18.1. Graham's growth lemma.


Lemma 18.1 (Graham growth). *Suppose $\text{sums}(S)$ contains a run of ℓ consecutive integers, and γ is a positive integer with $\gamma \leq \ell$. Then $\text{sums}(S \cup \{\gamma\})$ (one new copy of γ) contains a run of $\ell + \gamma$ consecutive integers.*

Proof. Let $[u, u + \ell - 1] \subseteq \text{sums}(S)$. The enlarged multiset's sums include both the old run and the old run shifted by γ :

$$[u, u + \ell - 1] \cup [u + \gamma, u + \gamma + \ell - 1].$$

Since $\gamma \leq \ell$, the shifted run starts no later than $u + \ell$, so the two runs overlap or abut, forming $[u, u + \ell + \gamma - 1]$. $\square$

A run, once at least as long as a generator, can be extended by that generator forever at a price of one copy per γ gained. This ratchet is what lets us prove SHARP for a *smaller* maximum and then climb.



18.2. Discarding redundant elements. Call an alphabet G *irreducible* if $|G| = 3$, or $|G| \geq 4$ and G is *minimal*: no proper subset of size ≥ 3 has gcd equal to 1.

Lemma 18.2 (Graham reduction). *If SHARP holds for all irreducible alphabets, it holds for all alphabets.*

Proof. Induct on $|G|$. Let $|G| \geq 4$ be non-minimal: some $e^\dagger \in G$ is *redundant*, meaning $G \setminus \{e^\dagger\}$ still has gcd 1 (and size ≥ 3). Two cases.

If $e^\dagger \neq M$: the smaller alphabet has the same maximum M , so its SHARP witness — a multiset of size $\leq M - 1$ covering M consecutive integers — is literally also a witness for G .

If $e^\dagger = M$: let $M' := \max(G \setminus \{M\})$. By induction, a multiset of size $\leq M' - 1$ drawn from $G \setminus \{M\}$ covers a run of M' consecutive integers. Now ratchet with $a := \min G$: the run has length $M' > a$, so each added copy of a extends it by a (Lemma 18.1, the hypothesis holding throughout since runs only grow); after $\lceil (M - M')/a \rceil \leq M - M'$ copies the run has length $\geq M$. Total size $\leq (M' - 1) + (M - M') = M - 1$. $\square$

18.3. Cheap pairs. Three lemmas now dispose of every triple whose two smaller elements are arithmetically convenient. Throughout, G has $\min G = a \geq 2$ and $\max G = M$.

Lemma 18.3 (Small coprime pair). *If some coprime pair $\alpha, \gamma \in G$ has $\alpha + \gamma \leq M - 1$, then $m(G) \leq M - 1$.*

Proof. Take $x := M - \alpha$ copies of α and $\alpha - 1$ copies of γ — budget exactly $(M - \alpha) + (\alpha - 1) = M - 1$. The representatives $0, \gamma, 2\gamma, \dots, (\alpha - 1)\gamma$ cover all residues modulo α (coprimality),

with top height $S = (\alpha - 1)\gamma$, so by the frame mechanism (Lemma 17.3 with $\nu = \alpha$) the sums cover $[S, \alpha x]$. Its length:

$$\alpha x - S + 1 = \alpha(M - \alpha) - (\alpha - 1)\gamma + 1 \geq \alpha(M - \alpha) - (\alpha - 1)(M - \alpha - 1) + 1 = M,$$

using $\gamma \leq M - \alpha - 1$; the expansion, line by line, is item (A1) of Appendix A — it lands exactly on M , with nothing to spare. $\square$

Lemma 18.4 (Pair at the boundary). *If some coprime pair $\alpha < \beta$ in G has $\alpha + \beta \in \{M, M + 1\}$, then $m(G) \leq M - 1$.*

Proof. For a parameter $t \geq 0$, take $x := \beta - 1$ copies of α and $y := \alpha - 1 + t$ copies of β . Lemma 17.1 applies and covers $[C, \alpha x + \beta y - C]$ with $C = (\alpha - 1)(\beta - 1)$; the length works out (item (A2) of Appendix A) to

$$\alpha x + \beta y - 2C + 1 = \alpha + \beta - 1 + t\beta,$$

at budget $x + y = \alpha + \beta - 2 + t$. If $\alpha + \beta = M + 1$: take $t = 0$; length M , budget $M - 1$. If $\alpha + \beta = M$: $t = 0$ gives length $M - 1$ — *one short* — so take $t = 1$: length $M - 1 + \beta \geq M$, budget $M - 1$ exactly. The budget absorbs precisely one extra copy; nothing is wasted, nothing is spare. $\square$

Example 18.5. $(\alpha, \beta, M) = (2, 3, 5)$, the $\alpha + \beta = M$ case. At $t = 0$ the multiset $\{2, 2, 3\}$ has sums $\{0, 2, 3, 4, 5, 7\}$: longest run $[2, 5]$, length $4 < 5$, insufficient exactly as the bookkeeping predicts. At $t = 1$, $\{2, 2, 3, 3\}$ has sums $\{0, 2, 3, 4, 5, 6, 7, 8, 10\} \supseteq [2, 8]$: seven consecutive, at budget $4 = M - 1$.

Lemma 18.6 (Non-coprime pair, $|G| = 3$). *Let $G = \{a, b, M\}$ with $d := \gcd(a, b) \geq 2$. (Then $\gcd(d, M) = 1$ automatically, since $\gcd(G) = 1$.) Then $m(G) \leq M - 1$.*

Proof. Write $a = d\alpha$, $b = d\beta$, $\gcd(\alpha, \beta) = 1$, $\alpha < \beta$. The multiset:

$$x := \beta - 1 \text{ copies of } a, \quad y := M - d - \beta + 1 \text{ copies of } b, \quad d - 1 \text{ copies of } M,$$

of size $(\beta - 1) + (M - d - \beta + 1) + (d - 1) = M - 1$ exactly.

The scaled grid. Sums using only a 's and b 's are d times sums from the (α, β) -grid. The hypotheses of Lemma 17.1: $x = \beta - 1$ exactly, and $y \geq \alpha - 1$ because $M > b = d\beta$ forces $M \geq d\beta + 1$, so $y \geq d\beta + 1 - d - \beta + 1 = (d - 1)(\beta - 1) + 1 \geq \beta > \alpha - 1$. With $C := (\alpha - 1)(\beta - 1)$, the covered set is the arithmetic progression $d \cdot [C, \alpha x + \beta y - C]$, of step d and

$$L := \alpha x + \beta y - 2C + 1 \text{ terms.}$$

The grid is long enough: $L \geq M + 1$. L grows with M (through y), so the worst case is the smallest legal $M = d\beta + 1$, where the claim rearranges (item (A3) of Appendix A, with $u := \beta(d - 1)$) to

$$(\beta - 1)(u - \alpha + 2) \geq u + 1,$$

which at $d = 2$ (i.e. $u = \beta$) reduces to $\beta(\beta - \alpha) + \alpha - 3 \geq 0$ — true, since $\beta(\beta - \alpha) \geq \beta \geq 2$ and $\alpha \geq 1$ — and only strengthens as d (hence u) grows, item (A3) again.

Bridging the d residues. The $d - 1$ copies of M contribute $0, M, 2M, \dots, (d - 1)M$, which — since $\gcd(d, M) = 1$ — run through *all* d residues modulo d . Adding them to the step- d progression fills every gap: for $n \in [dC + (d - 1)M, dC + (L - 1)d]$, pick $\kappa \in [0, d - 1]$ with $\kappa M \equiv n - dC \pmod{d}$; then $n - \kappa M$ is a step- d progression member (it is $\equiv dC$, and lies in the progression's range). The solid interval has length $(L - 1)d - (d - 1)M + 1 \geq M$, which is equivalent to $L \geq M + 1$ — just proved; the equivalence, including the small integrality step, is item (A4) of Appendix A. $\square$

18.4. Minimal alphabets of four or more elements.

Lemma 18.7 (Structure of minimal alphabets). *Let $G = \{g_1, \dots, g_m\}$ be minimal with $m \geq 4$, and set $d_i := \gcd(G \setminus \{g_i\})$. Then: each $d_i \geq 2$; the d_i are pairwise coprime; $d_j \mid g_i$ whenever $j \neq i$; consequently every $g_i \geq 30$ when $m = 4$ and ≥ 210 when $m \geq 5$; and every pair of elements of G has $\gcd \geq 2$.*

Proof. $d_i \geq 2$ is the definition of minimality (the subset $G \setminus \{g_i\}$, of size ≥ 3 , must not have $\gcd 1$). For $i \neq j$, any common divisor of d_i and d_j divides every element of $(G \setminus \{g_i\}) \cup (G \setminus \{g_j\}) = G$, hence divides $\gcd(G) = 1$. That $d_j \mid g_i$ for $j \neq i$ is immediate ($g_i \in G \setminus \{g_j\}$). So g_i is divisible by the product of the $m - 1$ pairwise coprime numbers $\{d_j\}_{j \neq i}$, each ≥ 2 ; the least possible such product is $2 \cdot 3 \cdot 5 = 30$ ($m = 4$) or $2 \cdot 3 \cdot 5 \cdot 7 = 210$ ($m \geq 5$). Finally, a coprime pair inside G would sit inside a proper size-3 subset of $\gcd 1$ (add any third element), contradicting minimality. $\square$

Lemma 18.8 (Minimal alphabets). *If G is minimal with $|G| \geq 4$, and SHARP holds for all alphabets of maximum $< M$, then $m(G) \leq M - 1$.*

Proof. (First pass: the two-sentence summary is — rescale the alphabet without its smallest element, apply SHARP inductively at the smaller maximum, then ratchet, scale and bridge as in Lemma 18.6; the rest is the count.)

Descend. Remove $a := \min G$. By minimality $d := \gcd(G \setminus \{a\}) \geq 2$; note $\gcd(a, d) = 1$ (a common factor would divide all of G). The rescaled alphabet $H := (G \setminus \{a\})/d$ has $\gcd(H) = 1$, $|H| \geq 3$, and maximum $M/d < M$ — so SHARP applies to it by hypothesis: a run of M/d consecutive integers within H 's subset sums, at cost $\leq M/d - 1$.

Extend. Let $a' := \min H$ and ratchet the run (Lemma 18.1, inside H) up to length

$$\ell := \left\lceil \frac{M - 1 + (d - 1)a}{d} \right\rceil + 1,$$

at an extra cost of at most $\lceil (\ell - M/d)/a' \rceil$ copies of a' .

Scale and bridge. Multiplying the run's multiset by d (i.e. using the original elements of $G \setminus \{a\}$) turns the run into a step- d progression of ℓ terms; the d residues are bridged by $0, a, 2a, \dots, (d - 1)a$ — all residues modulo d , as $\gcd(a, d) = 1$ — yielding a solid interval of length $(\ell - 1)d - (d - 1)a + 1 \geq M$ by the choice of ℓ . Total budget:

$$T \leq \left(\frac{M}{d} - 1 \right) + \left\lceil \frac{\ell - M/d}{a'} \right\rceil + (d - 1).$$

Count. From $a < \min(G \setminus \{a\}) = d a'$, i.e. $a \leq d a' - 1$:

$$\ell - \frac{M}{d} \leq \frac{(d - 1)a}{d} + 2 < (d - 1)a' + 2,$$

so the middle term is $\leq d + 1$ and $T \leq M/d + 2d - 1$. The requirement $T \leq M - 1$ rearranges to $M(1 - \frac{1}{d}) \geq 2d$, i.e. $M \geq 2d^2/(d - 1)$. For $d \geq 3$ this is at most $3d$, and $M \geq 3d$ because $G \setminus \{a\}$ contains three distinct multiples of d . For $d = 2$ it demands $M \geq 8$, — and every element of G is ≥ 30 (Lemma 18.7). $\square$

18.5. The hard core. Assemble the reductions. By Lemma 18.2 we may take G irreducible; Lemma 18.8 absorbs the minimal $|G| \geq 4$ case (its appeal to smaller maxima will be organized as a strong induction in §24); so $G = \{a, b, M\}$ is a triple, with $a \geq 2$ after the $1 \in G$ observation of §16. If $\gcd(a, b) \geq 2$: Lemma 18.6. So $\gcd(a, b) = 1$; and now Lemmas 18.3




FIGURE 28. Routing a hard-core triple. Conditions are read top to bottom and the first match wins, so by the time you are reading Case E you may assume $a \nmid M$, $a \nmid b + M$ and $e \neq h$ — which is exactly why $\eta \notin \{0, 1, a-1\}$ there. Each branch carries a representative triple; these are the running examples of §§19–23.

and 18.4 applied to the pair (a, b) dispose of every triple with $a + b \leq M + 1$. What survives is the *hard core*:

$$G = \{a, b, M\}, \quad \gcd(a, b) = 1, \quad h := M - b \in [1, a - 2],$$

(the condition $a + b \geq M + 2$ is $h \leq a - 2$), which forces $a \geq 3$. In words: the two large generators crowd within $a - 2$ of the top, and the small generator is at least 3. The standing notation of §17.3 is in force; add to it the mod- a data

$$\bar{e} := e \bmod a \in [1, a - 1], \quad \bar{\mu} := M \bmod a, \quad \eta := \bar{\mu} \cdot \bar{e}^{-1} \bmod a \quad (\text{when } \bar{\mu} \neq 0),$$

where $\bar{e}^{-1}$ is the inverse modulo a (existing since $\gcd(e, a) = \gcd(b - a, a) = \gcd(b, a) = 1$; in particular $\bar{e} \neq 0$). Three observations, each a one-liner worth doing:

$$\eta = 0 \iff a \mid M; \quad \eta = a - 1 \iff a \mid b + M; \quad \eta = 1 \text{ is impossible.}$$

(For the last: $\eta = 1$ would mean $\bar{\mu} = \bar{e}$, i.e. $a \mid M - e$; and the complete identity is

$$M - e = M - (b - a) = (M - b) + a = h + a,$$

so $a \mid M - e$ forces $a \mid h$ — impossible, since $1 \leq h \leq a - 2$.) The first two conditions are disjoint on the hard core (both would give $a \mid b$).

Every hard-core triple is now routed through the ordered six-branch tree of §16.2 — first matching condition wins, as Figure 28 shows — and §§19–23 close the branches one by one.

Big idea. Why is the hard core hard? Because b and M crowd within $a - 2$ of each other at the top: the two big generators are nearly interchangeable, their small difference h and offset e do all the arithmetic work, and whether the residues modulo a can be covered *cheaply enough* comes down to fine divisibility relations between a , e , and h . The six branches are six answers to one question: *who pays for the residue cover?* — b alone (D), residue-neutral pairs (P), an arithmetic progression’s regularity (L), the invariant η (E), or a finite certificate where every formula runs out (T, B).

Exercise 18.9 (CORE). First check each of these *is* a hard-core triple, then route it through the ordered tree (compute e, h, μ, g, η): $(5, 12, 15)$, $(3, 4, 5)$, $(5, 7, 9)$, $(13, 22, 27)$, $(5, 8, 9)$, $(5, 16, 17)$. You should hit all six branches, in the order D, P, L, E, T, B — several of

these are the coming sections' running examples. Then find a triple that is *not* hard-core and say which reduction lemma catches it (try $(12, 25, 36)$).

Exercise 18.10 (CORE). Show that the hard core is nonempty for every $a \geq 3$ (find its smallest member for each a), and that $(3, 4, 5)$ is the smallest hard-core triple of all.

19. CASES D AND L: TWO CLEAN FORMULAS

The first two branches are warm-ups with real content: each is closed by one explicit multiset whose budget check is a two-line inequality.

19.1. Case D: $a \mid M$. *Who lives here:* triples whose smallest generator divides the largest — e.g. $(5, 12, 15)$. *The trick:* when $a \mid M$, the copies of M are just extra copies of a in disguise (each worth $q = M/a$ of them), so b can be left to do the residue work alone. The three questions: *residues* — $a - 1$ copies of b sweep them all; *solid run* — copies of a and M pad at “half price”, since M is a multiple of a ; *budget* — the padding is so cheap that $2a + q - 3 \leq M - 1$ with room to spare unless $q = 2$.

Lemma 19.1 (Case D). *Let (a, b, M) be a hard-core triple with $a \mid M$; write $M = qa$ ($q \geq 2$). Put $x_{\text{eff}} := \lceil (M - 1 + (a - 1)b)/a \rceil$ and $z := \lceil (x_{\text{eff}} - (q - 1))/q \rceil$. Then the multiset of*

$$a - 1 \text{ copies of } b, \quad q - 1 \text{ copies of } a, \quad z \text{ copies of } M$$

has size $\leq M - 1$ and its subset sums contain M consecutive integers.

Proof. Coverage. The copies of b give the representatives $0, b, 2b, \dots, (a - 1)b$ — all residues modulo a , coprimality again — with top height $S = (a - 1)b$. For the sliding phase we need enough “effective copies of a ”: the pads $i a + \kappa M = a(i + q\kappa)$ with $i \leq q - 1$, $\kappa \leq z$ realize au for every $u \in [0, (q - 1) + qz]$ — the blocks $[q\kappa, q\kappa + q - 1]$ abut as κ steps — and $(q - 1) + qz \geq x_{\text{eff}}$ by the choice of z . Now run the frame argument (Lemma 17.3's proof verbatim, with the effective copies in place of literal ones): every $n \in [S, S + M - 1]$ is (residue representative) $+ au$ for some $u \leq \lceil (M - 1 + S)/a \rceil = x_{\text{eff}}$.

Budget. Since $b \leq M - 1$,

$$\frac{M - 1 + (a - 1)b}{a} \leq \frac{M - 1 + (a - 1)(M - 1)}{a} = M - 1,$$

so $x_{\text{eff}} \leq M - 1 = qa - 1$, whence

$$z \leq \left\lceil \frac{qa - 1 - (q - 1)}{q} \right\rceil = \left\lceil \frac{q(a - 1)}{q} \right\rceil = a - 1.$$

The total is then

$$(a - 1) + (q - 1) + z \leq 2a + q - 3 \leq qa - 1 = M - 1,$$

the last step being exactly $(a - 1)(q - 2) \geq 0$: true for all $q \geq 2$, $a \geq 3$, with equality only at $q = 2$ — the tightest members of the branch. $\square$

Example 19.2. $(a, b, M) = (5, 12, 15)$: $q = 3$, $S = 4 \cdot 12 = 48$, $x_{\text{eff}} = \lceil 62/5 \rceil = 13$, $z = \lceil 11/3 \rceil = 4$. Multiset: four 12s, two 5s, four 15s — size $10 \leq 14 = M - 1$. The four 12s hit residues $0, 2, 4, 1, 3$ modulo 5; the pads realize $5u$ for $u \in [0, 14] \ni 13$; and the sums cover $[48, 62] \supseteq$ fifteen consecutive integers.

19.2. Case L: equal spacings, $e = h$. *Who lives here:* triples in arithmetic progression — e.g. $(5, 7, 9)$. *The trick:* regular spacing makes the offsets $je + k\mu$ into consecutive multiples of g , so the staircase's frames are as long as they can be. The three questions: *residues* — handled inside the staircase, whose levels are offset by a ; *solid run* — the frame-merge condition holds, with equality at the extremes; *budget* — exactly $a + 2g - 1 = M - 1$, spent to the last unit.

When $e = h$ the triple is an arithmetic progression $\{a, a + g, a + 2g\}$ with $g = e = h$ (so $g \leq a - 2$ and $\gcd(a, g) = 1$), and the staircase machinery of §17 closes it directly.

Lemma 19.3 (Case L). *For $G = \{a, a + g, a + 2g\}$ with $\gcd(a, g) = 1$ and $1 \leq g \leq a - 2$, the multiset*

$$(x, y, z) = \left(\left\lfloor \frac{a-2}{2} \right\rfloor + 2g, \quad 1, \quad \left\lceil \frac{a-2}{2} \right\rceil \right)$$

(copies of a, b, M respectively) has budget exactly $a + 2g - 1 = M - 1$ and covers M consecutive integers.

Proof. Here $e = g, \mu = 2g$, so $(e', \mu') = (1, 2)$, $C' = 0$, and with $y = 1$: $V' = 1 + 2z, c = z + 1$. The budget is $x + 1 + z = (a - 2) + 2g + 1 = M - 1$ in all parity cases. Three sub-cases, one per applicable staircase clause.

a even. Then $z = \frac{a-2}{2}$ and $x - c = 2g - 1 \geq g$, so $x \geq c + g$; the frame-merge condition $V' - C' = 2z + 1 = a - 1 \geq a - 1$ holds *with equality*. The extended form (Lemma 17.4(d)) covers a solid interval of length $(x - c - 2g + 2)a + g(2z + 1) + 1 = a + g(a - 1) + 1 \geq a + 2g = M$ (i.e. $g(a - 3) \geq -1$: always).

a odd, $g \geq 2$. Then $z = \frac{a-1}{2}$, $x - c = 2g - 2 \geq g$, and $V' - C' = a \geq a - 1$ with room to spare. The extended form covers length $ga + 1 \geq a + 2g = M$ (i.e. $(g - 1)a \geq 2g - 1$: true for $g \geq 2, a \geq 3$).

a odd, $g = 1$. Then $x = c$ exactly, and the short merge (Lemma 17.4(c)) applies: level c covers $ca + [0, a]$ (as $V' = 2z + 1 = a$), level $c + 1$ covers $(c + 1)a + [1, a]$, and the merge condition $V' \geq a + \max(C', 1) - 1 = a$ holds with equality; the run $[ca, (c + 1)a + a]$ has length $2a + 1 \geq a + 2 = M$. $\square$

Note the pattern, which will repeat: the constructions are not delicate to *find* — they are the obvious staircases — but their budgets close against $M - 1$ with equality at the extremes. The margin throughout Part IV is zero at the corners; that is what “sharp” means concretely.

Exercise 19.4 (CORE). Run Case L on $(5, 7, 9)$ ($g = 2$; odd a): write out the multiset, the levels, and the covered interval, and reconcile with Remark 17.5 — why does the Case-L multiset succeed where $\{5, 5, 5, 7\}$ failed?

20. CASE P: PAIRS AS MOVERS

Who lives here: triples whose two big generators sum to a multiple of the small one — e.g. $(3, 4, 5)$, the smallest hard-core triple of all. *The trick:* modulo a the generators b and M cancel, so a (b, M) pair is a pure long-distance *mover*. The three questions: *residues* — unpaired copies of b (or of M) set the class, at most $\lfloor a/2 \rfloor$ of either; *solid run* — movers and copies of a slide within the class; *budget* — exactly $M - 1$, verified through four corner inequalities.

Now the line $a \mid b + M$: here $b + M = ra$ for an integer r , and $r \geq 3$ — because $a < b < M$ gives $b + M > 2a$. This branch has the prettiest mechanism in Part IV. Modulo a , the two big generators are *negatives of each other*: $M \equiv -b$. So one copy of b plus one copy of M is residue-neutral — a *mover*: it costs 2 elements and advances the total by exactly $b + M = ra$,




FIGURE 29. The mover of Case P. One b and one M together go nowhere on the residue clock ($b + M \equiv 0 \pmod{a}$) while leaping ra up the number line: a pure translation along the walk's own residue class, at a cost of two elements.

i.e. r grid units of a , without touching the residue. Unpaired copies of b (or of M) set the residue; movers and copies of a then slide the frame (Figure 29).

Lemma 20.1 (Case P). *Let (a, b, M) be a hard-core triple with $b + M = ra$, $r \geq 3$. Then the multiset*

$$(x, y, z) = \left(r - 1, \left\lceil \frac{M-r}{2} \right\rceil, \left\lfloor \frac{M-r}{2} \right\rfloor \right)$$

(copies of a, b, M), of budget exactly $(r - 1) + (M - r) = M - 1$, has subset sums containing M consecutive integers.

Proof. Coordinates. A selection of i copies of a , j of b , κ of M is worth

$$ia + jb + \kappa M = ia + wb + \kappa ra, \quad w := j - \kappa$$

(substituting $M = ra - b$): κ full movers, plus w unpaired copies of b if $w > 0$ (or $|w|$ unpaired copies of M if $w < 0$, by the symmetric substitution). Set

$$n_b := \left\lfloor \frac{a}{2} \right\rfloor, \quad n_M := \left\lfloor \frac{a-1}{2} \right\rfloor \quad (\text{so } n_b + n_M = a - 1)$$

— the most unpaired copies of b , respectively of M , we will ever spend (we avoid the letters p, q , busy elsewhere in this paper) — and declare the run to start at $\text{low} := \max(n_b b, n_M M)$. We must realize every $n \in [\text{low}, \text{low} + M - 1]$.

Three preliminary facts.

- (P1) $M \geq a + r - 1$. From $b < M$: $ra = b + M < 2M$, so $2M \geq ra + 1$; and $a, r \geq 3$ give $a(r - 2) \geq 3(r - 2) \geq 2r - 3$, whence $2M \geq ra + 1 = a(r - 2) + 2a + 1 \geq 2a + 2r - 2$.
- (P2) $M \leq a(r - 1) - 1$. From $b > a$: $M = ra - b < a(r - 1)$.
- (P3) $y \geq n_b$ and $z \geq n_M$. By (P1), $M - r \geq a - 1$, so $z = \lfloor (M - r)/2 \rfloor \geq \lfloor (a - 1)/2 \rfloor = n_M$ and $y = \lceil (M - r)/2 \rceil \geq \lceil (a - 1)/2 \rceil = n_b$.

The two branches. Fix a target n . Since $\gcd(a, b) = 1$ there is a unique $w \in [0, a - 1]$ with $wb \equiv n \pmod{a}$ (the covering lemma, Lemma 7.6).

Branch 1: $w \leq n_b$ — spend w unpaired bs. Then $wb \leq n_b b \leq \text{low} \leq n$, so $A := (n - wb)/a$ is a nonnegative integer. Split $A = i + r\kappa$ with $i := A \bmod r \leq r - 1 = x$ and $\kappa := \lfloor A/r \rfloor$: take i copies of a , κ movers, and the w unpaired bs — that is, i copies of a , $w + \kappa$ copies of b , κ copies of M , worth exactly n (check: $ia + (w + \kappa)b + \kappa M = ia + wb + \kappa ra = aA + wb = n$). Availability requires

$$\kappa \leq z \quad \text{and} \quad w + \kappa \leq y.$$

For the first, since $aA \leq n \leq \text{low} + M - 1$, it suffices that $\text{low} + M - 1 \leq a(rz + x)$ (then $A \leq rz + x$, so $\kappa \leq z$). For the second, note $a(A + rw) = (n - wb) + w(b + M) = n + wM \leq (\text{low} + M - 1) + n_b M$, so it suffices that

$$\text{low} + M - 1 + n_b M \leq a(ry + x). \quad (\text{C}_y)$$

Branch 2: $w > n_b$ — spend $v := a - w \in [1, n_M]$ unpaired M s. Since $M \equiv -b$, $vM \equiv -vb \equiv wb \equiv n \pmod{a}$, and $vM \leq n_M M \leq \text{low} \leq n$: so $A' := (n - vM)/a$ is a nonnegative integer, split as $A' = i' + r\kappa'$, giving i' copies of a , κ' copies of b , $v + \kappa'$ copies of M . Availability: $\kappa' \leq y$, which follows from $\text{low} + M - 1 \leq a(ry + x)$ (weaker than (C_y)); and $v + \kappa' \leq z$, for which — from $a(A' + rv) = n + vb \leq (\text{low} + M - 1) + n_M b$ — it suffices that

$$\text{low} + M - 1 + n_M b \leq a(rz + x). \quad (\text{C}_z)$$

(C_z also implies the interim condition of Branch 1.)

Four extremes. (First pass: skippable — the mechanism is complete; what follows audits the two conditions at their four corner cases.) Each of (C_y) , (C_z) splits in two according to which term achieves $\text{low} = \max(n_b b, n_M M)$; using $n_b + n_M = a - 1$ and $b + M = ra$, the four resulting inequalities are

$$\begin{aligned} (\text{E1}) \quad & (a - 1)b + M - 1 \leq a(rz + x) && [\text{C}_z, \text{low} = n_b b] \\ (\text{E2}) \quad & n_M ra + M - 1 \leq a(rz + x) && [\text{C}_z, \text{low} = n_M M] \\ (\text{F1}) \quad & n_b ra + M - 1 \leq a(ry + x) && [\text{C}_y, \text{low} = n_b b] \\ (\text{F2}) \quad & aM - 1 \leq a(ry + x) && [\text{C}_y, \text{low} = n_M M] \end{aligned}$$

(for (E1): $n_b b + n_M b = (a - 1)b$; for (F1): $n_b b + n_b M = n_b ra$; for (F2): $n_M M + n_b M + M = aM$). Verifying these four suffices for every w and every n , because w entered only through $w \leq n_b$ and $v \leq n_M$ — the extremes used.

(F2). It suffices that $M \leq ry + r - 1$. From $y \geq (M - r)/2$ and $r \geq 3$: $ry + r - 1 \geq \frac{3}{2}(M - r) + r - 1 = \frac{1}{2}(3M - r - 2) \geq M \iff M \geq r + 2$, and (P1) with $a \geq 3$ gives $M \geq a + r - 1 \geq r + 2$.

(F1). $n_b ra \leq ar y$ by (P3), so (F1) reduces to $M - 1 \leq ax = a(r - 1)$, which follows from (P2).

(E2). Identical: $n_M ra \leq ar z$ by (P3), and the remainder is (P2) again.

(E1) — the binding one. Substitute $b = ra - M$ and use $z \geq (M - r - 1)/2$; it suffices that

$$2[(a - 1)(ra - M) + M - 1] \leq ar(M - r - 1) + 2a(r - 1).$$

Expand both sides. On the left,

$$2[(a - 1)(ra - M) + M - 1] = 2ra^2 - 2ra - 2aM + 4M - 2;$$

on the right,

$$ar(M - r - 1) + 2a(r - 1) = arM - ar^2 - ar + 2ar - 2a = arM - ar^2 + ar - 2a.$$

Moving everything to the right and collecting the M terms, the inequality rearranges to

$$0 \leq M(ar + 2a - 4) - ar^2 - 2ra^2 + 3ar - 2a + 2.$$

The coefficient of M is positive, so we may insert (P1)'s bound $2M \geq ra + 1$; after multiplying through by 2, the right side is at least

$$(ra + 1)(ar + 2a - 4) - 2ar^2 - 4ra^2 + 6ar - 4a + 4 = a(ar^2 - 2ar - 2r^2 + 3r - 2)$$

(the intermediate expansion is item (A5) of Appendix A), so (E1) holds whenever

$$ar(r - 2) \geq 2r^2 - 3r + 2. \quad (*)$$

For $r \geq 4$: $a \geq 3$ gives $ar(r-2) \geq 3r^2 - 6r$, and $3r^2 - 6r \geq 2r^2 - 3r + 2 \iff r^2 - 3r - 2 \geq 0$, true from $r = 4$ on. For $r = 3$: (*) reads $3a \geq 11$, i.e. $a \geq 4$.

So (*) can fail only at $(a, r) = (3, 3)$ — where $b + M = 9$ with $3 < b < M$ forces $(a, b, M) = (3, 4, 5)$, and we check that one triple directly: $(x, y, z) = (2, 1, 1)$, and (E1) reads $2 \cdot 4 + 5 - 1 = 12 \leq 3(3 + 2) = 15$. True. (Concretely the multiset is $\{3, 3, 4, 5\}$ of Example 15.2, whose sums cover the run $[3, 12]$ — our very first example was the tightest triple on the P-line.) $\square$

Case P needs no certificate table and no computer: one closed-form multiset, four inequalities, one directly-checked triple. (The Lean formalization proves this case by a slightly different route — instantiating small r separately — a difference of route, not of content.)

Exercise 20.2 (CORE). Play out Case P on $(3, 4, 5)$ by hand: for each $n \in [\text{low}, \text{low} + 4]$, find the branch, the residue count w , and the selection realizing n . Where is the mover used?

Exercise 20.3 (ROUTINE). Derive the display in (E1) from (C_z) with $\text{low} = n_b b$: check $n_b b + n_M b = (a - 1)b$ and trace each substitution. (This is the kind of bookkeeping the formalization checks so humans don't have to — but doing it once by hand is how you learn what the machine is sparing you.)

21. CASE E: THE η -BOX

Who lives here: everything the divisibility branches missed, in the roomy regime where both a and μ are at least 12 — e.g. $(13, 22, 27)$. *The trick:* one invariant, η , measures how many steps of b it takes to manufacture one step of M modulo a ; a box of that shape covers the residues, and roominess pays for it. The three questions: *residues* — the η -box $(t - 1, Z)$, whose offsets sweep an interval of length $\geq a$; *solid run* — the frame lemma, as always; *budget* — one inequality, $M(a - \sigma) \geq a\sigma$, which both 12s exist to guarantee.

The remaining triples have $\eta \notin \{0, 1, a - 1\}$ — neither divisibility holds — and split by size. This section handles the roomy regime: both a and $\mu = M - a$ at least 12. Here a single box works, built directly from the invariant η , and the whole case is one budget computation. The case operates entirely modulo a : recall $b \equiv \bar{e}$ and $M \equiv \bar{\mu} = \eta \bar{e} \pmod{a}$.

Lemma 21.1 (Case E). *Let (a, b, M) be a hard-core triple with $a \nmid M$, $a \nmid b + M$, $a \geq 12$ and $\mu \geq 12$. Put*

$$t := \min(\eta, a - \eta) \in [2, a/2], \quad Z := \left\lceil \frac{a - t}{t} \right\rceil, \quad \sigma := t + Z.$$

Then the box $(Y, Z) = (t - 1, Z)$ covers all residues modulo a , and its budget is $\leq M - 1$.

Proof. First, the asserted range $t \in [2, a/2]$: $t \leq a/2$ is the definition of a minimum of η and $a - \eta$; and $t \geq 2$ because the earlier branches removed $\eta \in \{0, a - 1\}$ (cases D and P) while $\eta = 1$ is impossible on the hard core (§18.5) — so $2 \leq \eta \leq a - 2$, and both η and $a - \eta$ are at least 2.

Coverage. Work modulo a . If $t = \eta$: a box pair (j, k) contributes $jb + kM \equiv \bar{e}(j + \eta k) = \bar{e}(j + tk)$, and over $j \leq t - 1$, $k \leq Z$ the quantity $j + tk$ takes every value in $[0, t(Z + 1) - 1]$ — the t -blocks abut — which has length $t(Z + 1) \geq a$ by the choice of Z . If $t = a - \eta$: then $M \equiv -t\bar{e}$, the contribution is $\bar{e}(j - tk)$, and $j - tk$ covers the interval $[-tZ, t - 1]$, again of length $\geq a$. Either way $j + tk$ (or $j - tk$) sweeps a full residue system, and multiplying by the *unit* $\bar{e}$ permutes residues: all of $\mathbb{Z}_a$ is covered.




FIGURE 30. The η -box for $(a, b, M) = (13, 22, 27)$, where $\eta = 3$. The box of $Y = t - 1 = 2$ copies of b and $Z = 4$ copies of M realizes the quantities $j + tk$ (left), which sweep the whole run $0, 1, \dots, 14$ — shaded are the first 13, already a complete set of residues modulo 13. Multiplying by the unit $\bar{e}$ permutes the residue clock without collapsing it, so every class is covered (right). The t -blocks abutting as k steps is the whole mechanism: choosing $t = \min(\eta, a - \eta)$ is what keeps the box small enough to afford.

Budget. Every representative height is at most $S_0 := (t - 1)b + ZM$, and the box size is $K := t - 1 + Z = \sigma - 1$. It suffices to show

$$M - 1 + S_0 \leq a(M - 1 - K),$$

for then $\lceil (M - 1 + S)/a \rceil \leq M - 1 - K$ and the budget is $\leq M - 1$. Using $b \leq M - 1$: $M - 1 + S_0 \leq t(M - 1) + ZM \leq \sigma M - t$, while the right side is $a(M - \sigma)$; so it suffices that $\sigma M + a\sigma \leq aM + t$, and — discarding $t \geq 0$ — that

$$M(a - \sigma) \geq a\sigma.$$

Now bound σ : $Z \leq (a - 1)/t$ gives $\sigma \leq t + (a - 1)/t =: f(t)$, and f peaks at an endpoint of $[2, a/2]$ — no calculus needed: for $s < t$,

$$f(t) - f(s) = (t - s) \left(1 - \frac{a - 1}{ts} \right),$$

so f decreases as long as $ts < a - 1$ and increases once $ts > a - 1$; on any interval its maximum is therefore at one of the two ends. Evaluating, $\sigma \leq \max(f(2), f(a/2)) \leq \max(\frac{a+3}{2}, \frac{a}{2} + 2) = \frac{a}{2} + 2$. (In particular $\sigma < a$, so the quantity $a - \sigma$ we are about to divide by is positive.) Hence

$$\frac{a\sigma}{a - \sigma} \leq \frac{a(a + 4)}{a - 4} = a + 8 + \frac{32}{a - 4} \leq a + 12 \quad (a \geq 12),$$

and $M = a + \mu \geq a + 12$ delivers exactly this. Both hypotheses $a \geq 12$ and $\mu \geq 12$ are spent in this last line. $\square$

Example 21.2. $(a, b, M) = (13, 22, 27)$: $e = 9$, $h = 5$, $\mu = 14$, $\bar{e} = 9$, $\bar{\mu} = 1$, and $\eta = 9^{-1} \bmod 13 = 3$ (as $9 \cdot 3 = 27 \equiv 1$). So $t = 3$, $Z = \lceil 10/3 \rceil = 4$: box $(2, 4)$ — drawn in Figure 30. The offsets $\bar{e}(j + 3k)$ with $j \leq 2$, $k \leq 4$ have $j + 3k$ sweeping $[0, 14] \supseteq [0, 12]$. Corner height $S_0 = 2 \cdot 22 + 4 \cdot 27 = 152$; budget $\leq \lceil (26 + 152)/13 \rceil + 6 = 14 + 6 = 20 \leq 26 = M - 1$, with plenty to spare — Case E is the *comfortable* regime, which is exactly why it tolerates a one-size-fits-all box.

Exercise 21.3 (CORE). Why work with $t = \min(\eta, a - \eta)$ rather than η itself? Find a triple where η is close to a , and compare the budgets of the two candidate boxes. Then verify the claim $f(2) = 2 + (a - 1)/2 \leq (a + 3)/2$ used above.

22. CASE T: THIN μ , AND THE FIRST TABLE

Who lives here: triples whose generators crowd together at the top ($\mu \leq 11$) while a runs free — e.g. (5, 8, 9). *The trick:* finitely many shapes (e, h) , each a *line* parameterized by a alone, on which one staircase formula works for every a — and only the budget, growing more slowly than the target, can ever fail. The three questions: *residues* and *solid run* — the staircase, unconditionally, for every a ; *budget* — fails for exactly 172 triples, each of which carries a replacement certificate.

Now $\mu = e + h \leq 11$ (with $e \neq h$, neither divisibility, and a unbounded). Since $e \leq 10$ and $h \leq 10$, there are finitely many pairs (e, h) ; each pair defines a *line* of triples $(a, a+e, a+e+h)$ parameterized by a alone. On each line we run one explicit staircase; its *coverage* works for every a , and only its *budget* can fail — and only finitely often, which is where the certificates of Table A come in. Note $\mu' \geq 3$ throughout this case: if $\mu' = 2$ then $e' = 1$, giving $e = g$ and $\mu = 2g$, i.e. $h = \mu - e = g = e$ — but $e \neq h$ here.

Construction T.. The multiset will be a staircase, and its z will be a maximum of two branches with two distinct jobs: a *constant branch* maintaining the staircase's standing hypothesis $z \geq e' - 1$, and a *ceiling branch* buying enough run length. With that read ahead of time, take $y := \mu' - 1$ and:

- *Variant A* ($g = 1$; here $e' = e$, $\mu' = \mu$, and write $C = C'$, $V = V'$):

$$z := \max\left(e - 1, \left\lceil \frac{\max(a, \mu) - 1 + 2C - ye}{\mu} \right\rceil\right), \quad x := c + 1 \quad (c = y + z).$$

- *Variant $g \geq 2$:*

$$z := \max\left(e' - 1, \left\lceil \frac{a + \lceil(\mu - 1)/g\rceil + 2C' - ye'}{\mu'} \right\rceil\right), \quad x := c + g.$$

Exactly one variant applies to each line, according to its g .

Lemma 22.1 (Case T, coverage). *On every Case-T line, Construction T's subset sums contain M consecutive integers. Its budget is $2c + 1$ (variant A) or $2c + g$ (variant $g \geq 2$).*

Proof. Variant A. The ceiling branch gives $z\mu \geq \max(a, \mu) - 1 + 2C - ye$, so

$$V - C = ye + z\mu - 2C \geq \max(a, \mu) - 1 \geq a - 1:$$

the two-frame merge condition. With $x = c + 1$, Lemma 17.4(b) covers $[ca + C, (c + 1)a + V]$, of length $a + (V - C) + 1 \geq a + \max(a, \mu) \geq a + \mu = M$. Budget $x + y + z = 2c + 1$.

Variant $g \geq 2$. The ceiling branch gives $V' - C' \geq a + \lceil(\mu - 1)/g\rceil \geq a + (\mu - 1)/g$. With $x = c + g$, the *base form* of Lemma 17.4(d) — no merge hypothesis needed — covers a solid interval of length

$$g(V' - C') - (g - 1)a + 1 \geq ga + (\mu - 1) - (g - 1)a + 1 = a + \mu = M.$$

Budget $x + y + z = 2c + g$. □

So coverage is unconditional; everything now rides on the inequality

$$\text{budget} \leq M - 1 = a + \mu - 1.$$

The budget grows like $2z \approx 2a/\mu'$, i.e. with slope $2/\mu' \leq 2/3$ in a , while the target grows with slope 1 — so far out on every line the margin is positive and widening. Failures can only be *early* on a line.

Example 22.2 (The line $(e, h) = (2, 1)$). Here $\mu = 3$, $g = 1$, $(e', \mu') = (2, 3)$, $C = 2$: the triples $(a, a + 2, a + 3)$. Variant A takes $y = 2$ and $z = \max(1, \lceil (a - 1)/3 \rceil)$, budget $2c + 1$ against target $a + 2$:

a	z	c	budget $2c+1$	target $a+2$	margin
7	2	4	9	9	0 (tight)
13	4	6	13	15	2
19	6	8	17	21	4
25	8	10	21	27	6

The margin grows at rate $\approx 1/3$; on this lucky line it never goes negative. Other lines are less lucky: on $(e, h) = (5, 1)$, the triples $(4, 9, 10)$, $(7, 12, 13)$ and $(8, 13, 14)$ (among others) have construction budgets *exceeding* $M - 1$, and each needs a replacement certificate.

Proposition 22.3 (T-tail). *Across all Case-T lines, the applicable variant's budget exceeds $M - 1$ for precisely 172 triples, all with $a \leq 29$; each of these carries an explicit box certificate — a triple (x, Y, Z) with $Y + Z \leq a - 1$ and budget $\leq M - 1$, verified by the two-line computation of Lemma 17.3 — listed as Table A of the parent repository. For $a > 3000$ no failures occur; the range $a \leq 3000$ was checked exhaustively by computer.*


Proof, with the computation's role made explicit. Two separate claims.

The finite range $a \leq 3000$. For each of the finitely many lines and each $a \leq 3000$, the applicable variant's budget is a closed formula; a computer evaluated all of them exactly (integer arithmetic, no floating point) and found the budget to exceed $M - 1$ for exactly the 172 triples recorded in Table A, the largest at $a = 29$. This is a finite computation, redone independently three ways (§26); each failing triple's replacement certificate is itself re-verified mechanically. Nothing here asks you to trust a search heuristic: the claims are “this explicit formula, at these explicit inputs, exceeds that explicit bound” and “this explicit multiset covers that explicit run”, both checkable by hand for any single row — 172 rows is merely more than a human wants to do.

The tail $a > 3000$. (First pass: skippable.) Here no computer is needed. On a fixed line, once $a \geq 3000$ the ceiling branch of z dominates its constant branch (the ceiling's argument exceeds $(3000 - 1 - 100)/11 > 260$, while the constant branch is at most 9), so the budget equals its ceiling-branch value, which is bounded above by a linear function $\beta(a)$ of slope $2/\mu' \leq 2/3$; the target $\tau(a) = a + \mu - 1$ has slope 1. Therefore the margin $\tau - \beta$ is nondecreasing in a (rate $\geq 1/3$), and it suffices that it is positive at $a = 3000$ on every line — a final finite evaluation in exact rational arithmetic. Which line is worst? Reading β 's formula, its only large term is $2a/\mu'$, so the margin $\tau - \beta$ is smallest when μ' is smallest; among Case-T lines $\mu' \geq 3$, with equality on $(e, h) = (2, 1)$ and its relatives. So it is enough to evaluate there and note that every other line has $\mu' \geq 3$ as well, hence a margin at least as large at $a = 3000$ — no separate cross-line search is needed, and nothing beyond this paragraph is asked of the computer. Worth doing once by hand: on $(e, h) = (2, 1)$ we have $\mu' = 3$, $e' = 2$, $C' = 2$, $g = 1$, so

$$\beta(3000) = 2 \left[2 + \frac{3000 + 3 + 4 - 4}{3} + 1 \right] + 1 = 2 [3 + 1001] + 1 = 2009, \quad \tau(3000) = 3002,$$

margin $993 > 0$. (Note what is being compared: the target against the *linear over-bound* β , not against the exact budget — the exact budget at $a = 3000$ on this line is 2005, margin 997. Comparing against β costs four units of margin and buys the monotonicity that finishes the proof.) Positive there, nondecreasing after: positive forever. $\square$



The shape of Case T (schematic): the construction's budget grows with slope at most $2/3$, the target with slope 1. Early on the budget can poke above the target — those finitely many failures (dots) are Table A — but once the target line is ahead it stays ahead forever. The computer's entire role is confined to the left of the dashed line.

Big idea. The logical shape of Case T is the pattern of every computer-assisted step in this work: the *infinite* part is a slope comparison a student can check (budget grows at rate $\leq 2/3$, target at rate 1), and the *finite* part is large but mindless. The machine never supplies an idea — only stamina. §26 describes the checking apparatus in full.

Exercise 22.4 (CORE). Verify the $a = 7$ row of Example 22.2: build Construction T's multiset for $(7, 9, 10)$, list its levels, confirm coverage of 10 consecutive integers and budget exactly $9 = M - 1$ (tight, but not failing).

Exercise 22.5 (CORE). The triple $(5, 8, 9)$ (line $(e, h) = (3, 1)$, $g = 1$, $\mu = 4$) is one of the 172 failures. Compute Construction T's budget and watch it exceed $M - 1 = 8$; then verify its Table-A replacement certificate $(x, Y, Z) = (5, 2, 1)$: the multiset $\{5, 5, 5, 5, 5, 8, 8, 9\}$ has budget 8, its parts $8j + 9k$ hit all residues mod 5, and its sums cover $[17, 25]$ — nine consecutive integers.

23. CASE B: FINITE BASES, LIFTED FOREVER

Who lives here: the mirror image of Case T — a tiny smallest generator ($a \leq 11$) with the others far above it ($\mu \geq 12$) — e.g. $(5, 16, 17)$. *The trick:* bounded a means only finitely many *residue patterns*, and moving up within one pattern costs less than it earns, so one certificate per pattern suffices forever. The three questions: *residues* and *solid run* — a box certificate, verified once per pattern; *budget* — climbs by at most a per lift while the target climbs by exactly a , so a margin once won is never lost.

The last branch: $a \leq 11$ with $\mu \geq 12$. Now a is bounded but b (hence e , hence M) is not — each small a carries infinitely many triples. The finiteness that saves us is *modular*: group the triples into *classes*


$$[a; \bar{e}, h], \quad \bar{e} \in [1, a - 1], \quad h \in [1, a - 2],$$

by the residue of e modulo a and the exact value of h . Within a class, members differ only in $\lambda := \lfloor e/a \rfloor$, and stepping λ by one is exactly the map $(a, b, M) \mapsto (a, b + a, M + a)$ of the λ -lift (Lemma 17.6). So:

certify the smallest Case-B member of each class by an explicit box; the λ -lift then certifies every larger member, forever.

There are finitely many classes ($a \leq 11, \bar{e} < a, h \leq a - 2$), and the enumeration — machine-checked, §26 — yields exactly 178 classes whose members reach Case B. Their certified bases are Table B of Appendix D.

Example 23.1 (A class and its lift). The class $[5; 1, 1]$: triples with $a = 5, e \equiv 1 \pmod{5}, h = 1$. Its base in Table B is $(a, b, M) = (5, 16, 17)$ with box $(x, Y, Z) = (10, 1, 2)$: the parts $16j + 17k$ ($j \leq 1, k \leq 2$) have residues $\{0, 1, 2, 3, 4\}$ modulo 5 — check: $0, 16 \equiv 1, 17 \equiv 2, 33 \equiv 3, 34 \equiv 4, 50 \equiv 0$ — and the budget is $10 + 1 + 2 = 13 \leq 16 = M - 1$. The class members are $(5, 21, 22), (5, 26, 27), (5, 31, 32), \dots$, and Lemma 17.6 hands the *same* box up the whole ladder: the residues never change, each lift grows the budget by at most $Y + Z + 1 = 4$ while the target $M - 1$ grows by exactly 5. The margin widens forever; one finite check settles an infinite family.



The λ -ladder for the class $[5; 1, 1]$, with exact budgets computed at each rung: the target climbs by $a = 5$ per lift, the budget by at most $Y + Z + 1 = 4$ (here in fact 3). Lines that never re-cross: one finite check, an infinite family certified.

Where the tables live — and what a row looks like. A document that leans on $172 + 178$ certificates owes the reader the certificates. They are in the same repository as this paper, in two machine-readable files (`paper/data/table-A.csv`, `paper/data/table-B.csv`) and type-set in full in Appendix D; every row is verifiable by the two-line frame-lemma computation, and §26 describes the three independent systems that have verified all of them. To make the format concrete, here are the first rows of each, verbatim:

Table A row $(a, b, M; x, Y, Z)$	budget	vs. $M - 1$
$(4, 9, 10; 7, 1, 1)$	9	9 (tight)
$(5, 8, 9; 5, 2, 1)$	8	8 (tight; §22's exercise)
$(7, 12, 13; 7, 3, 1)$	11	12
$(8, 13, 14; 7, 3, 2)$	12	13
Table B class $[a; \bar{e}, h] \rightarrow$ base $(a, b, M; x, Y, Z)$	budget	vs. $M - 1$
$[4; 1, 1] \rightarrow (4, 17, 18; 13, 1, 1)$	15	17
$[5; 1, 1] \rightarrow (5, 16, 17; 10, 1, 2)$	13	16 (Example 23.1)
$[5; 1, 2] \rightarrow (5, 16, 18; 11, 2, 1)$	14	17

Two administrative points, each with a one-line proof idea.

Class membership of cases D and P is λ -invariant. Both conditions ($a \mid M; a \mid b + M$) depend only on residues modulo a , which the lift preserves. So a class is “a D-class”, “a P-class”, or neither, wholesale — no class straddles branches as λ climbs, except for one wrinkle:

The $e = h$ wrinkle. The condition $e = h$ (Case L) is *not* λ -invariant — e grows with λ while h stays fixed — but it can only hold at the class’s smallest member ($e \geq a > h$ once $\lambda \geq 1$). Exactly six classes (all on the diagonal $\bar{e} = h$, with $a \in \{9, 10, 11\}$) have a $\lambda = 0$ member that satisfies $\mu \geq 12$ but is diverted to Case L; those classes are simply *re-based* at their first genuine Case-B member, from which the lift proceeds as usual.

Why most bases sit at $\lambda \geq 1$. A class’s $\lambda = 0$ member often has $\mu = e + h < 12$ (small e), landing in Case T instead; the class enters Case B only once e has grown past $12 - h$. Example 23.1’s base has $\lambda = 2$ for exactly this reason. The enumeration accounts for every class and every member — completeness of the routing is part of what the harnesses check.

Exercise 23.2 (CORE). Verify the base certificate of the class $[5; 1, 1]$ end to end: residues, heights, budget, and the covered run of 17 consecutive integers. Then perform one lift by hand — write the certificate’s guarantee for $(5, 21, 22)$ and check the budget bound $13 + 4 \leq 21 - 1$... and note that the lemma’s bound is generous: recompute the exact budget at $(5, 21, 22)$ directly.

24. THE PROOF OF SHARP, AND ITS SHARPNESS

Proof of Theorem 16.1. By strong induction on M : proving the statement for maximum M , we may assume it for every alphabet with a smaller maximum. (Exactly one place consumes this: Lemma 18.8 passes to the rescaled alphabet H of maximum $M/d < M$. No hard-core branch invokes the hypothesis, and no branch re-enters Lemma 18.8 at a larger maximum, so the recursion terminates.)

If $1 \in G$: done (§16). By Lemma 18.2, reduce to irreducible G ; by Lemma 18.8, dispose of minimal $|G| \geq 4$; by Lemma 18.6, assume the triple $\{a, b, M\}$ has $\gcd(a, b) = 1$; by Lemmas 18.3–18.4, assume $a + b \geq M + 2$: the hard core. Route:

if $a \mid M$: D $\rightarrow$ else if $a \mid b + M$: P $\rightarrow$ else if $e = h$: L $\rightarrow$ else if $a, \mu \geq 12$: E $\rightarrow$ else if $\mu \leq 11$: T $\rightarrow$ else B.

Each branch delivers a multiset of size $\leq M - 1$ covering M consecutive integers: Lemmas 19.1, 20.1, 19.3, 21.1, 22.1 with Proposition 22.3 (Table A), and Table B with Lemma 17.6.

Exhaustiveness and consistency of the routing. D and P are disjoint on the hard core (both would force $a \mid b$). A triple with $e = h$ and $a \mid b + M$ would have $b + M = 2a + 3e$, so $a \mid 3e$; with $\gcd(a, e) = 1$ this forces $a \mid 3$, i.e. $a = 3$, $e = h = 1$: the single triple $(3, 4, 5)$, which the ordered routing sends to P (and which L would also handle — no conflict, the router just picks first). Every triple surviving to the last three branches has $\eta \notin \{0, 1, a - 1\}$, and the conditions “ $a \geq 12 \wedge \mu \geq 12$ ”, “ $\mu \leq 11$ ”, “ $a \leq 11 \wedge \mu \geq 12$ ” partition the remaining (a, μ) -space. Every hard-core triple lands in exactly one branch. $\square$

The debt of Part III is paid: SHARP holds, so Proposition 15.5 holds unconditionally, and with it the non-existence half. It remains only to record sharpness — the checks were exhibited in §16:

Proposition 24.1 (Non-improvability). $m(\{3, 4, 5\}) = 4 = M - 1$: the constant $M - 1$ cannot be improved to $M - 2$. And for $G = \{2, 3\}$ ($\gcd 1$, $M = 3$), no multiset of size ≤ 2 covers 3 consecutive integers: the hypothesis $|G| \geq 3$ cannot be dropped. $\square$

Exercise 24.2 (CORE). The routing sends $(3, 4, 5)$ to P, yet L’s construction also covers it. Compute both multisets and compare. Then find a triple where the D-condition and the E-size-conditions both hold, and explain why the ordered routing makes the overlap harmless.

Part V. Assembly, and what the computer checked

25. ASSEMBLY: PROOF OF THE DICHOTOMY

All the pieces exist; this section only has to bolt them together.

Proof of Theorem A. Part (1), existence. This is Theorem 6.1 (proved there for all $k \geq 2$, hence for the problem's $k \geq 3$): if $d_2 \geq k + 1$, the announcement $r = 192 d_2$ wins against every lacunary B of that ratio.

Part (2), strong non-existence. Let $d_2 \leq k$ (with $k \geq 3$) and let $r_1, r_2, \dots$ be any prescribed sequence of ratio lower bounds. Build the certificate B of Lemma 8.6 for the sequence $R(i) := r_i$; it satisfies $b_{i+1} \geq r_i b_i$ for every i . Let A be any (d_1, d_2) -sequence. By Proposition 15.5 — now unconditional, since Part IV proved SHARP (Theorem 16.1) — A is tail-covering. Retracing which case fired, for the record:

The walk's tail alphabet (rescaled)	Case closed by
one letter	Lemma 9.3
two letters, gap word eventually periodic	Lemma 10.1
two letters, k even and $d_2 = k$	Proposition 14.6
two letters, some $W_2 \geq 2$ (k odd or $d_2 \leq k-1$)	Lemmas 12.4(a), 12.2
two letters, balanced	Theorem 13.8, then Lemma 10.1 or 14.1
three or more letters	Lemma 15.3 + SHARP

By Lemma 8.6(2), $(kA) \cap B \neq \emptyset$. As A was arbitrary, the single B defeats every admissible walk; as (r_i) was arbitrary, this is the strong form.

The dichotomy. For any $k \geq 3$ and $1 \leq d_1 < d_2$, exactly one of $d_2 \geq k + 1$ and $d_2 \leq k$ holds; part (1) gives a winning strategy in the first case, part (2) denies one in the second (a winning r would in particular defeat the certificate built for the constant sequence $r_i \equiv r$). So the answer to Erdős problem 1112 is: *yes precisely when $d_2 \geq k + 1$.* $\square$

Every promise of the honest ledger (§2.4) has now been redeemed: gap (1) in Part II, gaps (2) and (3) in Part III, gap (4) in Part IV. The remaining section is not part of the proof; it is about how the proof is *checked*.

Exercise 25.1 (CHALLENGE — grand tour). Fix $k = 3$, $(d_1, d_2) = (1, 3)$ — the losing parameter set of §1.10 — and the ratio sequence $r_i = 10^i$. Follow the entire non-existence machine for the specific walk with gap word 1, 2, 1, 1, 2, 1, 2, $\dots$ (the Fibonacci word's letters +1): which case of the table fires, what modulus results, and where in the certificate B 's construction does the collision get guaranteed? (Nothing new to prove — the point is to watch every part of Parts III–IV touch one concrete example.)

26. WHAT THE COMPUTER CHECKED, AND HOW YOU CAN TOO

This proof is computer-assisted, and this paper has flagged every place the assistance occurs: the 172 Table-A certificates and the $a \leq 3000$ scan (Case T), and the 178 Table-B base certificates with the class enumeration (Case B) — those four items are the load-bearing finite layer. (The parent paper's harnesses additionally run corroborating enumerations that this document's proofs never lean on, such as listing all minimal alphabets with $M \leq 130$; corroboration is welcome, but the ledger above is complete without it.) This closing section explains what, exactly, was computed, why you need not take anyone's word for it — and what such checking can and cannot promise. Nothing here is needed to *understand* the proof; it is needed to *trust* it, which is a different and equally honorable activity.

26.1. Certificates, not searches. Every computed object in this work is a *certificate*: a small tuple whose correctness is a finite, mindless computation. A Table-A or Table-B row $(a, b, M; x, Y, Z)$ asserts two things — the multiset of x copies of a , Y of b , Z of M has size $\leq M - 1$, and its subset sums contain M consecutive integers — and *you* can check any row with pencil and paper in minutes (Exercises in §§22–23 did exactly this). How the rows were *found* is irrelevant to their truth: a certificate severs the proof from the search that produced it. The same holds for the scans: “this formula exceeds that bound at exactly these 172 inputs” is a statement any programmer can re-derive in an afternoon, in any language.

To be precise about packaging: the complete human-readable proof record is *this paper alone*. All 350 rows are printed in Appendix D, generated from the same canonical data that the Lean kernel independently decides, and exported as CSV with recorded digests for readers who would rather compute than read. Nothing load-bearing is held elsewhere.

26.2. Three independent checks. The repository verifies the finite layer three separate ways, sharing no code:

- (1) *Two Python harnesses.* One reconstructs every designated witness across all 83,251 hard-core triples with $M \leq 120$ and verifies each in exact integer arithmetic; the other is a from-scratch re-implementation working only from the lemma statements. Both fail closed — any invalid row, missing class, or deviant count aborts with an error — and both report identical branch tallies (2392 D, 2605 P, 1420 L, 71,421 E, $3254 + 172$ T, 1987 B) — tallies which sum to exactly 83,251, the number of triples: the books visibly close. Agreement between independent implementations is strong evidence against ordinary coding slips — two people are unlikely to mistype the same thing — but it is evidence, not proof, and it says nothing at all about a shared misunderstanding of the mathematical design, since both programs were written from the same case analysis. That is what the proofs are for.
- (2) *The table generator.* The typeset tables in the certificate supplement are *generated* from one canonical data file by a script that re-verifies every row (budget and coverage, by exact bitmask arithmetic) on every build, so no printed row can silently diverge from the data that is checked.
- (3) *The Lean kernel.* The entire proof — not just the finite layer — is formalized in Lean 4 and machine-checked; the certificate tables are transcribed into the formal development and re-decided by Lean’s kernel.

26.3. What a Lean formalization is. Lean is a *proof assistant*: a programming language in which definitions, theorem statements and proofs are all written formally, and a small program — the *kernel* — checks that every proof step follows from the rules of logic and the axioms. The human (or, in this work, largely the AI systems noted below) supplies the proof; the kernel is the referee that cannot be sweet-talked. Three facts summarize the state of this development:


- It is *complete*: all three headline theorems — the dichotomy, the explicit ratio $192 d_2$, and the strong non-existence form — are proved, over frozen definitions, from start to finish.
- It is *gap-free*: no `sorry` (Lean’s “I owe you a proof” placeholder) anywhere, and the axiom audit reports only Lean’s three standard foundational axioms.
- Its finite checks run *in the kernel*: no appeal to compiled external code (no `native_decide`), so even the $172 + 178$ certificate rows and the $a \leq 3000$ scan are validated by the same small trusted checker as every logical step.

Anyone can reproduce this: clone the repository, run `lake build`, and watch it compile with no complaints (the repository README gives the exact commands, and links a public log of a complete run).

26.4. What machine-checking does not promise. Honesty requires the other half. A formal verification establishes: *these formal statements follow from these axioms*. It does not establish:

- *That the formal statement means what the English problem says.* If the Lean definition of “lacunary” or “gap window” subtly mismatched the problem, the kernel would happily verify the wrong theorem. This correspondence must be audited by humans reading both — which is why §26.5 walks through each modelling choice (for instance: over the natural numbers, careless subtraction truncates at zero, and a naively phrased gap condition would silently admit decreasing walks; the formalization phrases the window additively to dodge exactly this). The definitions are few and short by design, to make that audit as small as possible.
- *Novelty or attribution.* Whether SHARP is genuinely new is a question about the literature, addressed in §1.9 by a documented search — no computer can settle it.
- *That nothing outside the trusted base is wrong.* The kernel, the pinned compiler and library versions — these are trusted. The trust base is small and standard, but it is not zero; no honest account says otherwise.

The layers of trust, drawn as a stack — each layer checks the one above and must itself be trusted or audited:



One story worth retelling: an early draft of the staircase lemma omitted the frame-merge hypothesis, and the *formalization process* caught it — the counterexample (5, 7, 9) of Remark 17.5 is now permanently enshrined in the development as a guard. The moral runs both ways: human-written mathematics slips, and machine-checking catches real slips — and what the machine certifies is exactly the statement it was given, neither more nor less.

26.5. The formal statement, and why it is faithful. A formal proof certifies only the statement it was given, so the statement deserves to be read. The definitions live in a frozen file (`lean/Erdos1112.lean`) that the development imports but never edits, and the three theorems are stated once each, directly over those definitions:

```
def IsLacunaryWith (r : ℕ) (b : ℕ → ℕ) : Prop :=
  0 < b 0 ∧ StrictMono b ∧ ∀ i, r * b i ≤ b (i + 1)

def HasGapsIn (d₁ d₂ : ℕ) (a : ℕ → ℕ) : Prop :=
  0 < a 0 ∧ ∀ i, a i + d₁ ≤ a (i + 1) ∧ a (i + 1) ≤ a i + d₂

def kFoldSumset (k : ℕ) (a : ℕ → ℕ) : Set ℕ :=
  { n | ∃ f : Fin k → ℕ, n = ∑ j, a (f j) }
```

```

def RatioWorks (k d1 d2 r : ℕ) : Prop :=
  ∀ b : ℕ → ℕ, IsLacunaryWith r b →
    ∃ a : ℕ → ℕ, HasGapsIn d1 d2 a ∧
      Disjoint (kFoldSumset k a) (Set.range b)

def Question (k d1 d2 : ℕ) : Prop := ∃ r : ℕ, RatioWorks k d1 d2 r

```

Six modelling choices carry the weight, and each is a place where a formalization could quietly prove the wrong theorem.

- (1) **Sets as increasing enumerations.** A and B are infinite sets of positive integers, encoded by the functions that list them in order. `StrictMono` is what makes such a function enumerate a *set* rather than a multiset; it is imposed on B directly and *derived* for A from the gap condition when $d_1 \geq 1$.
- (2) **The gap window, phrased additively.** Over the natural numbers, subtraction is truncated at zero — so a subtractive phrasing would be a trap. An isolated upper bound “ $a_{i+1} - a_i \leq d_2$ ” would read $0 \leq d_2$ on a *decreasing* step, which is true, and would silently admit walks that go backwards. Writing the window as $a_i + d_1 \leq a_{i+1} \wedge a_{i+1} \leq a_i + d_2$ avoids this entirely. (The lower bound is not the trap: an underflowed gap gives $d_1 \leq 0$, false, since $d_1 \geq 1$.)
- (3) **The sumset, with repetition.** `kFoldSumset` ranges over an arbitrary index function f with *no* injectivity requirement, so summands may repeat — the intended kA of Definition 1.6. Requiring f injective would model the distinct-summand sumset, a strictly smaller set and a different problem.
- (4) **Quantifier order.** `Question` is $\exists r \forall B \exists A$, with A chosen after B — the game of §1.6, not the easier $\exists A \forall B$.
- (5) **The ratio over $\mathbb{N}$ rather than $\mathbb{Z}$.** The problem says “an integer r ”. Passing to $\mathbb{N}$ loses no existence content, and this is now proved rather than argued: the frozen file also carries the integer form `QuestionInt` together with the machine-checked equivalence `question_iff_questionInt`, and the headline theorem in integer form `erdos_1112_int`. So the dichotomy is verified for the problem’s literal phrasing.
- (6) **A genuine, non-vacuous biconditional.** Both directions are separately witnessed, and both sides are inhabited: none of the predicates is silently trivial. `IsLacunaryWith r` is satisfiable for every r (take $b_i = (r + 2)^i$), `HasGapsIn` is satisfiable (take $a_i = 1 + d_1 i$), and `kFoldSumset k a` is nonempty while `Set.range b` is infinite — so disjointness is a real constraint, not a vacuous one.

Scope. Every headline theorem carries the problem’s own hypothesis $k \geq 3$. The existence half is proved here in prose for all $k \geq 2$ (§3); the formalization deliberately carries $k \geq 3$ throughout, so the $k = 2$ existence statement stands as an unformalized supplementary observation. Nothing in the mathematical argument of this paper is delegated to the formalization.

Trust base. At the commit recorded in §29, built with the pinned toolchain `leanprover/lean4:v4.27.0` and the Mathlib revision fixed by `lake-manifest.json`, `#print axioms` reports for each of the three headline theorems exactly

[propext, Classical.choice, Quot.sound]

— Lean’s three standard foundational axioms, and nothing else. In particular there is no `sorryAx` (no gaps) and no `Lean.ofReduceBool` (the axiom `native_decide` would introduce). Every finite check, including the 350 certificate rows and the $a \leq 3000$ scan, is therefore

discharged by the trusted kernel rather than by compiled code. The audit is itself the gap check: a `sorry` anywhere would surface in that list.

26.6. What machine-checking still does not settle. Two honest caveats. “Depends only on three axioms” is a statement about the *logical* axiom report; the operative trusted base is larger, and includes the Lean kernel and elaborator, the Mathlib source at the pinned revision, and the dependency snapshot. And an axiom audit cannot detect a mismatch between a numbered result in this paper and the declaration that purports to formalize it. That correspondence is asserted in §E and must be checked by reading both — which is one more reason for a paper written to be read.

27. PROVENANCE AND AUTHORSHIP

This work was produced with substantial assistance from automated systems, and we record the division of labour.

The mathematical constructions, case analyses and proofs, together with the Lean formalization and a first draft of this paper, were produced by Claude (Anthropic; models Fable 5, Opus 4.8 and Opus 5) acting as the reasoning agent, under the author’s direction. GPT-5.5 (OpenAI) and Gemini 3.1 (Google DeepMind) were used for advice and for adversarial review of drafts.

The author selected the problem and fixed its formalization; set the intermediate targets and overall strategy; decided which candidate routes to pursue and which to abandon; audited each construction and proof; chose the formal statement; and made the final mathematical judgements. He takes full responsibility for the content, including any error that remains. The automated systems are tools, not authors: they meet no authorship criterion, cannot take responsibility for the work, and are not listed as authors.

Formal verification settles correctness relative to the formal statement only; it does not establish novelty, attribution, or exposition. The one thing no formalization can settle on the reader’s behalf — that the formal statement means what the informal problem says — is set out decision by decision in §26.5.

If you are a student reading this: the theorem is true because of the proofs you have just read and can check — every one of them, with enough patience, by hand — and not because any person or machine says so. That is the whole point of mathematics, and it is why a paper like this one can exist at all.

28. DECLARATIONS

Funding. None.

Competing interests. The author declares no competing interests.

Author contributions. Sole author. The division of labour between the author and the automated systems used is set out in §27.

Use of generative AI. Automated systems played a substantial role in producing this work; the full account — models, tasks, division of labour, and verification — is §27. The author directed the work, audited the output, and takes full responsibility for the content; the systems are tools, not authors. Every mathematical claim used in the proof of Theorem A has been checked by at least one of the routes of §26, and the dichotomy is machine-checked in Lean 4 within the scope stated in §26.5.

A note on repository metadata. The version-control history of the accompanying repository contains **Co-Authored-By** trailers naming the automated systems. These record which tooling produced a commit; they are not a claim of scholarly authorship, which is as stated above.

Data and code availability. See §29.

29. DATA AVAILABILITY AND REPRODUCTION

Everything needed to check this paper is contained in one commit of the repository https://github.com/beetree/math_erdos_1112. That commit carries this $\text{\LaTeX}$ source, the Lean 4 development [Lean4, Mathlib], the machine-readable certificate exports (data/table-A.csv, data/table-B.csv, data/SHA256SUMS), the table generator and re-verifier gen-tables.py, the two independent Python harnesses, and the documented prior-art search (paper/novelty-search.md).

The machine-checked proof. A clean checkout builds with the pinned toolchain:

```
lake exe cache get           # fetch the pinned Mathlib
lake build                   # compile every file
lake env lean Erdos1112Proof/AxiomsCheck.lean # the 3-axiom report, per theorem
```

Continuous integration re-runs the pinned build, the axiom audit, the table re-verification and the correspondence check of §E on every push.

The finite layer. The complete certificate tables are typeset in the Appendix D and exported as CSV with recorded SHA-256 digests. Running make regenerates both from the one canonical source and re-checks every row (budget $\leq M - 1$, and coverage of M consecutive integers, by exact bitmask arithmetic, aborting on any failure), so no listing — printed, exported, or kernel-decided — can silently diverge from the others. Any single row can also be checked by hand in a few minutes, by the method of §26.

APPENDIX A. EVERY EXPANSION, IN FULL

The main text keeps ideas in the foreground and refers routine expansions here. Each item is pure algebra — no cleverness, only care.

(A1) Lemma 18.3: the run has length $\geq M$. With $x = M - \alpha$, $S = (\alpha - 1)\gamma$ and $\gamma \leq M - \alpha - 1$:

$$\begin{aligned} \alpha x - S + 1 &= \alpha(M - \alpha) - (\alpha - 1)\gamma + 1 \\ &\geq \alpha(M - \alpha) - (\alpha - 1)(M - \alpha - 1) + 1 \\ &= [\alpha M - \alpha^2] - [\alpha M - \alpha^2 - M + 1] + 1 \quad \text{since } (\alpha - 1)(M - \alpha - 1) = \alpha M - \alpha^2 - M + 1 \\ &= M. \end{aligned}$$

(The middle expansion: $(\alpha - 1)(M - \alpha - 1) = \alpha M - \alpha^2 - \alpha - M + \alpha + 1 = \alpha M - \alpha^2 - M + 1$.)

(A2) Lemma 18.4: the covered length. With $x = \beta - 1$, $y = \alpha - 1 + t$, $C = (\alpha - 1)(\beta - 1) = \alpha\beta - \alpha - \beta + 1$:

$$\begin{aligned} \alpha x + \beta y - 2C + 1 &= \alpha(\beta - 1) + \beta(\alpha - 1 + t) - 2(\alpha\beta - \alpha - \beta + 1) + 1 \\ &= \alpha\beta - \alpha + \alpha\beta - \beta + t\beta - 2\alpha\beta + 2\alpha + 2\beta - 2 + 1 \\ &= \alpha + \beta - 1 + t\beta. \end{aligned}$$

(A3) Lemma 18.6: the worst-case rearrangement. At $M = d\beta + 1$: $y = M - d - \beta + 1 = (d - 1)(\beta - 1) + 1$. Then, with $C = (\alpha - 1)(\beta - 1)$,

$$\begin{aligned} L &= \alpha x + \beta y - 2C + 1 = \alpha(\beta - 1) + \beta[(d - 1)(\beta - 1) + 1] - 2(\alpha - 1)(\beta - 1) + 1 \\ &= (\beta - 1)[\alpha + \beta(d - 1) - 2\alpha + 2] + \beta + 1 = (\beta - 1)[\beta(d - 1) - \alpha + 2] + \beta + 1. \end{aligned}$$

The claim $L \geq M + 1 = d\beta + 2$ is therefore

$$(\beta - 1)[\beta(d - 1) - \alpha + 2] \geq d\beta + 2 - \beta - 1 = \beta(d - 1) + 1,$$

i.e. $(\beta - 1)(u - \alpha + 2) \geq u + 1$ with $u := \beta(d - 1)$, as stated. At $d = 2$ ($u = \beta$): expand $(\beta - 1)(\beta - \alpha + 2) = \beta^2 - \alpha\beta + 2\beta - \beta + \alpha - 2$, so the claim is $\beta^2 - \alpha\beta + \beta + \alpha - 2 \geq \beta + 1$, i.e. $\beta(\beta - \alpha) + \alpha - 3 \geq 0$. Monotonicity in u : the difference $(\beta - 1)(u - \alpha + 2) - (u + 1)$ has u -coefficient $(\beta - 1) - 1 = \beta - 2 \geq 0$ for $\beta \geq 2$; and if $\beta = 2$ then $\alpha = 1$ and the difference is $(u + 1) - (u + 1) = 0$ — equality, which suffices.

(A4) Lemma 18.6: the bridge length.

$$(L - 1)d - (d - 1)M + 1 \geq M \iff (L - 1)d \geq dM - 1 \iff L \geq M + 1 - \frac{1}{d}.$$

Since L is an integer and $0 < \frac{1}{d} < 1$, the number $M + 1 - \frac{1}{d}$ is not an integer and lies strictly between M and $M + 1$; so the condition is exactly $L \geq M + 1$.

(A5) Case P, inequality (E1). Start from the sufficient form

$$2[(a - 1)(ra - M) + M - 1] \leq ar(M - r - 1) + 2a(r - 1).$$

Left side: $2(a - 1)ra - 2(a - 1)M + 2M - 2 = 2ra^2 - 2ra - 2aM + 4M - 2$. Right side: $arM - ar^2 - ar + 2ar - 2a = arM - ar^2 + ar - 2a$. Moving everything right, the inequality is $0 \leq arM + 2aM - 4M - ar^2 - 2ra^2 + 3ar - 2a + 2 = M(ar + 2a - 4) - ar^2 - 2ra^2 + 3ar - 2a + 2$, as displayed in the text. Now insert $2M \geq ra + 1$ (valid since $ar + 2a - 4 > 0$), multiplying the target by 2:

$$\begin{aligned} 2M(ar + 2a - 4) - 2ar^2 - 4ra^2 + 6ar - 4a + 4 \\ &\geq (ra + 1)(ar + 2a - 4) - 2ar^2 - 4ra^2 + 6ar - 4a + 4 \\ &= [a^2r^2 + 2a^2r - 4ar + ar + 2a - 4] - 2ar^2 - 4ra^2 + 6ar - 4a + 4 \\ &= a^2r^2 - 2a^2r - 2ar^2 + 3ar - 2a = a(ar^2 - 2ar - 2r^2 + 3r - 2). \end{aligned}$$

So (E1) holds whenever $ar^2 - 2ar - 2r^2 + 3r - 2 \geq 0$, i.e. whenever $ar(r - 2) \geq 2r^2 - 3r + 2$ — the condition (*) of the text.

APPENDIX B. HINTS AND ANSWERS FOR THE EXERCISES

One entry per exercise, in reading order, titled by its section and topic. Entries come in four kinds: a *hint* leaves you work; an *answer* settles the matter; CHALLENGE marks something genuinely hard and safely skipped; DISCUSSION marks an open-ended prompt with no single clean answer — an invitation to experiment, not a debt. If an exercise has parts, the entry follows them.

Part I. §1, walks. (a) Any aperiodic rule works, e.g. steps 3, 2, 3, 3, 2, 3, 3, 3, 2, ... (one 2 after ever more 3s); nothing obliges a walk to repeat. (b) Sum the $n - 1$ gap bounds $d_1 \leq g_i \leq d_2$.

§1, sumsets. (a) All members are odd, and sums of three odd numbers are odd; the answer is every odd number from 3 to 25. (b) Each summand is $\geq a_1$. (c) $ka_{n+1} - ka_n = kg_n \leq kd_2$.

§1, lacunary growth. Induction: $b_i \geq rb_{i-1}$. Then $b_1r^{i-1} \leq b_i \leq N$ gives $i \leq 1 + \log_r(N/b_1) \leq 1 + \log_2 N$.

Exercise 1.14. (a) $kA = \{kc + sd : s \geq k\}$ (Example 2.4); the Adversary picks $b_1 = kc + kd$, then b_{i+1} a far-later element of the same progression. (b) CHALLENGE — this is Part III's

certificate trick; the honest answer at this stage is “not obviously, since different progressions own different classes”.

Exercise 2.2. Add $x_t - 1 < \lfloor x_t \rfloor \leq x_t$ over t .

Exercise 2.3. (a) Only sums of terms ≤ 19 matter; there are finitely many, and none hits 18 or 25. (b) Any sum using a term > 25 exceeds 25. (c) 32 and 39 (the pattern is $18 + 7j$: daylight recurs once per two strides here).

Exercise 2.5. $i_t \geq 1$ forces $s \geq k$; conversely $i_1 = s - k + 1$, $i_2 = \dots = i_k = 1$.

Part II. §3, floors. (a) 5, -3 , -4 . (b) $x = y = 0.5$: $\lfloor 1 \rfloor = 1 \neq 0$; in general write $x = \lfloor x \rfloor + \{x\}$ and split on whether $\{x\} + \{y\} < 1$. (c) $m + n$ is no longer an integer, so the defining property does not transfer.

§3, Beatty walks. (a) $\lfloor id \rfloor = id$. (b) Gap 4 exactly at odd i : $\lfloor (i+1) \cdot \frac{7}{2} \rfloor - \lfloor i \cdot \frac{7}{2} \rfloor$ depends on the parity of i . (c) Telescope: the sum of the first n gaps is $\lfloor (n+1)\gamma \rfloor - \lfloor \gamma \rfloor$; each gap is $\lfloor \gamma \rfloor$ or $\lfloor \gamma \rfloor + 1$, so the number of big gaps is the total minus $n\lfloor \gamma \rfloor$.

Exercise 4.5. (a) $s = 8$: $3.5 \cdot 8 = 28 \in [25, 28]$. (b) 28 sits on the *closed* right endpoint; the half-open window means membership of 25 in $3A$ would require $s\gamma \in [25, 28)$ for some s ; the only candidate is $3.5 \cdot 8 = 28$, which the half-open interval excludes — so the dodge succeeds. The *closed* certification interval $[25, 28]$ does contain 28, so the certificate fails: conservative, not wrong. (c) The smallest multiple u of 3.5 with $u \geq t$ satisfies $u < t + 3.5$; multiples of 3.5 are integers or half-integers, and $(t+3, t+3.5)$ contains neither, so $u \leq t+3$: every closed window is hit. (d) ANSWER. The unsafe intervals meeting $(3, 4)$ are exactly those with $s = 12, \dots, 16$:

$$\left[\frac{46}{12}, \frac{49}{12}\right], \left[\frac{46}{13}, \frac{49}{13}\right], \left[\frac{46}{14}, \frac{49}{14}\right], \left[\frac{46}{15}, \frac{49}{15}\right], \left[\frac{46}{16}, \frac{49}{16}\right],$$

i.e. approximately $[3.833, 4.083]$, $[3.538, 3.769]$, $[3.286, 3.500]$, $[3.067, 3.267]$, $[2.875, 3.063]$. The certified slopes are what is left of $(3, 4)$ — *four* open gaps, not one:

$$\left(\frac{49}{16}, \frac{46}{15}\right), \left(\frac{49}{15}, \frac{46}{14}\right), \left(\frac{49}{14}, \frac{46}{13}\right), \left(\frac{49}{13}, \frac{46}{12}\right),$$

approximately $(3.063, 3.067)$, $(3.267, 3.286)$, $(3.500, 3.538)$ and $(3.769, 3.833)$. The text’s $\gamma = 3.8$ lies in the last and widest of them — and note how the gaps widen toward the right, exactly as §5 exploits by aiming s_1 at the top of the interval.

§4, $t = 20$. (a) $s = 6, 7$ give $[10/3, 23/6]$ and $[20/7, 23/7]$; all other unsafe intervals miss $(3, 4)$. (b) Any $\gamma \in (23/6, 4)$, e.g. $\gamma = 3.9$: then $b = 20 + 3c \notin 3A_{\gamma, c}$.

Exercise 4.8. Cross-multiply: $(t+k)s \geq t(s+1) \iff ks \geq t$.

Exercise 5.4. (a) $10000 \geq 512$ and $0.5 \geq 0.0064$. (b) About $t/\gamma^2 \cdot |I| \approx 300$ of them. (c) $s_1 = 250000$, $J = (10^6+3)/250001 \dots 4$, width = $1/250001$.

§5, hypotheses. $t \geq 32d_2^2$ enters at (2), in the width chain (via $t/(32(d_2 - \frac{1}{2})) \geq d_2 + \frac{1}{2}$), and via the input-width hypothesis in check (iii). For small t try $t = 20$, $d_2 = 4$: J is wide but pokes below lo .

§6, trace. $j_0 = 1$ ($b_1 = 800 \geq 64$), $c = 801$, $t_2 = 800^2 - 3 \cdot 801 = 637,597$; then $s_1 = \lceil t_2/(d_2 - \frac{1}{4}) \rceil$ etc. — any slope in the resulting middle third qualifies.

§6, eventual lacunarity. Only the choice of j_0 changes: pick it past i_0 as well.

Part III. Exercise 7.9. (a) Each step preserves common divisors ($d \mid x, y \Rightarrow d \mid x - qy$); $\gcd(2026, 715) = \gcd(715, 596) = \gcd(596, 119) = \gcd(119, 1) = 1$. (b) $21 \cdot 5 - 8 \cdot 13 = 105 - 104 = 1$. (c) $8^{-1} \equiv 8 \pmod{21}$ ($64 = 63 + 1$); $s \equiv 8 \cdot 5 = 40 \equiv 19 \pmod{21}$.

Exercise 7.10. (a) 0, 5, 3, 1, 6, 4, 2: all seven. (b) 0, 6, 4, 2 — only the even classes. (c) Multiply by the inverse; solutions $s \equiv \bar{a}c$.

§8, the 3,5 example. (a) Two gaps sum to 8; a prefix of the gap word is m full $(3,5)$ -blocks plus possibly a leading 3. (b) Fix all five indices even (contributing $8m_i$); “sufficiently large” is $5a_1 + 8 \cdot 0$ onward within the class. (c) Sums are $\equiv 3j \pmod{8}$ with $j \leq 5$, and $2 \notin \{0, 3, 6, 1, 4, 7\}$.

§8, certificate values. $b_2 = 11$, $b_3 = 1102$, $b_4 = 1,102,001$; ratios $11 \geq 10$, $1102/11 \geq 100$, etc.

§10, $(2,3,2)$. Prefix sums modulo 7 are $\{0, 2, 5\}$; sums of three of them hit all seven classes — so $3A - 3a_1$ eventually contains *every* class modulo 7, i.e. all large integers.

§11, balance. (b): windows 00 and 11. Final claim: compare both windows to the initial window of the same length.

§12, 0011... $q = 0, 0, 1, 2, 2, 2, 3, 4, \dots$; $\sigma_0 = 4$: $q_0 + q_4 = 2$, $q_2 + q_2 = 0$, so $W_2(4) = 2$. For $s = 12$: configurations $(4, 4, 4)$ (value 6) and $(2, 2, 8)$ (value 4).

§12, W_2 vs width. Pad with $k - 2$ zero indices ($q_0 = 0$): a two-index spread survives into k indices. Strict inequality: almost any word with $k = 3$, e.g. 0011... at $s = 12$ has width ≥ 3 .

§13, examples. (a) $\alpha = \frac{1}{3}$; D_n cycles through $0, -\frac{1}{3}, -\frac{2}{3}$ shifted... compute: $q_{3m} = m$, so $D_{3m} = 0$; $w_n = 1$ for all n , and Step 4’s argument ends immediately at period 3. (b) $\{j\alpha\}$ for $j = 0..5$: two land within $\frac{1}{5}$; Dirichlet gives $n_0 = 5$, $\theta = |5\sqrt{2} - 7| \approx 0.071$, so $L = 5 \cdot 16 = 80$. CHALLENGE to tighten: the true minimal L is far smaller.

§13, irrationality. Once in “at most one exceptional index” (two would give $(n_1 - n_2)\alpha \in \mathbb{Z}$), once in $\alpha \notin \{0, 1\}$. For rational α the exceptional indices recur periodically and the tail-pass fails — consistent with Step 4 taking over.

Exercise 14.8. $q_{10} = 4$, $q_{26} = 10$, $q_{36} = 14$: sums 72 and 28. Another valid triple: $(10, 23, 39)$, with $4 + 9 + 15 = 28$.

§14, why not the sweep. CHALLENGE. W_s sits inside a real window of length k (Step 0), so $\text{width}(s) \leq k - 1$ always; for a Sturmian word the two extreme integers of the window are attained only when fractional parts pile up near 0 (resp. near 1) *simultaneously*, which balancedness forbids infinitely often — the band is genuinely too narrow at the corner $\delta' = d_2 = k$.

§15, small m ’s. (a) $\{1, 2\}$: sums $0, 1, 2, 3$; $\{2, 3, 4\}$: sums $0, 2, 3, 4, 5, 6, 7, 9 \supseteq [2, 7]$. (b) Use $S = \{2, 3, 4\}$ from (a): its sums cover $[2, 7]$, amply containing $M = 4$ consecutive values. Slots at three positions with gaps 2, 3, 4 inside the first two periods; nothing to park ($k - 1 - m = 0$); dial ϕ down by steps ≤ 4 into the block.

§15, parking and leaping. Parking at $p_0 = 0$ is legal (indices ≥ 0 are allowed in Lemma 9.1) and changes only base. Leap example: block $\{4, 5\}$ ($M = 3$ would need length 3), dial multiples of 3: ϕ passes $6 \rightarrow 3$, skipping both.

Part IV. §16, small cases. (a) Answer: $m(\{1, 2, M\}) = \lceil M/2 \rceil$. Construction: one 1 and $j := \lceil (M-2)/2 \rceil$ twos — their sums fill $[0, 2j+1]$ solidly (twos reach the even numbers, the 1 adjusts parity), a run of $2j+2 \geq M$, at size $1+j = \lceil M/2 \rceil$. The matching lower bound is a small CHALLENGE: show a run of M consecutive sums needs both parities (hence the 1, for even M) and total reach $\geq M-1$ from the 1s and 2s alone — copies of M cannot glue a run longer than the $\{1, 2\}$ -part supplies. (b) DISCUSSION, genuinely open-ended: $2^s \geq M$ forces $s \geq \log_2 M$, but with only three generators the sums form a rigid three-parameter grid; experiment with $G = \{a, M-1, M\}$ and watch $\min G$ control how far below $M-1$ the truth sits. No clean closed answer is intended here.

Exercise 17.2. Coprimality enters choosing j with $j\beta \equiv n$; for $(2, 4)$ all grid values are even — no odd number is ever covered, and no interval of length ≥ 2 either.

§17, (5, 7, 9) and $\{5^{(x)}, 7\}$. The sums of x fives and one seven occupy only residues $\{0, 2\}$ modulo 5: no five consecutive integers, for any x . Missing ingredient: a box covering *all* residues.

§17, the lift's hypothesis. ANSWER. $Y + Z \leq a - 1$ caps the budget growth at $Y + Z + 1 \leq a$ per lift, matched by the target's growth of exactly a . A concrete instance of the failure: take $a = 5$, and a box with $Y + Z = 5 = a$, say $(Y, Z) = (3, 2)$, on the class base $(5, 16, 17)$. The representative heights are now bounded by $S_0 = 3 \cdot 16 + 2 \cdot 17 = 82$ rather than 34, so the budget is $\lceil (16 + 82)/5 \rceil + 5 = 20 + 5 = 25$ against a target $M - 1 = 16$ — already failing, and the lift makes it worse: at $(5, 21, 22)$ the same box gives $S_0 = 3 \cdot 21 + 2 \cdot 22 = 107$, budget $\lceil (21 + 107)/5 \rceil + 5 = 26 + 5 = 31$ against target 21. The budget rose by $6 = Y + Z + 1$, the target by $5 = a$: the margin lost one unit on the step, and will lose one on every step forever. With $Y + Z \leq a - 1$ that leak is exactly plugged.

§18, routing. $(5, 12, 15)$: D. $(3, 4, 5)$: P. $(5, 7, 9)$: L. $(13, 22, 27)$: E. $(5, 8, 9)$: T. $(5, 16, 17)$: B. And $(12, 25, 36)$ is not hard-core: $a + b = 37 = M + 1$, so Lemma 18.4 catches it.

§18, smallest members. $(a, a + 1, a + 2)$ is hard-core for every $a \geq 3$ ($\gcd(a, a + 1) = 1$, $h = 1 \leq a - 2$); nothing smaller than $(3, 4, 5)$ exists since $a \geq 3$ and $M \geq a + 2$.

§19, Case L on $(5, 7, 9)$. $(x, y, z) = (5, 1, 2)$: budget $8 = M - 1$. Levels $c = 3$ through $c + g = 5$ cover solidly; the difference from $\{5, 5, 5, 7\}$ is exactly the extra frames that make the merge condition true.

§20, $(3, 4, 5)$. $\text{low} = \max(1 \cdot 4, 1 \cdot 5) = 5$; targets $5, \dots, 9$; e.g. $8 = 4 + 4$? No — follow the branches: $8 \equiv 2 \pmod{3}$, $w = 2 > n_b = 1$, so $v = 1$: $8 = 5 + 3$. The mover appears for target $9 = 4 + 5$.

§20, (E1) derivation. Mechanical; item (A5) shows the target form.

Exercise 21.3. Answer, on the concrete triple $(13, 24, 34)$, which routes to E with $\eta = 9$ (check: $\bar{e} = 11$, $\bar{\mu} = 8$, $11^{-1} \equiv 6$, $\eta = 48 \pmod{13} = 9$). Working from η directly would take the box $(8, 1)$: top height $8 \cdot 24 + 34 = 226$, budget $\lceil (33 + 226)/13 \rceil + 9 = 20 + 9 = 29$. Working from $t = \min(9, 4) = 4$ takes the box $(3, 3)$: top height $3 \cdot 24 + 3 \cdot 34 = 174$, budget $\lceil (33 + 174)/13 \rceil + 6 = 16 + 6 = 22$. Both pass here ($M - 1 = 33$), but the min is systematically cheaper — and it is what keeps $\sigma \leq a/2 + 2$, which the general threshold argument needs. Final claim: $2 + (a - 1)/2 = (a + 3)/2$: multiply out.

§22, $(7, 9, 10)$. $y = 2$, $z = 2$, $x = 5$: multiset $\{7^5, 9^2, 10^2\}$, nine elements; levels 4 and 5 cover $[28 + C, 35 + V]$ with the merge exact; budget $9 = M - 1$.

§22, $(5, 8, 9)$. Construction budget $11 > 8$ (the line $(3, 1)$ is unlucky at $a = 5$); the replacement certificate is checked by listing $8j + 9k$ residues $\{0, 3, 1, 4, 2\}$ and the run $[17, 25]$.

§23, the base and one lift. Residues of $16j + 17k$: $\{0, 1, 2, 3, 4\}$; heights ≤ 50 ; budget 13. At $(5, 21, 22)$ the exact budget is 16 (recompute $S = 44$, $x = 13$) — better than the lemma's guaranteed $13 + 4 = 17$.

§24, overlaps. Answer, and a pleasing one: both branches produce the *same* multiset. Case L at $(3, 4, 5)$ has $a = 3$, $g = 1$, so $(x, y, z) = (\lfloor \frac{1}{2} \rfloor + 2, 1, \lceil \frac{1}{2} \rceil) = (2, 1, 1)$ — the multiset $\{3, 3, 4, 5\}$, exactly Case P's witness. The two mechanisms converge at the smallest triple. D-and-E-sizes example: $(13, 25, 26)$ ($a \mid M$, both sizes ≥ 12); the router reads conditions in order, so only D fires — overlap costs nothing.

Part V. §25, grand tour. CHALLENGE, and the best single exercise in the document. Route: two letters $\{1, 2\}$, aperiodic balanced (Fibonacci) $\rightarrow$ Sturmian $\rightarrow$ modulus $m = 1$ from Lemma 14.1; the certificate meets kA at any listed pair $(1, 0)$ whose b_i clears the walk's threshold X_0 .

APPENDIX C. SYMBOL GLOSSARY

By part, in rough order of first appearance. A few letters are reused across parts (never within one); reuses are flagged.

Global.

k	number of summands	fixed throughout
$d_1 \leq d_2$	the gap window	fixed throughout
A, a_i, g_n	the walk, its terms, its gaps	
B, b_i, r	the Adversary's sequence, its ratio	
kA	the k -fold sumset, repetitions allowed	Definition 1.6
$\lfloor \cdot \rfloor, \lceil \cdot \rceil, \{ \cdot \}$	floor, ceiling, fractional part	§3, §5, §13

Part II.

γ	the slope (dial position)	$\in (d_2 - 1, d_2)$
c	the walk's offset	reused: $y+z$ in Part IV, run start in §15
$A_{\gamma,c}$	the Beatty walk	Definition 3.4
t, t_i	shifted target(s) $b - kc$	reused: $\min(\eta, a - \eta)$ in Case E
U_t, s_1, J, L, R	unsafe set; gap anatomy	§4.3–5
I_j, γ^*	the nest and its limit point	§6

Part III.

p_n	prefix sums of the gaps	§9
G_∞, T	tail alphabet; tail start	§9
$\delta < \delta', e$	the two letters; $e = \delta' - \delta$	reused: $e = b - a$ in Part IV
h, h_n	the binary gap word	reused: $h = M - b$ in Part IV
q_n	counting function of h	§11
$W_s, w^\pm(s)$, width	the reachable band	§12
$W_2(\sigma)$	two-index width	§12
α, D_n, β	slope, discrepancy, intercept	§13; α, β reused in Lemma 17.1
u/v	α in lowest terms	Step 4
$\theta, \psi, \zeta, \kappa$	depth, rung, margin, padding	§14
$m(G)$, sums	covering cost; subset sums	Definition 15.1

Part IV.

$a < b < M$	the triple	$\gcd(a, b) = 1$ on the hard core
e, h, μ	$b - a, M - b, M - a = e + h$	§17.3
g, e', μ', C', V'	$\gcd(e, \mu)$; divided forms; grid corners	§17.3
$(x, y, z), c = y + z$	copies of a, b, M ; box total	
$\mathcal{O}_t$	offsets at level t	staircase
$(Y, Z), S, \mathcal{B}$	box; top height; budget	frame lemma
$\bar{e}, \bar{\mu}, \eta$	mod- a data; $\eta = \bar{\mu} \bar{e}^{-1}$	§18.5
$t = \min(\eta, a - \eta), Z, \sigma$	Case E box data	
n_b, n_M	max unpaired bs, Ms (Case P)	
λ	$\lfloor e/a \rfloor$: rung on a class ladder (Case B)	

APPENDIX D. THE CERTIFICATE TABLES

These are the $172 + 178$ finite certificates that Cases T and B consume — the entire finite layer of the proof, printed in full so that this paper is self-contained.

Each entry is a triple (a, b, M) together with a *box certificate* (x, Y, Z) , asserting that the multiset S of x copies of a , Y copies of b and Z copies of M satisfies

- **budget:** $x + Y + Z \leq M - 1$, and
- **coverage:** $\text{sums}(S)$ contains M consecutive integers,

i.e. that S witnesses $m(G) \leq M - 1$ for $G = \{a, b, M\}$. The mechanism is the frame lemma (Lemma 17.3): the copies of b and M realize a modulo- a box covering every residue, and the x copies of a lift that residue cover to a solid run.

You can check any row of these tables by hand in a few minutes, and that is the point of printing them. Take the row $(5, 8, 9) \rightarrow (5, 2, 1)$, which is Exercise 22.5's worked instance: the multiset is $\{5, 5, 5, 5, 5, 8, 8, 9\}$, of size $5 + 2 + 1 = 8 = M - 1$, meeting the budget exactly. Its parts $8j + 9k$ with $j \leq 2, k \leq 1$ take the values $\{0, 8, 16, 9, 17\}$, whose residues modulo 5 are $\{0, 3, 1, 4, 2\}$ — all five. Lifting by the copies of 5 gives the run $[17, 25]$, which is $9 = M$ consecutive integers. That single computation is exactly what the checker performs on each row, and what the Lean kernel re-decides.

These tables are *generated*, not transcribed: `gen-tables.py` parses the canonical source `data/certificate-data.md`, re-verifies every row's budget and coverage by exact arbitrary-precision bitmask, aborts the build on any failure, and only then emits the L^AT_EX below. The same rows are exported as CSV with recorded digests and transcribed into Lean, where the kernel decides them (§26). No listing — printed, exported, or kernel-checked — can silently diverge from the others.

Table 1: The 172 exceptional Case-T triples (the T-tail lemma of the main paper), all with $a \leq 29$; the scan over $a \leq 3000$ finds no others. Format $(a, b, M) \rightarrow (x, Y, Z)$: the multiset of x copies of a , Y of b , Z of M .

$(4, 9, 10) \rightarrow (7, 1, 1)$	$(4, 13, 14) \rightarrow (10, 1, 1)$	$(5, 8, 9) \rightarrow (5, 2, 1)$	$(5, 9, 12) \rightarrow (7, 2, 1)$
$(5, 11, 12) \rightarrow (7, 1, 2)$	$(5, 11, 13) \rightarrow (8, 2, 1)$	$(5, 12, 14) \rightarrow (9, 1, 2)$	$(5, 13, 14) \rightarrow (8, 2, 1)$
$(5, 13, 16) \rightarrow (10, 1, 2)$	$(6, 11, 14) \rightarrow (9, 1, 2)$	$(6, 11, 15) \rightarrow (9, 2, 1)$	$(6, 13, 14) \rightarrow (9, 1, 2)$
$(6, 13, 15) \rightarrow (10, 2, 1)$	$(6, 13, 16) \rightarrow (10, 1, 2)$	$(7, 10, 12) \rightarrow (7, 3, 1)$	$(7, 10, 15) \rightarrow (8, 2, 2)$
$(7, 11, 12) \rightarrow (7, 2, 2)$	$(7, 11, 13) \rightarrow (7, 2, 2)$	$(7, 11, 16) \rightarrow (8, 3, 1)$	$(7, 12, 13) \rightarrow (7, 3, 1)$
$(7, 12, 15) \rightarrow (8, 2, 2)$	$(7, 13, 17) \rightarrow (9, 3, 1)$	$(7, 13, 18) \rightarrow (9, 2, 2)$	$(7, 15, 16) \rightarrow (9, 1, 3)$
$(7, 15, 17) \rightarrow (9, 2, 2)$	$(7, 15, 18) \rightarrow (10, 3, 1)$	$(7, 16, 17) \rightarrow (10, 2, 2)$	$(7, 16, 18) \rightarrow (11, 1, 3)$
$(8, 11, 12) \rightarrow (7, 3, 1)$	$(8, 11, 15) \rightarrow (9, 2, 2)$	$(8, 13, 14) \rightarrow (7, 3, 2)$	$(8, 13, 15) \rightarrow (8, 2, 2)$
$(8, 13, 17) \rightarrow (10, 2, 2)$	$(8, 15, 18) \rightarrow (9, 3, 2)$	$(8, 15, 19) \rightarrow (11, 2, 2)$	$(8, 17, 18) \rightarrow (11, 1, 3)$
$(8, 17, 19) \rightarrow (10, 2, 2)$	$(9, 13, 16) \rightarrow (8, 3, 2)$	$(9, 13, 20) \rightarrow (9, 4, 1)$	$(9, 14, 15) \rightarrow (8, 2, 2)$
$(9, 14, 16) \rightarrow (9, 4, 1)$	$(9, 14, 17) \rightarrow (9, 2, 3)$	$(9, 14, 20) \rightarrow (9, 3, 2)$	$(9, 16, 17) \rightarrow (9, 4, 1)$
$(9, 16, 19) \rightarrow (10, 3, 2)$	$(9, 17, 20) \rightarrow (11, 2, 3)$	$(9, 19, 20) \rightarrow (11, 1, 4)$	$(10, 13, 15) \rightarrow (9, 4, 1)$
$(10, 13, 18) \rightarrow (10, 3, 2)$	$(10, 13, 21) \rightarrow (9, 3, 2)$	$(10, 17, 18) \rightarrow (9, 3, 2)$	$(10, 17, 19) \rightarrow (9, 3, 2)$
$(10, 17, 21) \rightarrow (10, 2, 3)$	$(11, 14, 15) \rightarrow (7, 3, 4)$	$(11, 14, 18) \rightarrow (9, 5, 1)$	$(11, 14, 21) \rightarrow (10, 3, 2)$
$(11, 15, 16) \rightarrow (7, 3, 2)$	$(11, 15, 20) \rightarrow (9, 4, 2)$	$(11, 16, 18) \rightarrow (8, 4, 2)$	$(11, 16, 19) \rightarrow (10, 5, 1)$
$(11, 16, 20) \rightarrow (9, 3, 2)$	$(11, 17, 18) \rightarrow (8, 2, 3)$	$(11, 18, 19) \rightarrow (9, 4, 3)$	$(11, 18, 20) \rightarrow (11, 5, 1)$
$(11, 18, 21) \rightarrow (10, 2, 3)$	$(11, 19, 20) \rightarrow (9, 4, 2)$	$(11, 20, 21) \rightarrow (11, 5, 1)$	$(12, 17, 18) \rightarrow (10, 5, 1)$
$(12, 17, 20) \rightarrow (10, 3, 2)$	$(12, 17, 21) \rightarrow (10, 2, 3)$	$(12, 17, 23) \rightarrow (12, 4, 2)$	$(12, 19, 20) \rightarrow (10, 3, 2)$
$(12, 19, 21) \rightarrow (11, 2, 3)$	$(12, 19, 22) \rightarrow (11, 3, 4)$	$(12, 19, 23) \rightarrow (11, 4, 2)$	$(13, 16, 21) \rightarrow (10, 6, 1)$
$(13, 16, 24) \rightarrow (11, 2, 4)$	$(13, 17, 20) \rightarrow (9, 4, 2)$	$(13, 17, 24) \rightarrow (10, 4, 4)$	$(13, 18, 19) \rightarrow (9, 4, 2)$
$(13, 18, 20) \rightarrow (9, 3, 3)$	$(13, 18, 22) \rightarrow (11, 6, 1)$	$(13, 18, 24) \rightarrow (11, 3, 3)$	$(13, 19, 24) \rightarrow (10, 3, 3)$
$(13, 20, 21) \rightarrow (10, 2, 4)$	$(13, 20, 22) \rightarrow (10, 4, 2)$	$(13, 20, 23) \rightarrow (12, 6, 1)$	$(13, 20, 24) \rightarrow (11, 4, 2)$
$(13, 21, 22) \rightarrow (10, 4, 4)$	$(13, 21, 24) \rightarrow (11, 2, 4)$	$(13, 22, 23) \rightarrow (11, 3, 3)$	$(13, 22, 24) \rightarrow (13, 6, 1)$
$(13, 23, 24) \rightarrow (11, 4, 2)$	$(14, 17, 24) \rightarrow (12, 5, 2)$	$(14, 19, 20) \rightarrow (9, 3, 3)$	$(14, 19, 21) \rightarrow (12, 6, 1)$
$(14, 19, 22) \rightarrow (9, 5, 2)$	$(14, 19, 25) \rightarrow (10, 4, 2)$	$(14, 23, 24) \rightarrow (12, 3, 5)$	$(14, 23, 25) \rightarrow (11, 3, 4)$
$(15, 19, 24) \rightarrow (11, 2, 4)$	$(15, 22, 24) \rightarrow (11, 2, 4)$	$(15, 22, 25) \rightarrow (11, 4, 2)$	$(15, 22, 26) \rightarrow (13, 7, 1)$

Table 1: (continued).

$(15, 23, 24) \rightarrow (11, 2, 4)$	$(15, 23, 26) \rightarrow (12, 5, 4)$	$(16, 19, 27) \rightarrow (13, 6, 2)$	$(16, 21, 23) \rightarrow (10, 5, 2)$
$(16, 21, 24) \rightarrow (13, 7, 1)$	$(16, 21, 25) \rightarrow (10, 4, 3)$	$(16, 23, 24) \rightarrow (13, 7, 1)$	$(16, 23, 26) \rightarrow (11, 5, 2)$
$(16, 23, 27) \rightarrow (12, 3, 4)$	$(16, 25, 26) \rightarrow (10, 3, 4)$	$(16, 25, 27) \rightarrow (12, 2, 5)$	$(17, 21, 28) \rightarrow (11, 3, 4)$
$(17, 22, 23) \rightarrow (10, 5, 6)$	$(17, 22, 25) \rightarrow (10, 4, 3)$	$(17, 22, 26) \rightarrow (10, 6, 2)$	$(17, 22, 28) \rightarrow (13, 8, 1)$
$(17, 23, 24) \rightarrow (10, 3, 4)$	$(17, 24, 25) \rightarrow (10, 5, 2)$	$(17, 24, 26) \rightarrow (11, 4, 4)$	$(17, 24, 28) \rightarrow (11, 3, 4)$
$(17, 25, 28) \rightarrow (11, 6, 2)$	$(17, 26, 27) \rightarrow (11, 2, 5)$	$(17, 26, 28) \rightarrow (11, 4, 3)$	$(17, 27, 28) \rightarrow (12, 4, 3)$
$(18, 23, 27) \rightarrow (14, 8, 1)$	$(18, 23, 29) \rightarrow (12, 4, 3)$	$(18, 25, 26) \rightarrow (10, 5, 3)$	$(18, 25, 27) \rightarrow (15, 8, 1)$
$(18, 25, 28) \rightarrow (11, 3, 4)$	$(19, 24, 30) \rightarrow (11, 5, 3)$	$(19, 25, 30) \rightarrow (11, 4, 3)$	$(19, 26, 27) \rightarrow (10, 4, 4)$
$(19, 26, 28) \rightarrow (11, 3, 4)$	$(19, 26, 29) \rightarrow (12, 3, 4)$	$(19, 26, 30) \rightarrow (12, 6, 2)$	$(19, 27, 28) \rightarrow (12, 6, 2)$
$(19, 28, 30) \rightarrow (13, 3, 5)$	$(19, 29, 30) \rightarrow (13, 2, 6)$	$(20, 27, 28) \rightarrow (11, 3, 4)$	$(20, 27, 29) \rightarrow (12, 6, 2)$
$(20, 27, 30) \rightarrow (16, 9, 1)$	$(20, 27, 31) \rightarrow (12, 5, 4)$	$(20, 29, 30) \rightarrow (16, 9, 1)$	$(21, 29, 30) \rightarrow (10, 5, 4)$
$(21, 29, 32) \rightarrow (12, 3, 5)$	$(22, 29, 31) \rightarrow (11, 6, 3)$	$(22, 29, 32) \rightarrow (11, 5, 4)$	$(22, 29, 33) \rightarrow (17, 10, 1)$
$(22, 31, 32) \rightarrow (12, 5, 3)$	$(22, 31, 33) \rightarrow (18, 10, 1)$	$(23, 30, 31) \rightarrow (13, 7, 8)$	$(23, 30, 33) \rightarrow (13, 7, 2)$
$(23, 30, 34) \rightarrow (13, 4, 4)$	$(23, 31, 32) \rightarrow (11, 3, 5)$	$(23, 31, 34) \rightarrow (11, 6, 4)$	$(23, 32, 33) \rightarrow (13, 3, 5)$
$(23, 32, 34) \rightarrow (12, 4, 4)$	$(23, 33, 34) \rightarrow (13, 7, 2)$	$(24, 31, 35) \rightarrow (12, 4, 4)$	$(25, 33, 36) \rightarrow (14, 8, 2)$
$(25, 34, 35) \rightarrow (13, 4, 4)$	$(25, 34, 36) \rightarrow (13, 3, 6)$	$(26, 35, 36) \rightarrow (13, 3, 6)$	$(26, 35, 37) \rightarrow (13, 6, 3)$
$(27, 37, 38) \rightarrow (14, 6, 3)$	$(28, 37, 39) \rightarrow (14, 4, 5)$	$(29, 38, 39) \rightarrow (15, 9, 10)$	$(29, 39, 40) \rightarrow (14, 3, 7)$

Table 2: The 178 Case-B class bases ($a \leq 11$, $\mu \geq 12$). Format $[a; \bar{e}, h]$ followed by the base triple $(a, b, M) \rightarrow (x, Y, Z)$. The λ -lift lemma of the main paper extends each certificate to every member of its class with larger e .

$[4; 1, 1] (4, 17, 18) \rightarrow (13, 1, 1)$	$[5; 1, 1] (5, 16, 17) \rightarrow (10, 1, 2)$	$[5; 1, 2] (5, 16, 18) \rightarrow (11, 2, 1)$
$[5; 2, 2] (5, 17, 19) \rightarrow (12, 1, 2)$	$[5; 3, 1] (5, 18, 19) \rightarrow (11, 2, 1)$	$[5; 3, 3] (5, 18, 21) \rightarrow (13, 1, 2)$
$[5; 4, 3] (5, 14, 17) \rightarrow (10, 2, 1)$	$[6; 1, 1] (6, 19, 20) \rightarrow (13, 1, 2)$	$[6; 1, 2] (6, 19, 21) \rightarrow (14, 2, 1)$
$[6; 1, 3] (6, 19, 22) \rightarrow (14, 1, 2)$	$[6; 5, 3] (6, 17, 20) \rightarrow (13, 1, 2)$	$[6; 5, 4] (6, 17, 21) \rightarrow (13, 2, 1)$
$[7; 1, 1] (7, 22, 23) \rightarrow (13, 1, 3)$	$[7; 1, 2] (7, 22, 24) \rightarrow (13, 2, 2)$	$[7; 1, 3] (7, 22, 25) \rightarrow (14, 3, 1)$
$[7; 1, 4] (7, 15, 19) \rightarrow (11, 2, 2)$	$[7; 2, 1] (7, 23, 24) \rightarrow (14, 2, 2)$	$[7; 2, 2] (7, 23, 25) \rightarrow (15, 1, 3)$
$[7; 2, 4] (7, 16, 20) \rightarrow (11, 2, 2)$	$[7; 3, 2] (7, 17, 19) \rightarrow (11, 3, 1)$	$[7; 3, 3] (7, 17, 20) \rightarrow (12, 1, 3)$
$[7; 3, 5] (7, 17, 22) \rightarrow (12, 2, 2)$	$[7; 4, 1] (7, 18, 19) \rightarrow (11, 2, 2)$	$[7; 4, 2] (7, 18, 20) \rightarrow (11, 2, 2)$
$[7; 4, 4] (7, 18, 22) \rightarrow (13, 1, 3)$	$[7; 4, 5] (7, 18, 23) \rightarrow (12, 3, 1)$	$[7; 5, 1] (7, 19, 20) \rightarrow (11, 3, 1)$
$[7; 5, 3] (7, 19, 22) \rightarrow (12, 2, 2)$	$[7; 5, 5] (7, 19, 24) \rightarrow (14, 1, 3)$	$[7; 6, 3] (7, 20, 23) \rightarrow (13, 2, 2)$
$[7; 6, 4] (7, 20, 24) \rightarrow (13, 3, 1)$	$[7; 6, 5] (7, 20, 25) \rightarrow (13, 2, 2)$	$[8; 1, 1] (8, 25, 26) \rightarrow (16, 1, 3)$
$[8; 1, 2] (8, 25, 27) \rightarrow (14, 2, 2)$	$[8; 1, 3] (8, 17, 20) \rightarrow (12, 3, 1)$	$[8; 1, 4] (8, 17, 21) \rightarrow (12, 2, 2)$
$[8; 1, 5] (8, 17, 22) \rightarrow (11, 3, 2)$	$[8; 3, 1] (8, 19, 20) \rightarrow (12, 3, 1)$	$[8; 3, 3] (8, 19, 22) \rightarrow (14, 1, 3)$
$[8; 3, 4] (8, 19, 23) \rightarrow (14, 2, 2)$	$[8; 3, 6] (8, 19, 25) \rightarrow (12, 2, 2)$	$[8; 5, 1] (8, 21, 22) \rightarrow (11, 3, 2)$
$[8; 5, 2] (8, 21, 23) \rightarrow (12, 2, 2)$	$[8; 5, 4] (8, 21, 25) \rightarrow (15, 2, 2)$	$[8; 5, 5] (8, 21, 26) \rightarrow (16, 1, 3)$
$[8; 7, 3] (8, 23, 26) \rightarrow (13, 3, 2)$	$[8; 7, 4] (8, 23, 27) \rightarrow (16, 2, 2)$	$[8; 7, 5] (8, 15, 20) \rightarrow (11, 3, 1)$
$[8; 7, 6] (8, 15, 21) \rightarrow (10, 2, 2)$	$[9; 1, 1] (9, 28, 29) \rightarrow (16, 1, 4)$	$[9; 1, 2] (9, 19, 21) \rightarrow (12, 2, 2)$
$[9; 1, 3] (9, 19, 22) \rightarrow (12, 3, 2)$	$[9; 1, 4] (9, 19, 23) \rightarrow (12, 4, 1)$	$[9; 1, 5] (9, 19, 24) \rightarrow (13, 2, 2)$
$[9; 1, 6] (9, 19, 25) \rightarrow (12, 4, 2)$	$[9; 2, 1] (9, 20, 21) \rightarrow (12, 2, 2)$	$[9; 2, 2] (9, 20, 22) \rightarrow (13, 1, 4)$
$[9; 2, 3] (9, 20, 23) \rightarrow (13, 2, 3)$	$[9; 2, 4] (9, 20, 24) \rightarrow (13, 2, 2)$	$[9; 2, 6] (9, 20, 26) \rightarrow (13, 3, 2)$
$[9; 4, 2] (9, 22, 24) \rightarrow (13, 2, 2)$	$[9; 4, 3] (9, 22, 25) \rightarrow (13, 3, 2)$	$[9; 4, 4] (9, 22, 26) \rightarrow (15, 1, 4)$
$[9; 4, 6] (9, 22, 28) \rightarrow (13, 4, 2)$	$[9; 4, 7] (9, 22, 29) \rightarrow (14, 4, 1)$	$[9; 5, 1] (9, 23, 24) \rightarrow (13, 2, 2)$
$[9; 5, 2] (9, 23, 25) \rightarrow (14, 4, 1)$	$[9; 5, 3] (9, 23, 26) \rightarrow (14, 2, 3)$	$[9; 5, 5] (9, 23, 28) \rightarrow (16, 1, 4)$
$[9; 5, 6] (9, 23, 29) \rightarrow (14, 3, 2)$	$[9; 5, 7] (9, 14, 21) \rightarrow (10, 2, 2)$	$[9; 7, 1] (9, 25, 26) \rightarrow (14, 4, 1)$
$[9; 7, 3] (9, 25, 28) \rightarrow (15, 3, 2)$	$[9; 7, 5] (9, 16, 21) \rightarrow (11, 2, 2)$	$[9; 7, 6] (9, 16, 22) \rightarrow (10, 4, 2)$
$[9; 7, 7] (9, 25, 32) \rightarrow (18, 1, 4)$	$[9; 8, 3] (9, 26, 29) \rightarrow (16, 2, 3)$	$[9; 8, 4] (9, 17, 21) \rightarrow (11, 2, 2)$

Table 2: (continued).

$[9; 8, 5] (9, 17, 22) \rightarrow (11, 4, 1)$	$[9; 8, 6] (9, 17, 23) \rightarrow (11, 3, 2)$	$[9; 8, 7] (9, 17, 24) \rightarrow (12, 2, 2)$
$[10; 1, 1] (10, 21, 22) \rightarrow (13, 1, 4)$	$[10; 1, 2] (10, 21, 23) \rightarrow (11, 2, 3)$	$[10; 1, 3] (10, 21, 24) \rightarrow (11, 3, 2)$
$[10; 1, 4] (10, 21, 25) \rightarrow (14, 4, 1)$	$[10; 1, 5] (10, 21, 26) \rightarrow (14, 3, 2)$	$[10; 1, 6] (10, 21, 27) \rightarrow (13, 3, 2)$
$[10; 1, 7] (10, 21, 28) \rightarrow (14, 3, 3)$	$[10; 3, 1] (10, 23, 24) \rightarrow (12, 3, 3)$	$[10; 3, 2] (10, 23, 25) \rightarrow (15, 4, 1)$
$[10; 3, 3] (10, 23, 26) \rightarrow (16, 1, 4)$	$[10; 3, 5] (10, 23, 28) \rightarrow (16, 3, 2)$	$[10; 3, 6] (10, 23, 29) \rightarrow (14, 2, 3)$
$[10; 3, 8] (10, 23, 31) \rightarrow (14, 3, 2)$	$[10; 7, 1] (10, 27, 28) \rightarrow (14, 3, 2)$	$[10; 7, 2] (10, 27, 29) \rightarrow (14, 3, 2)$
$[10; 7, 4] (10, 27, 31) \rightarrow (15, 2, 3)$	$[10; 7, 5] (10, 17, 22) \rightarrow (12, 3, 2)$	$[10; 7, 7] (10, 27, 34) \rightarrow (20, 1, 4)$
$[10; 7, 8] (10, 17, 25) \rightarrow (12, 4, 1)$	$[10; 9, 3] (10, 19, 22) \rightarrow (11, 3, 3)$	$[10; 9, 4] (10, 19, 23) \rightarrow (11, 3, 2)$
$[10; 9, 5] (10, 19, 24) \rightarrow (13, 3, 2)$	$[10; 9, 6] (10, 19, 25) \rightarrow (13, 4, 1)$	$[10; 9, 7] (10, 19, 26) \rightarrow (11, 3, 2)$
$[10; 9, 8] (10, 19, 27) \rightarrow (12, 2, 3)$	$[11; 1, 1] (11, 23, 24) \rightarrow (13, 1, 5)$	$[11; 1, 2] (11, 23, 25) \rightarrow (12, 2, 3)$
$[11; 1, 3] (11, 23, 26) \rightarrow (12, 3, 2)$	$[11; 1, 4] (11, 23, 27) \rightarrow (13, 3, 4)$	$[11; 1, 5] (11, 23, 28) \rightarrow (14, 5, 1)$
$[11; 1, 6] (11, 23, 29) \rightarrow (14, 2, 4)$	$[11; 1, 7] (11, 23, 30) \rightarrow (13, 4, 2)$	$[11; 1, 8] (11, 23, 31) \rightarrow (14, 4, 3)$
$[11; 2, 1] (11, 24, 25) \rightarrow (12, 2, 4)$	$[11; 2, 2] (11, 24, 26) \rightarrow (15, 1, 5)$	$[11; 2, 3] (11, 24, 27) \rightarrow (12, 4, 2)$
$[11; 2, 4] (11, 24, 28) \rightarrow (13, 2, 3)$	$[11; 2, 5] (11, 24, 29) \rightarrow (13, 4, 3)$	$[11; 2, 6] (11, 24, 30) \rightarrow (13, 3, 2)$
$[11; 2, 8] (11, 24, 32) \rightarrow (15, 4, 2)$	$[11; 3, 1] (11, 25, 26) \rightarrow (12, 3, 4)$	$[11; 3, 2] (11, 25, 27) \rightarrow (12, 4, 3)$
$[11; 3, 3] (11, 25, 28) \rightarrow (16, 1, 5)$	$[11; 3, 4] (11, 25, 29) \rightarrow (15, 5, 1)$	$[11; 3, 6] (11, 25, 31) \rightarrow (14, 2, 3)$
$[11; 3, 7] (11, 25, 32) \rightarrow (15, 2, 4)$	$[11; 3, 9] (11, 14, 23) \rightarrow (9, 3, 2)$	$[11; 4, 1] (11, 26, 27) \rightarrow (12, 3, 2)$
$[11; 4, 2] (11, 26, 28) \rightarrow (13, 2, 4)$	$[11; 4, 4] (11, 26, 30) \rightarrow (17, 1, 5)$	$[11; 4, 5] (11, 26, 31) \rightarrow (14, 3, 4)$
$[11; 4, 6] (11, 26, 32) \rightarrow (14, 4, 2)$	$[11; 4, 8] (11, 15, 23) \rightarrow (10, 2, 3)$	$[11; 4, 9] (11, 15, 24) \rightarrow (10, 5, 1)$
$[11; 5, 2] (11, 27, 29) \rightarrow (13, 4, 2)$	$[11; 5, 3] (11, 27, 30) \rightarrow (16, 5, 1)$	$[11; 5, 4] (11, 27, 31) \rightarrow (14, 3, 2)$
$[11; 5, 5] (11, 27, 32) \rightarrow (18, 1, 5)$	$[11; 5, 7] (11, 16, 23) \rightarrow (10, 4, 3)$	$[11; 5, 8] (11, 16, 24) \rightarrow (11, 3, 2)$
$[11; 5, 9] (11, 16, 25) \rightarrow (11, 4, 2)$	$[11; 6, 1] (11, 28, 29) \rightarrow (13, 2, 3)$	$[11; 6, 2] (11, 28, 30) \rightarrow (14, 3, 4)$
$[11; 6, 3] (11, 28, 31) \rightarrow (14, 2, 4)$	$[11; 6, 4] (11, 28, 32) \rightarrow (15, 4, 3)$	$[11; 6, 6] (11, 28, 34) \rightarrow (19, 1, 5)$
$[11; 6, 7] (11, 17, 24) \rightarrow (10, 3, 2)$	$[11; 6, 8] (11, 17, 25) \rightarrow (11, 5, 1)$	$[11; 6, 9] (11, 17, 26) \rightarrow (11, 4, 2)$
$[11; 7, 1] (11, 29, 30) \rightarrow (14, 4, 3)$	$[11; 7, 2] (11, 29, 31) \rightarrow (17, 5, 1)$	$[11; 7, 3] (11, 29, 32) \rightarrow (15, 2, 3)$
$[11; 7, 5] (11, 18, 23) \rightarrow (10, 4, 2)$	$[11; 7, 6] (11, 18, 24) \rightarrow (11, 4, 2)$	$[11; 7, 7] (11, 29, 36) \rightarrow (20, 1, 5)$
$[11; 7, 9] (11, 18, 27) \rightarrow (13, 3, 2)$	$[11; 8, 1] (11, 30, 31) \rightarrow (14, 4, 2)$	$[11; 8, 2] (11, 30, 32) \rightarrow (15, 3, 2)$
$[11; 8, 4] (11, 19, 23) \rightarrow (11, 2, 4)$	$[11; 8, 5] (11, 19, 24) \rightarrow (11, 2, 3)$	$[11; 8, 7] (11, 19, 26) \rightarrow (12, 5, 1)$
$[11; 8, 8] (11, 30, 38) \rightarrow (21, 1, 5)$	$[11; 8, 9] (11, 19, 28) \rightarrow (12, 4, 3)$	$[11; 9, 1] (11, 31, 32) \rightarrow (17, 5, 1)$
$[11; 9, 3] (11, 20, 23) \rightarrow (11, 3, 4)$	$[11; 9, 5] (11, 20, 25) \rightarrow (11, 3, 2)$	$[11; 9, 6] (11, 20, 26) \rightarrow (12, 4, 3)$
$[11; 9, 7] (11, 20, 27) \rightarrow (12, 2, 3)$	$[11; 9, 8] (11, 20, 28) \rightarrow (12, 4, 2)$	$[11; 9, 9] (11, 31, 40) \rightarrow (22, 1, 5)$
$[11; 10, 3] (11, 21, 24) \rightarrow (11, 4, 3)$	$[11; 10, 4] (11, 21, 25) \rightarrow (11, 4, 2)$	$[11; 10, 5] (11, 21, 26) \rightarrow (12, 2, 4)$
$[11; 10, 6] (11, 21, 27) \rightarrow (13, 5, 1)$	$[11; 10, 7] (11, 21, 28) \rightarrow (13, 4, 2)$	$[11; 10, 8] (11, 21, 29) \rightarrow (12, 3, 2)$
$[11; 10, 9] (11, 21, 30) \rightarrow (13, 2, 3)$		

APPENDIX E. CORRESPONDENCE WITH THE LEAN DEVELOPMENT

Declarations live in the namespace `Erdos1112.Proof`, in files relative to `Erdos1112Proof/`, except the three headline theorems, which are in namespace `Erdos1112`. Every listed result is load-bearing: the development formalizes the full case tree rather than citing anything externally. The name-and-location layer of this table is checked mechanically (`scripts/check_correspondence.py`, re-run in continuous integration), so it cannot silently rot; the *statement-level* match — that each declaration says what the prose result says — remains a human judgement, and is where we invite scrutiny.

One row is marked ALT: for Case P this paper gives a single uniform argument for the whole line (§20) while the formalization instantiates small r separately and routes the rest through an auxiliary lemma. The formal statement proved is the one the assembly consumes; it is not weaker.

Result here	Lean declaration	File
Thm. A (dichotomy)	erdos_1112	Final.lean
Thm. 6.1	erdos_1112_existence_bound, existence_bound	Final.lean, Existence/Nested.lean
Strong non-existence	erdos_1112_strong_nonexistence	Final.lean
Integer form ($\mathbb{Z} \rightarrow \mathbb{N}$ bridge)	erdos_1112_int, question_iff_questionInt	Final.lean, Erdos1112.lean
Lem. 5.2 (free gap)	exists_safe_subinterval	Existence/FreeGap.lean
Lem. 8.6 (certificate)	cert_hit	NonEx/Certificate.lean
Lems. 9.3, 10.1	tailCoveringN_of_single_letter, ..._of_eventually_periodic	NonEx/GapWord.lean
Lem. 12.1 (the band)	Wset_interval	NonEx/TwoLetter/Core.lean
Lem. 12.2 (sweep)	sweep	NonEx/TwoLetter/Core.lean
Lem. 12.4 (width dichotomy)	width_of_unbalanced	NonEx/TwoLetter/Core.lean
Thm. 13.8 (balanced classification)	balanced_classification disc_osc_le_one, mechanical_tail	NonEx/TwoLetter/Balanced.lean NonEx/TwoLetter/MH/
Lem. 14.5 (boundary width)	width_even_boundary	NonEx/TwoLetter/Core.lean
Lem. 14.1 (Sturmian)	tailCovering_of_sturmian, walk_enters	NonEx/TwoLetter/Sturmian.lean NonEx/TwoLetter/MH/Walk.lean
Lem. 15.3 (slot lemma)	tailCovering_of_three_letters	NonEx/SlotLemma.lean
Lem. 17.1 (two-generator)	twoGen_interval	Sharp/TwoGen.lean
Lem. 17.3 (frame)	frame_lemma	Sharp/Frame.lean
Lem. 17.4 (staircase)	staircase_phase_base, staircase_phase_extended	Sharp/Staircase.lean
Lem. 17.6 (λ -lift)	FrameCert.lift	Sharp/Lift.lean
Lem. 18.2 (Graham reduction)	sharpAt_of_hardcore	Sharp/Graham.lean
Lem. 18.8 (minimal alphabets)	sharp_of_minimal	Sharp/Graham.lean
Lem. 19.1 (Case D)	caseD	Sharp/CaseD.lean
Lem. 20.1 (Case P) ALT	caseP_pair, caseP_r3, caseP_r4, caseP_large	Sharp/CaseP.lean
Lem. 19.3 (Case L)	caseL	Sharp/CaseL.lean
Lem. 21.1 (Case E)	caseE	Sharp/CaseE.lean
Lem. 22.1, Prop. 22.3 (Case T)	caseT, T_tail_line, tSuppT	Sharp/CaseT*.lean
Case B (+ descent)	caseB, rowFor, caseBComplete	Sharp/CaseB.lean, Sharp/CaseBClasses.lean
Tables A/B (data)	certTableA_ok, certTableB_ok	Sharp/TablesData.lean
Thm. 16.1 (subset-sum lemma)	sharp	Sharp/Main.lean
Axiom audit	three #print axioms commands	AxiomsCheck.lean

REFERENCES

- [AF88] N. Alon and G. Freiman, *On sums of subsets of a set of integers*, *Combinatorica* **8** (1988), no. 4, 297–306.
- [BHJ97] B. Bollobás, N. Hegyvári and G. Jin, *On a problem of Erdős and Graham*, *Discrete Math.* **175** (1997), no. 1–3, 253–257.
- [CFP21] D. Conlon, J. Fox and H. T. Pham, *Subset sums, completeness and colorings*, preprint, arXiv:2104.14766 (2021).
- [CMZ25] L. Chen, Y. Mao and G. Zhang, *Long arithmetic progressions in sumsets and subset sums: constructive proofs and efficient witnesses*, in: *Proc. 57th Annual ACM Symposium on Theory of Computing (STOC 2025)*, ACM, 2025, 2086–2097.
- [Ch00] Y.-G. Chen, *On sums and intersects of sequences*, *Discrete Math.* **223** (2000), no. 1–3, 351–354.
- [E1112] T. F. Bloom, *Erdős problem #1112*, <https://www.erdosproblems.com/1112>. Accessed 13 July 2026.
- [ErGr80] P. Erdős and R. L. Graham, *Old and new problems and results in combinatorial number theory*, *Monographies de L’Enseignement Mathématique* **28**, Université de Genève, Geneva, 1980, 128 pp.
- [Gr64] R. L. Graham, *Complete sequences of polynomial values*, *Duke Math. J.* **31** (1964), no. 2, 275–285.

- [GS20] A. Granville and G. Shakan, *The Frobenius postage stamp problem, and beyond*, Acta Math. Hungar. **161** (2020), no. 2, 700–718.
- [GW21] A. Granville and A. Walker, *A tight structure theorem for sumsets*, Proc. Amer. Math. Soc. **149** (2021), no. 10, 4073–4082.
- [Lean4] L. de Moura and S. Ullrich, *The Lean 4 theorem prover and programming language*, in: Automated Deduction — CADE 28, Lecture Notes in Computer Science **12699**, Springer, 2021, 625–635.
- [Lev98] V. F. Lev, *On consecutive subset sums*, Discrete Math. **187** (1998), no. 1–3, 151–160.
- [Lev03] V. F. Lev, *Blocks and progressions in subset sum sets*, Acta Arith. **106** (2003), no. 2, 123–142.
- [Lev22] V. F. Lev, *The structure of higher sumsets*, Proc. Amer. Math. Soc. **150** (2022), no. 12, 5165–5177.
- [Lo02] M. Lothaire, *Algebraic Combinatorics on Words*, Encyclopedia of Mathematics and its Applications **90**, Cambridge University Press, 2002.
- [Mathlib] The mathlib Community, *The Lean mathematical library*, in: Proc. 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2020), 367–381.
- [MH40] M. Morse and G. A. Hedlund, *Symbolic dynamics II. Sturmian trajectories*, Amer. J. Math. **62** (1940), no. 1, 1–42.
- [Na72] M. B. Nathanson, *Sums of finite sets of integers*, Amer. Math. Monthly **79** (1972), no. 9, 1010–1012.
- [RA05] J. L. Ramírez Alfonsín, *The Diophantine Frobenius Problem*, Oxford Lecture Ser. Math. Appl. **30**, Oxford Univ. Press, 2005.
- [Sa89] A. Sárközy, *Finite addition theorems, I*, J. Number Theory **32** (1989), no. 1, 114–130.
- [Sa94] A. Sárközy, *Finite addition theorems, II*, J. Number Theory **48** (1994), no. 2, 197–218.
- [SV06] E. Szemerédi and V. H. Vu, *Finite and infinite arithmetic progressions in sumsets*, Ann. of Math. (2) **163** (2006), no. 1, 1–35.
- [TaYa21] M. Tang and Q.-H. Yang, *On a problem of Erdős and Graham*, Publ. Math. Debrecen **99** (2021), no. 3–4, 485–493.

INDEPENDENT RESEARCHER

Email address: `johan.land@gmail.com`